

# Istio- Powered Observability

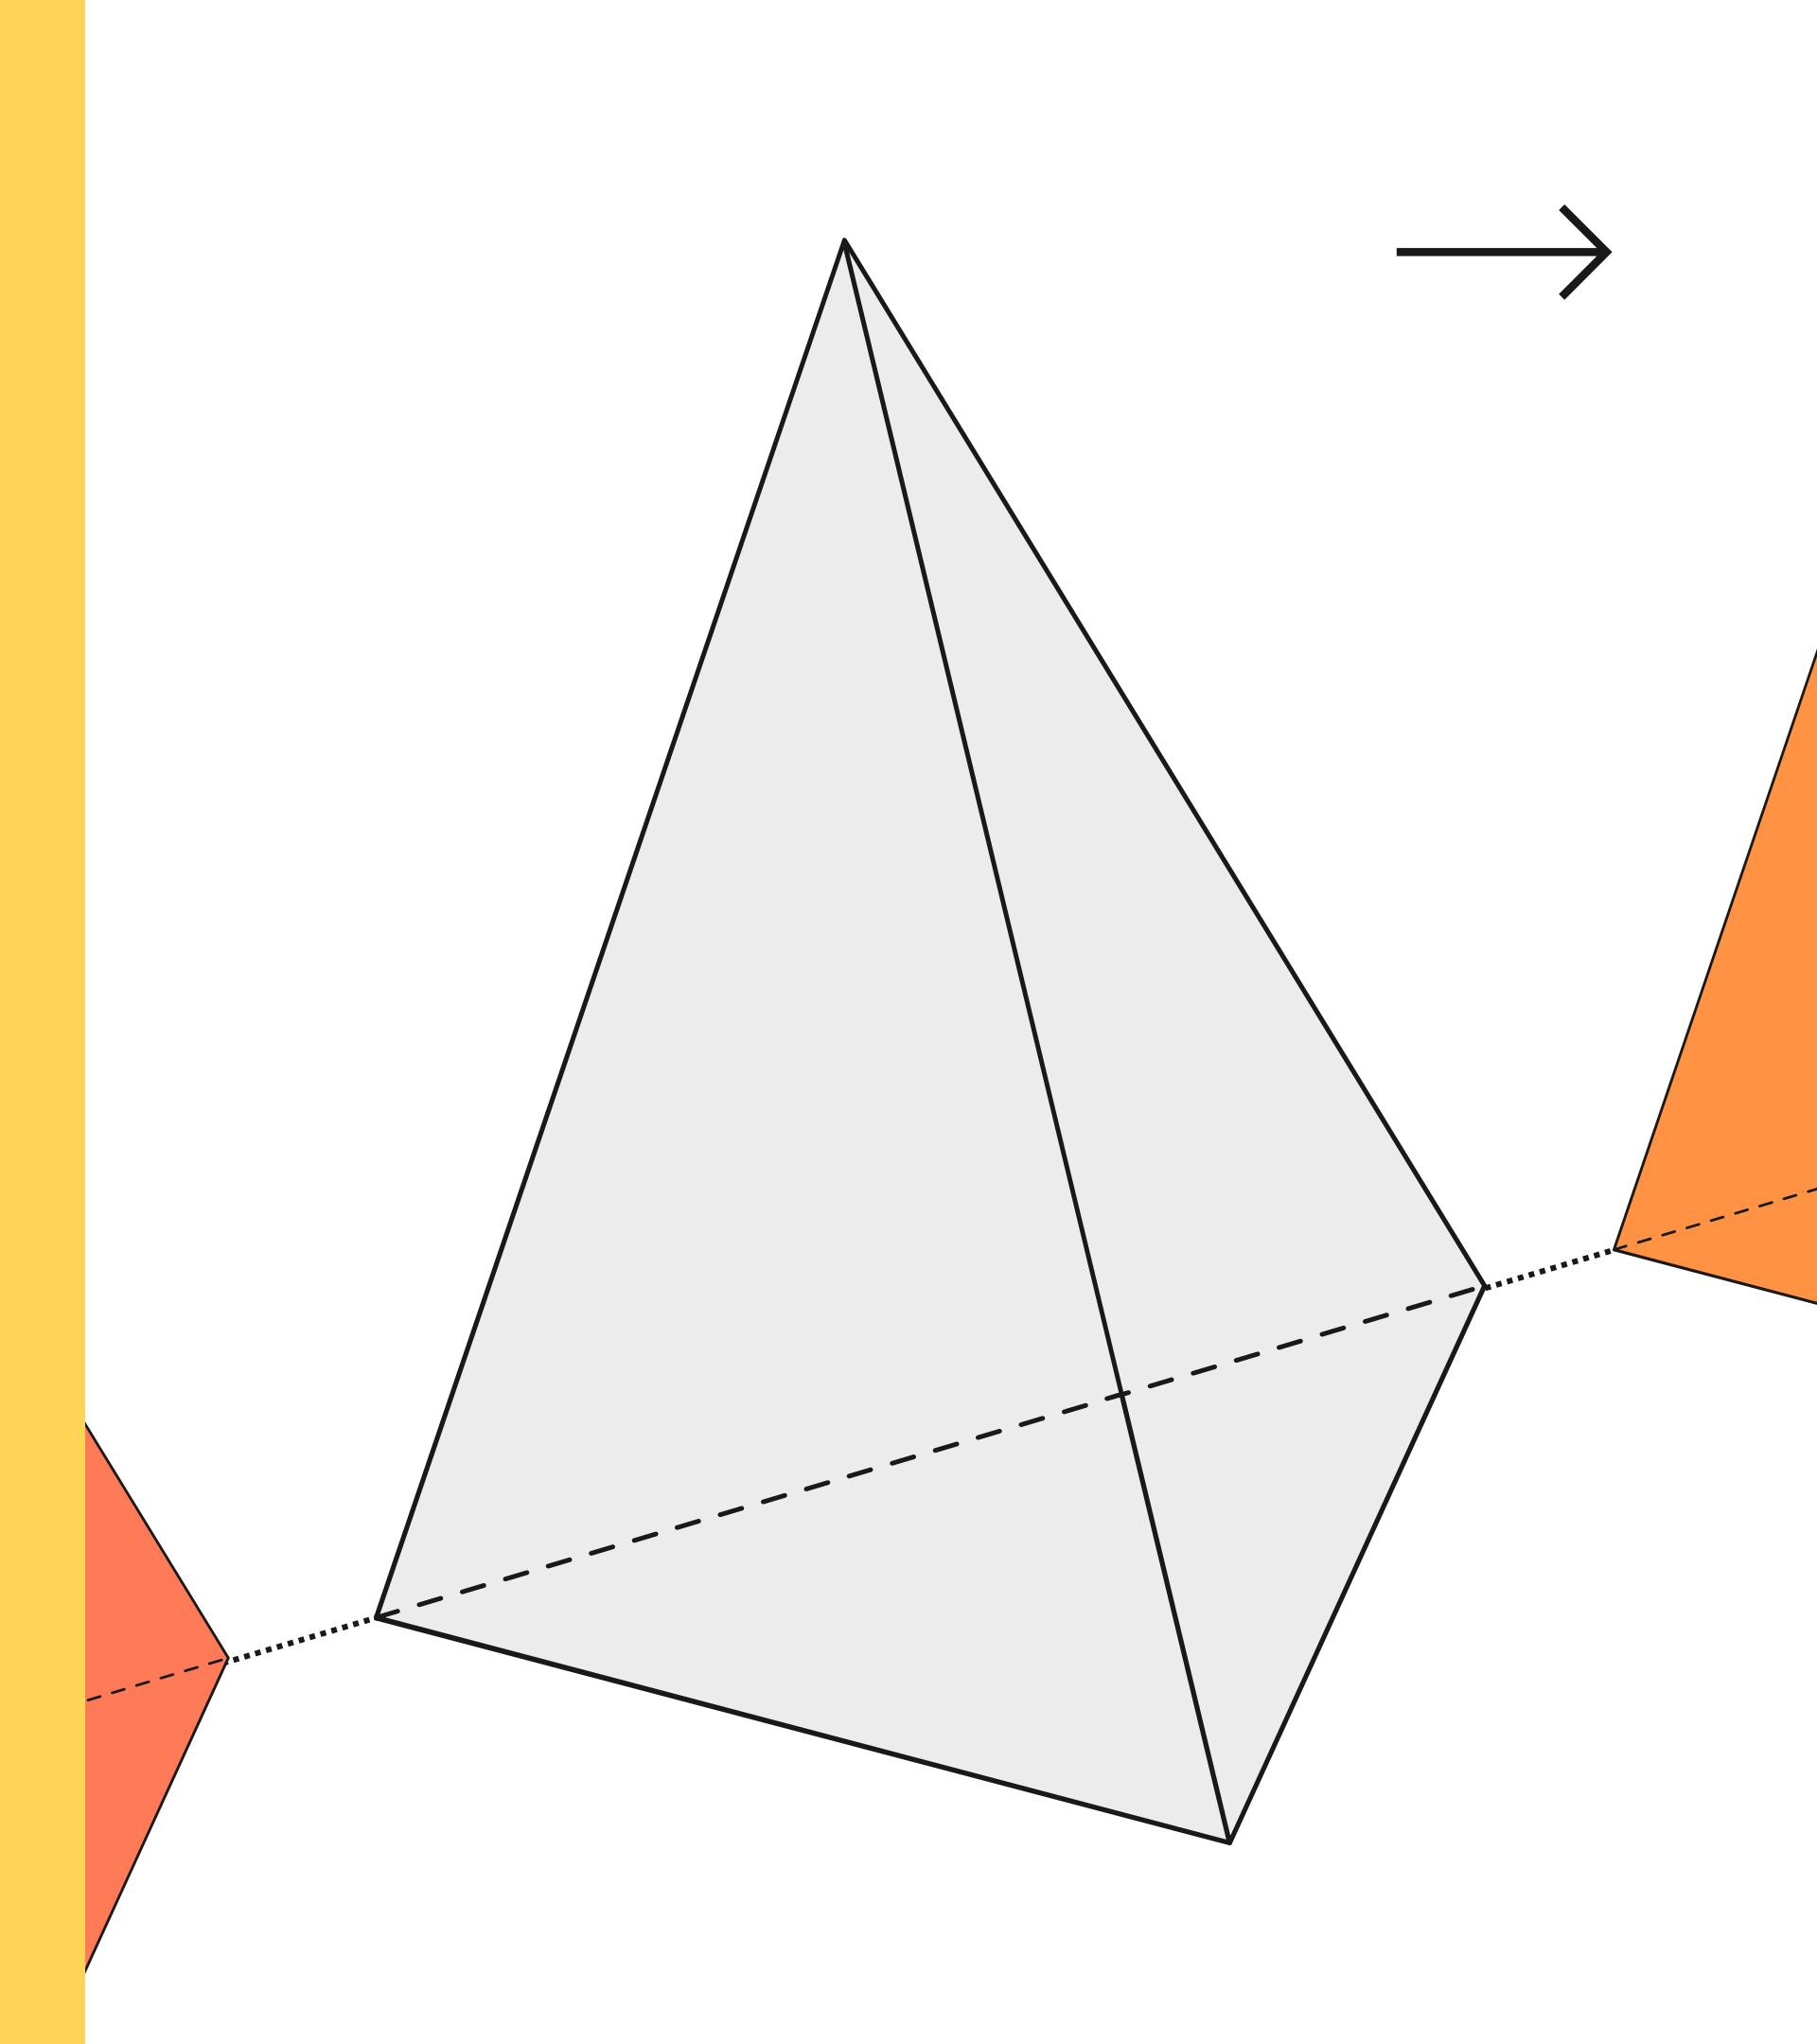
---

**NEERAJ PANDEY**

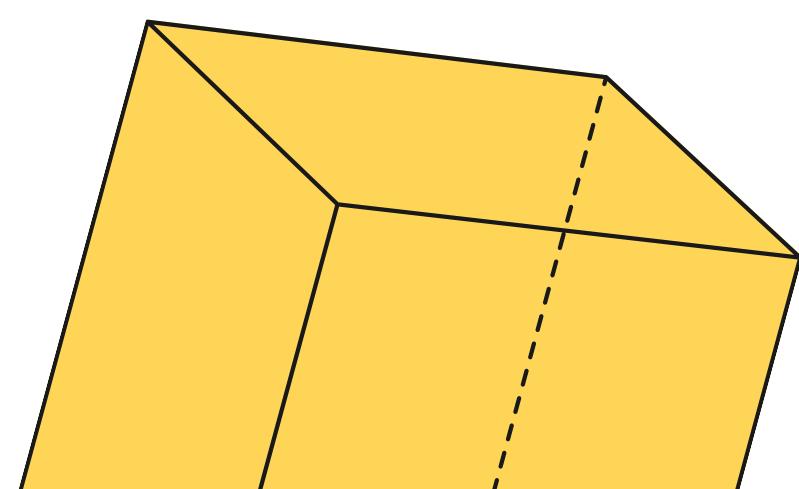
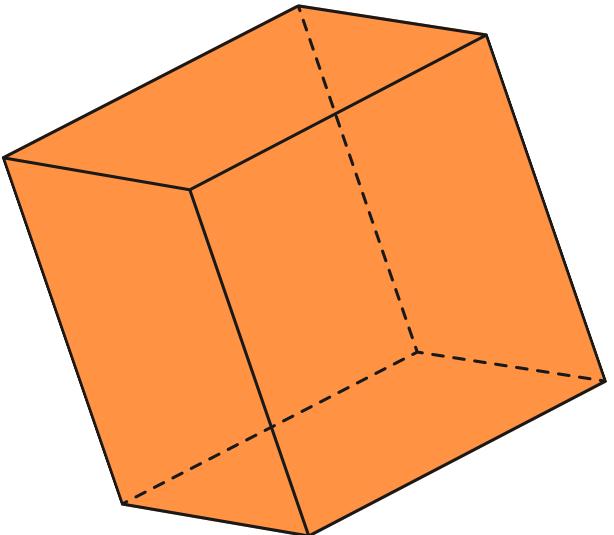
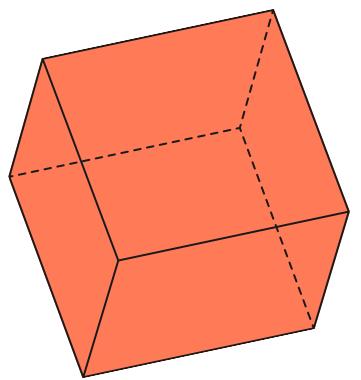
Founder Vivid Climate & LyntCube

---

**DEVOPSDAYS TW 2024**



# WHAT ARE WE LEARNING?



- Monoliths to Microservices
- Challenge of Orchestration
- Service Mesh
- Why and uses of Service Mesh
- Istio Service Mesh
- Istio and Observability Tools
- Implementation - Jaeger, Kiali, Grafana

# INTRODUCTION →

---

## NEERAJ PANDEY

<https://n4jp4y.com>

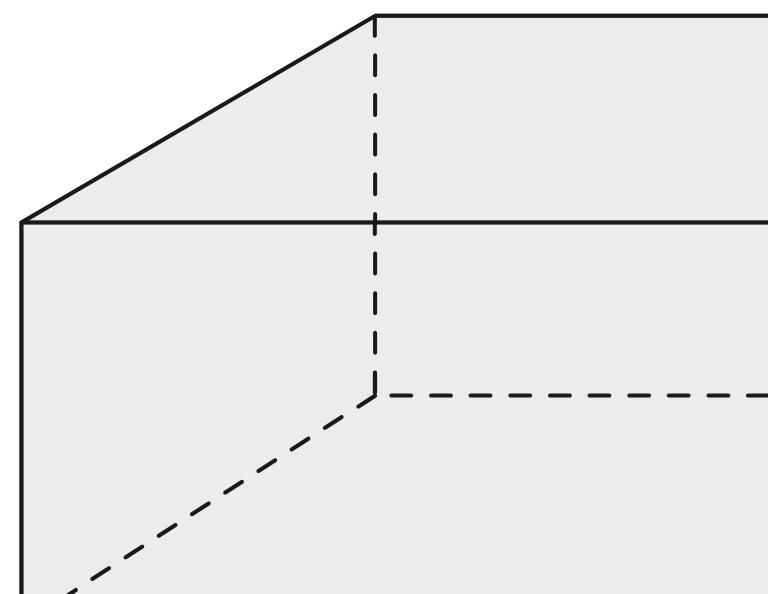
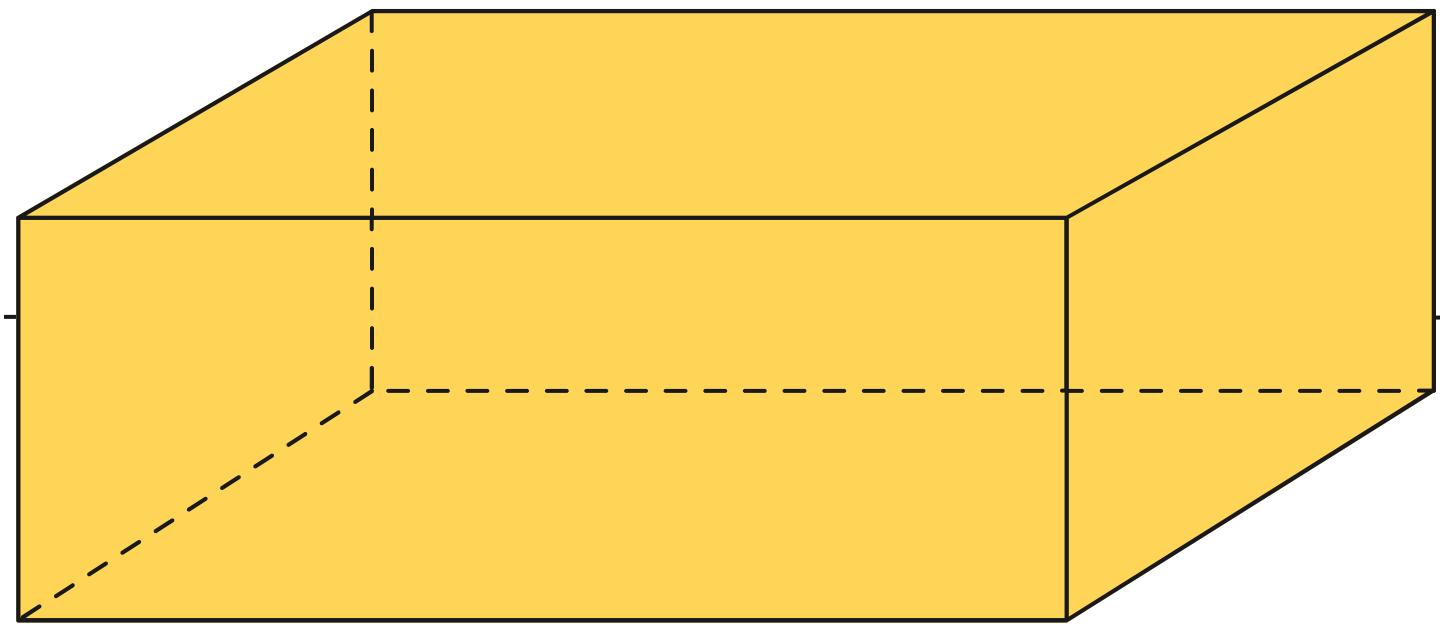
**n4jpy4** @twittter @linkedin

Vivid Climate - LyntCube Studio - Thrax Algo



# MONOLITHIC ERA

- Single codebase (Simplicity in Early Stages)
- Tight coupling of components
- Challenges with scalability and maintenance.



# Challenges of Monolithics

---

## SINGLE POINT OF FAILURE

Single issues can bring down the entire monolithic application, reducing resilience.

## DEPLOYMENT DILEMMA

Updating monoliths requires system-wide downtime, affecting user experience.

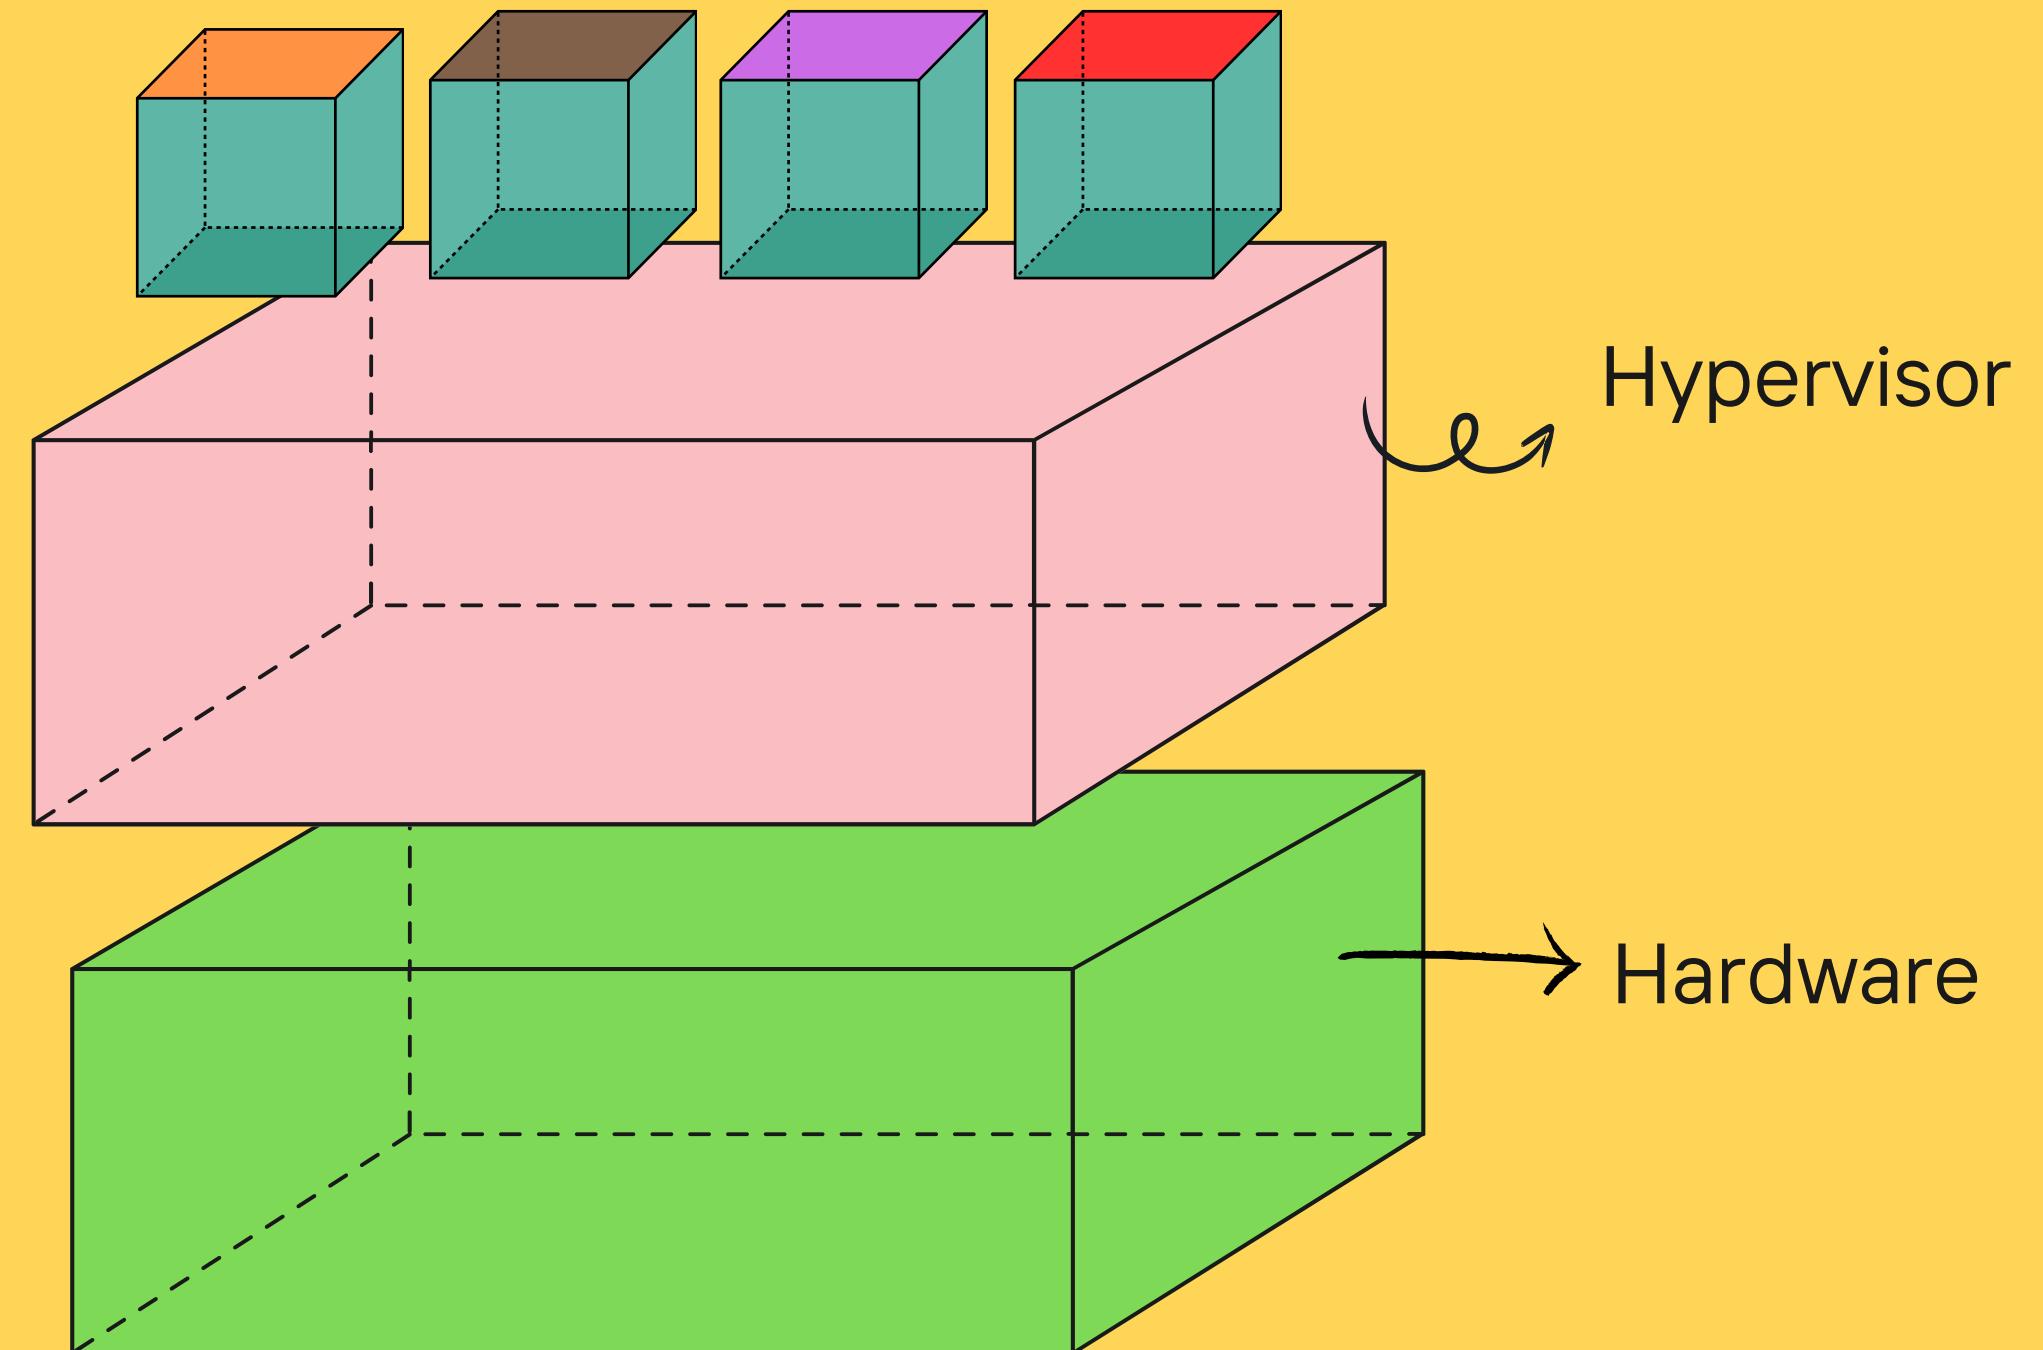
## SCALING STRUGGLES

Scaling monoliths is resource-intensive; can't scale individual components.

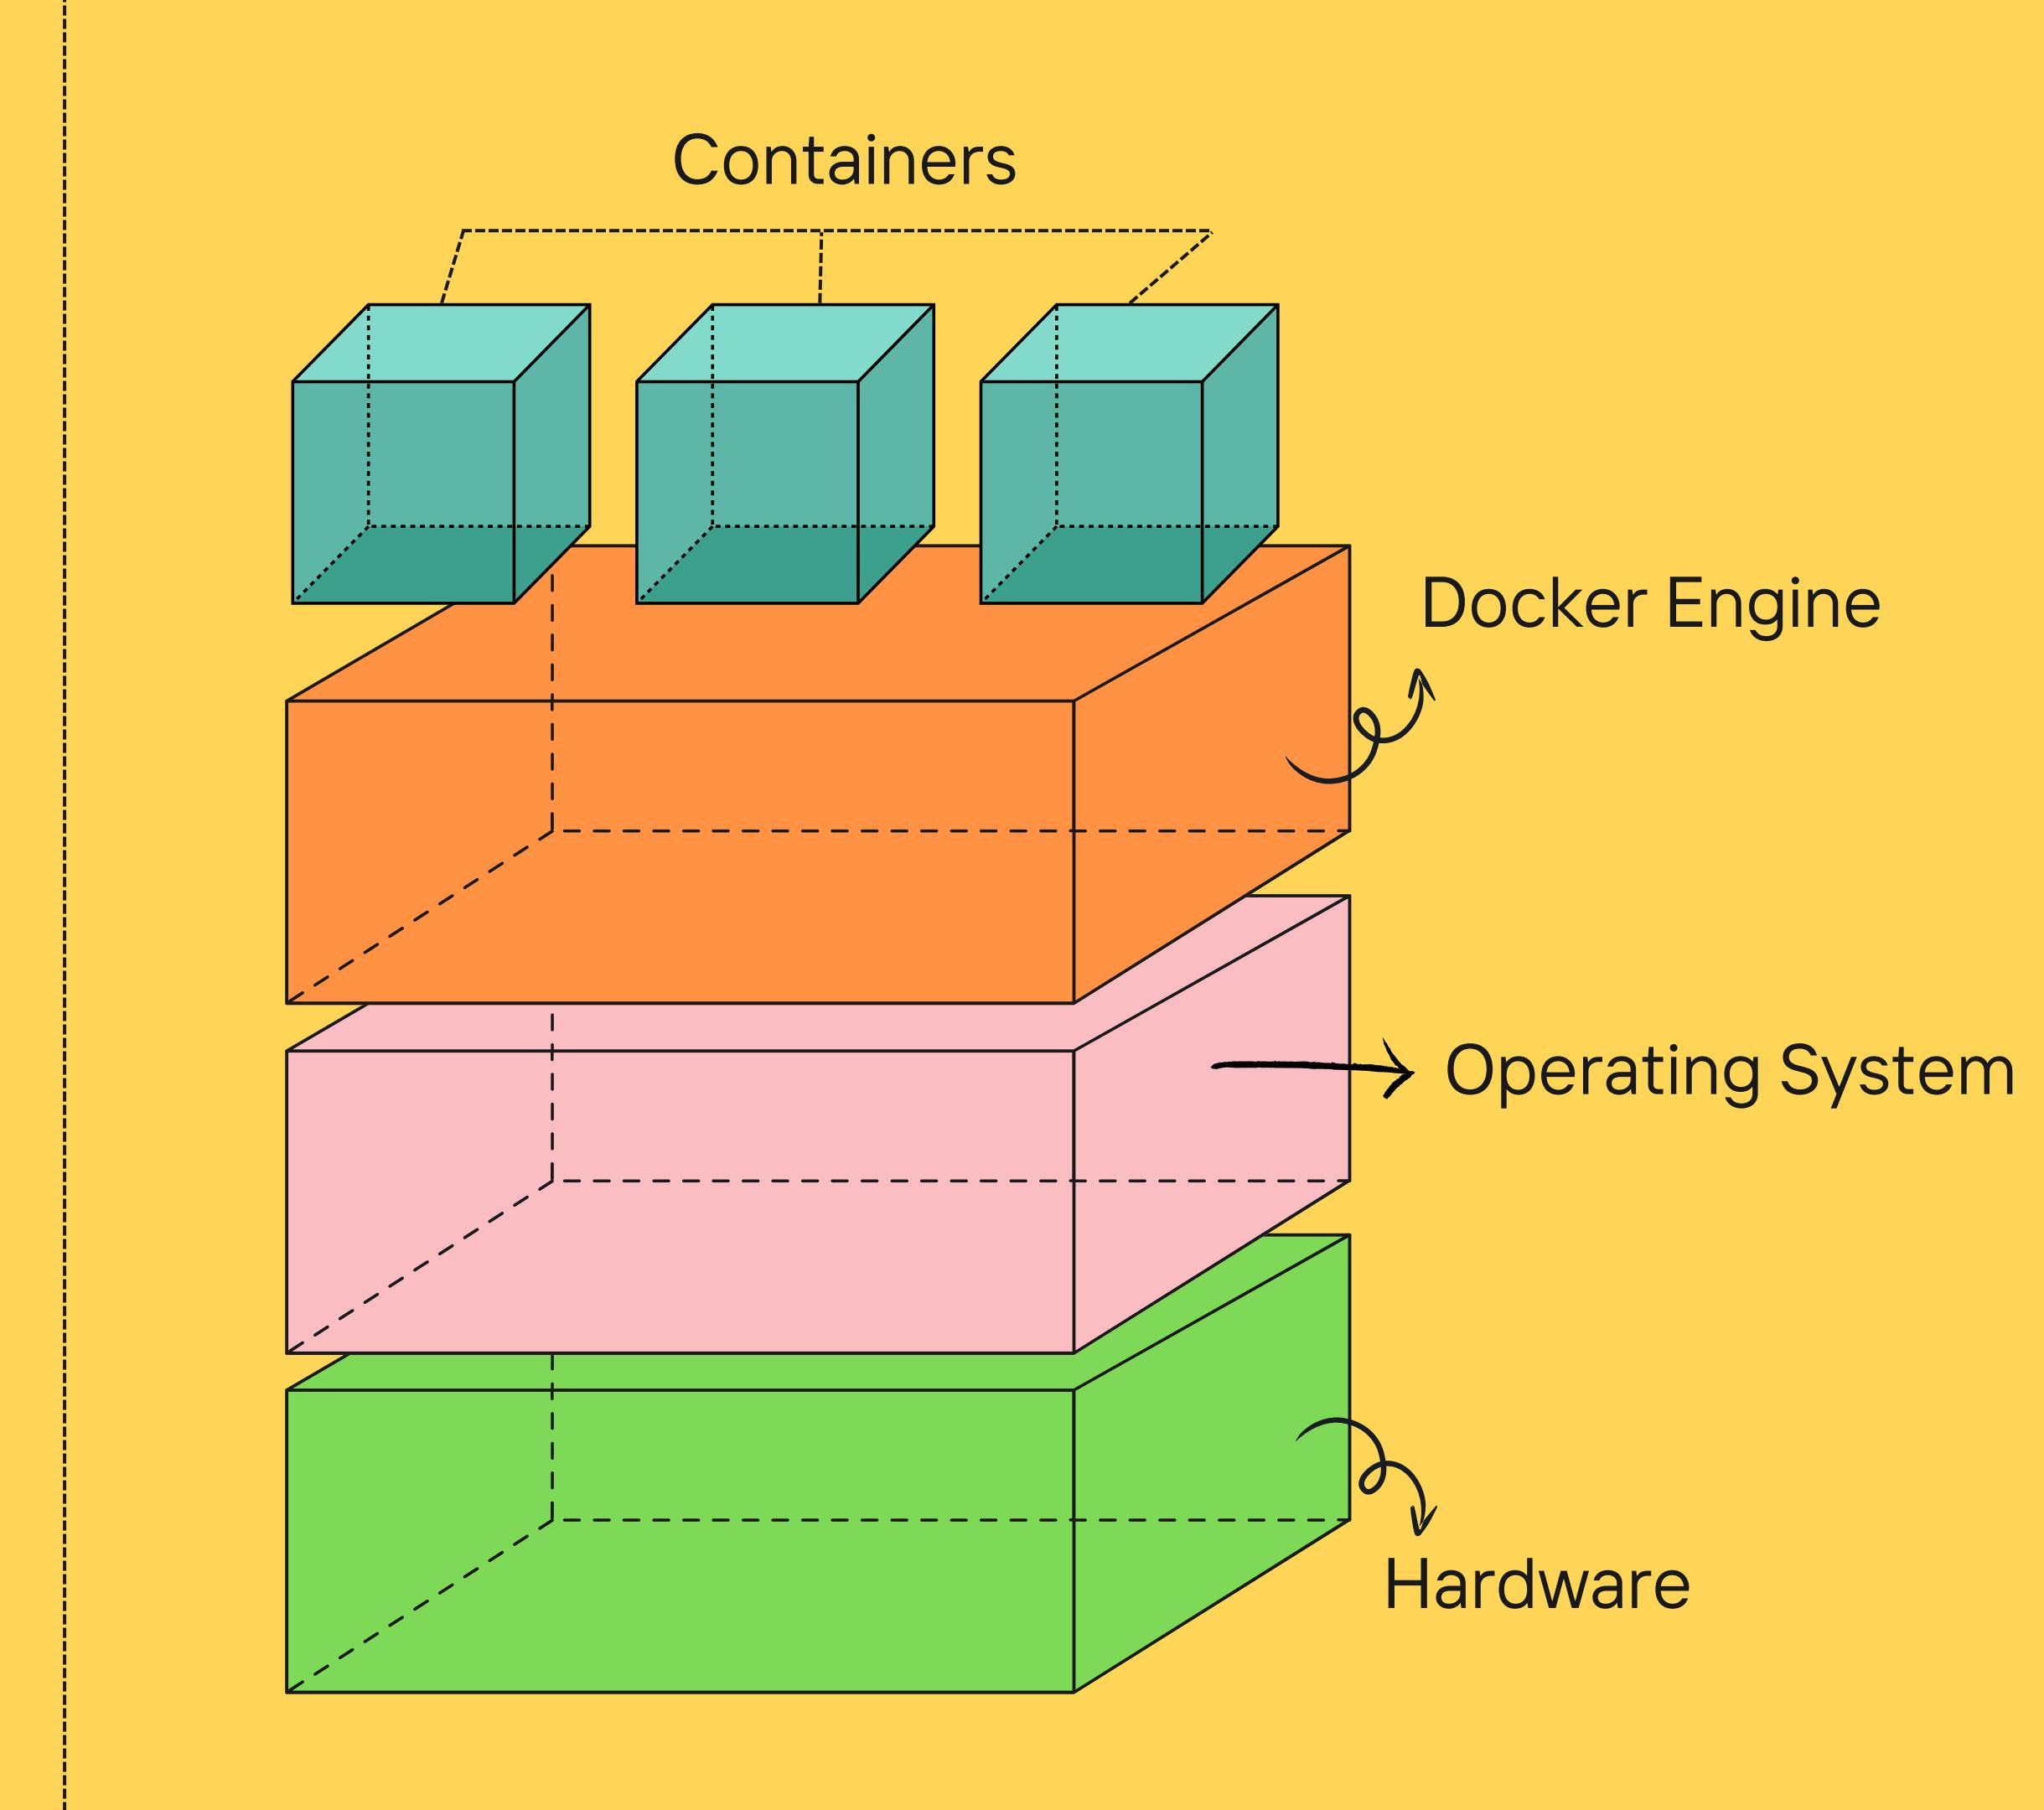
# VIRTUAL MACHINES



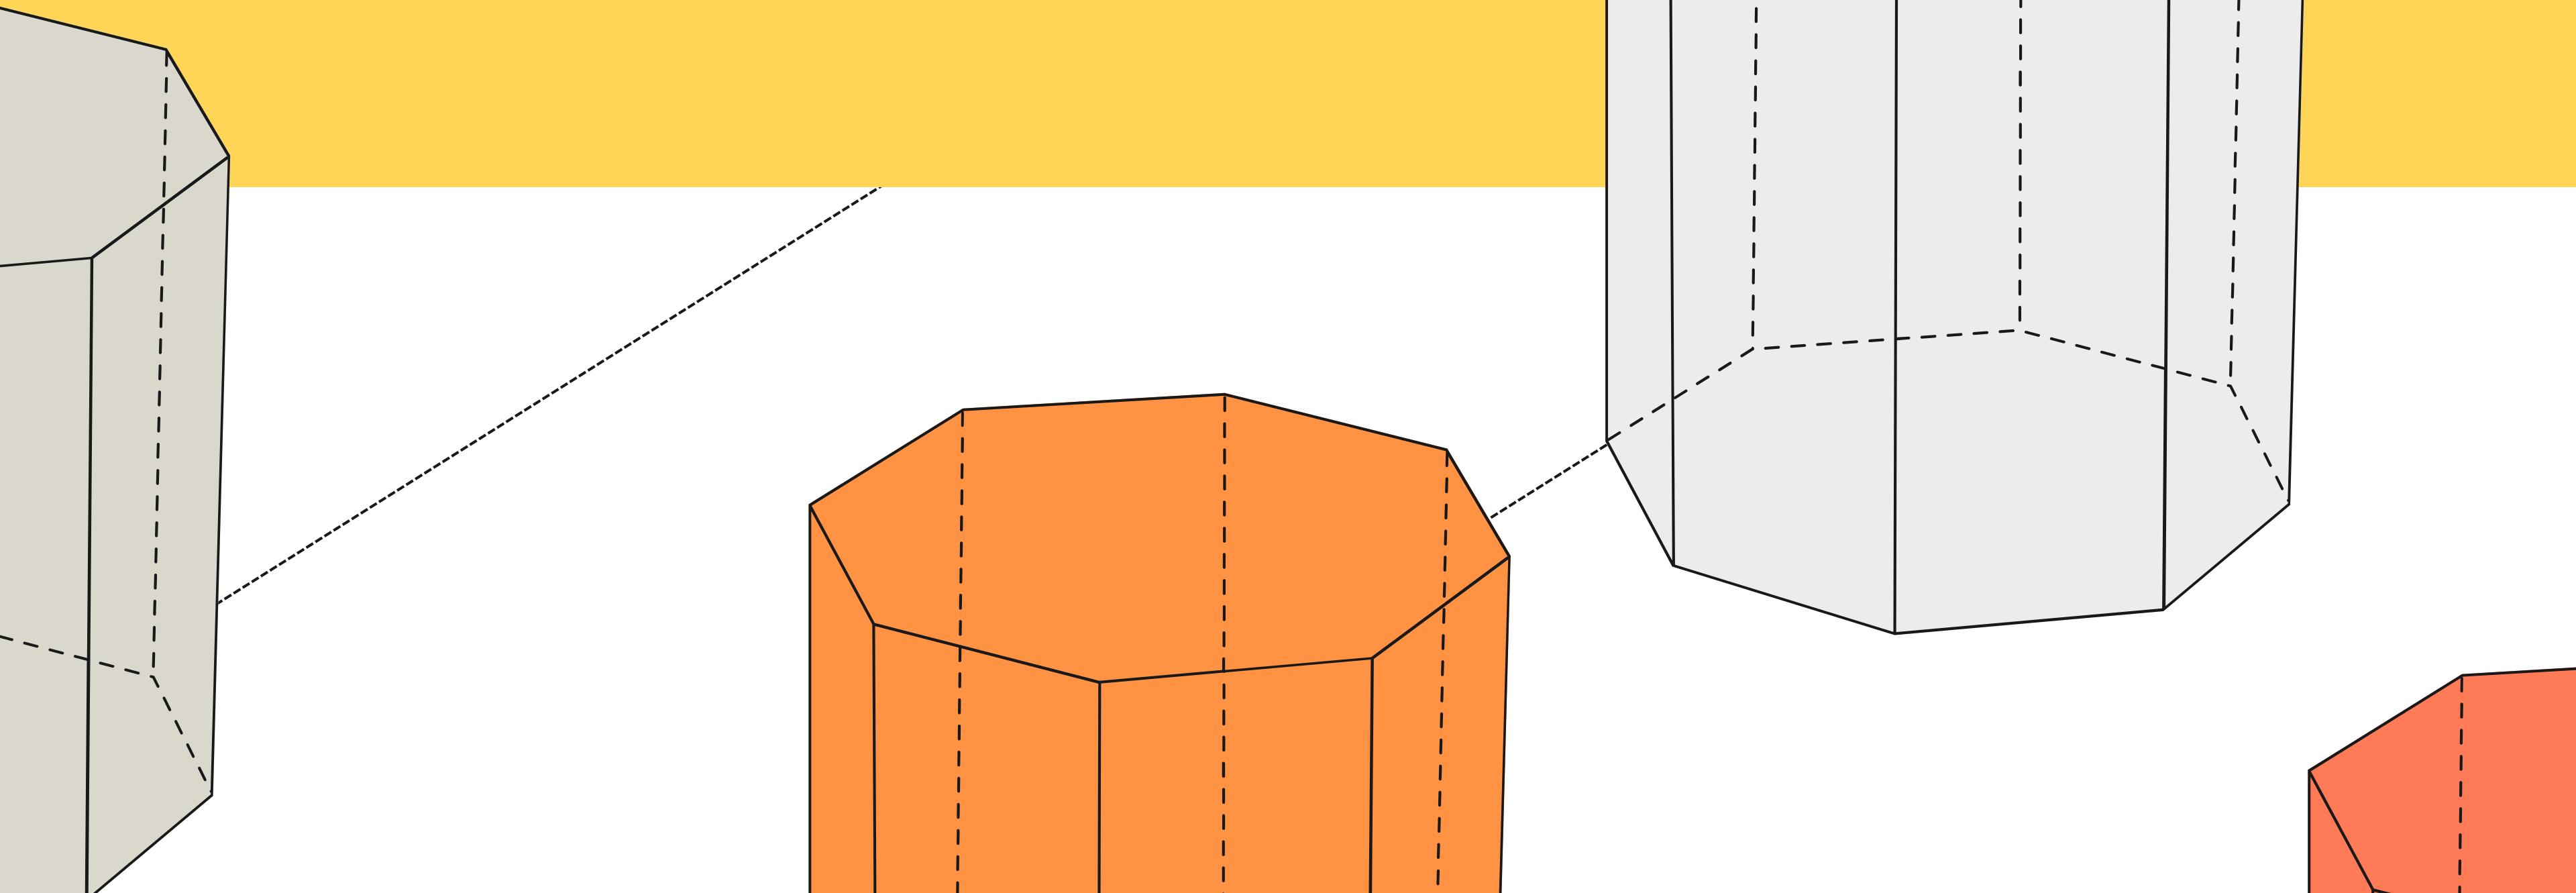
Operating Systems



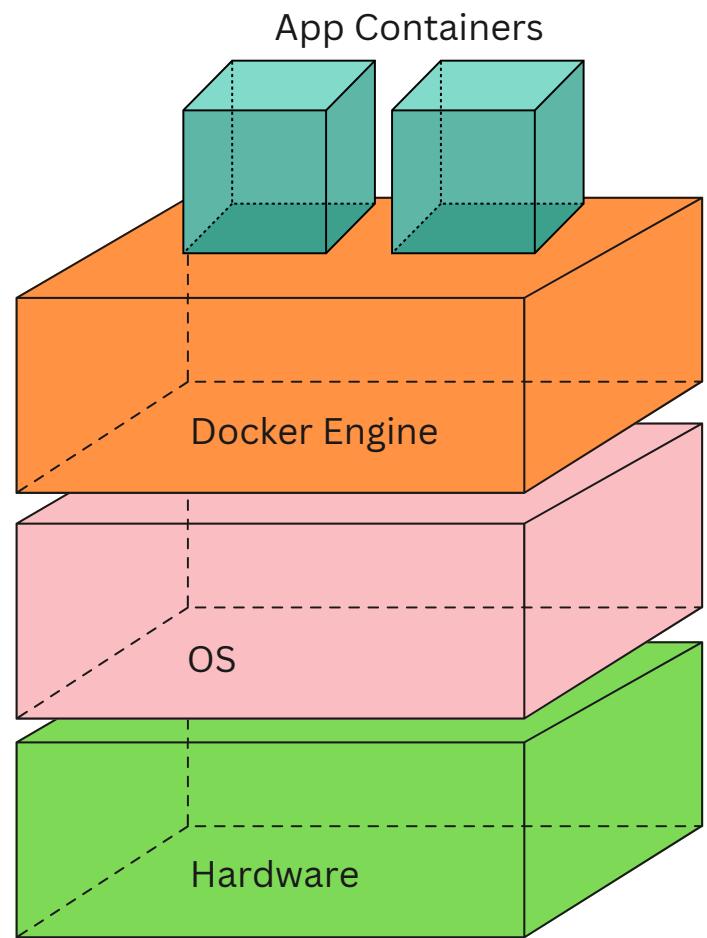
# DOCKER



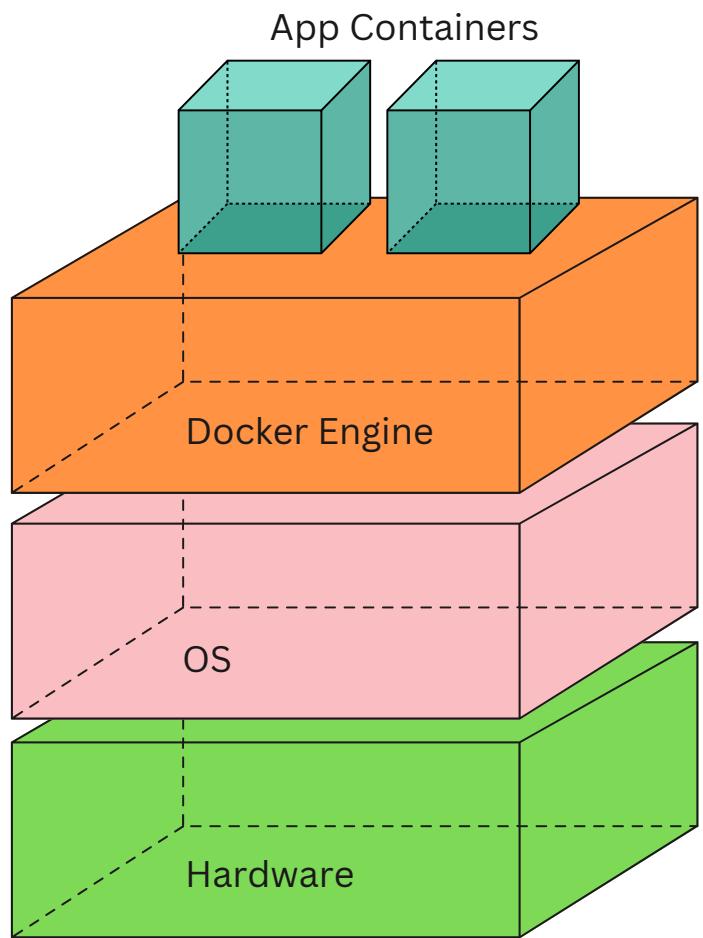
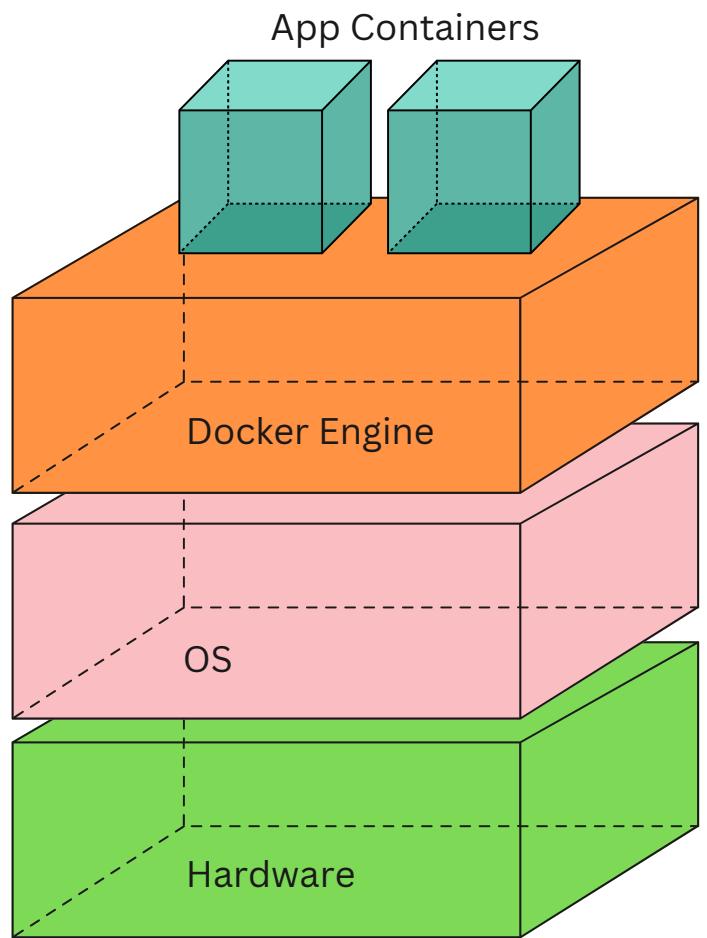
# DOCKER CHALLENGES



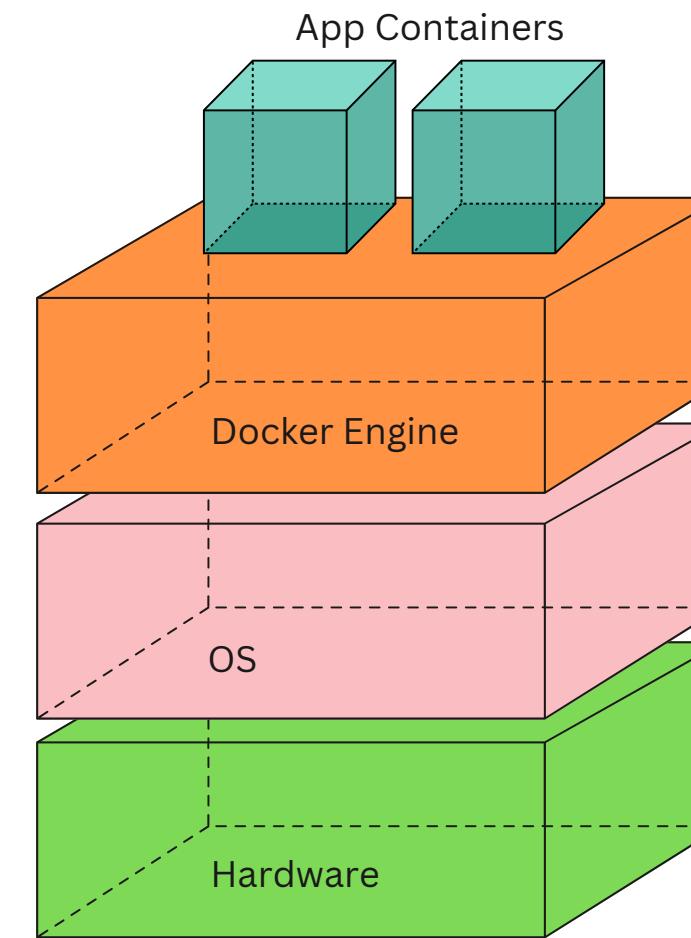
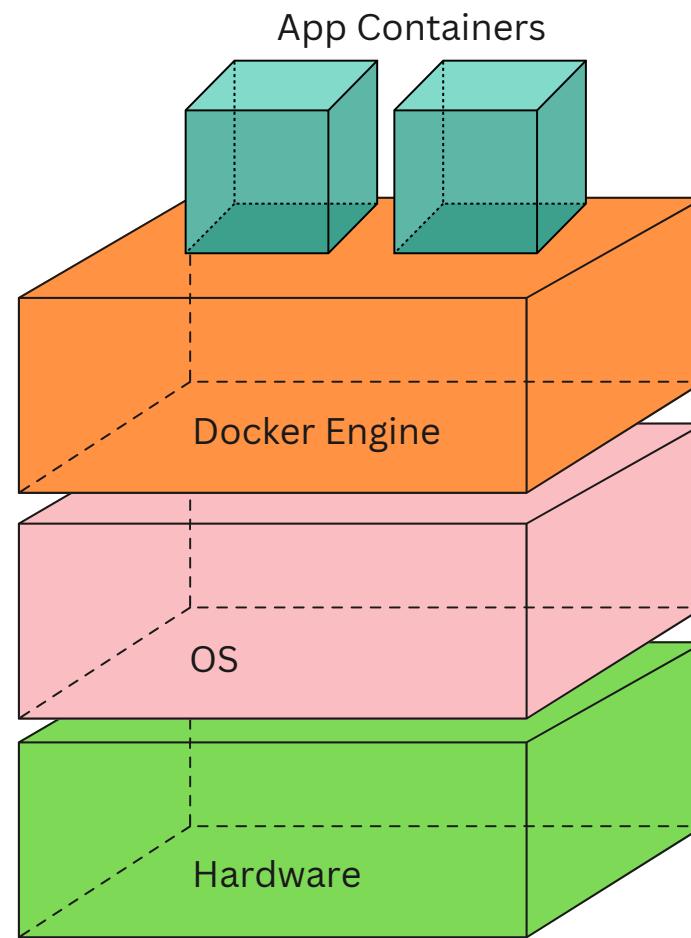
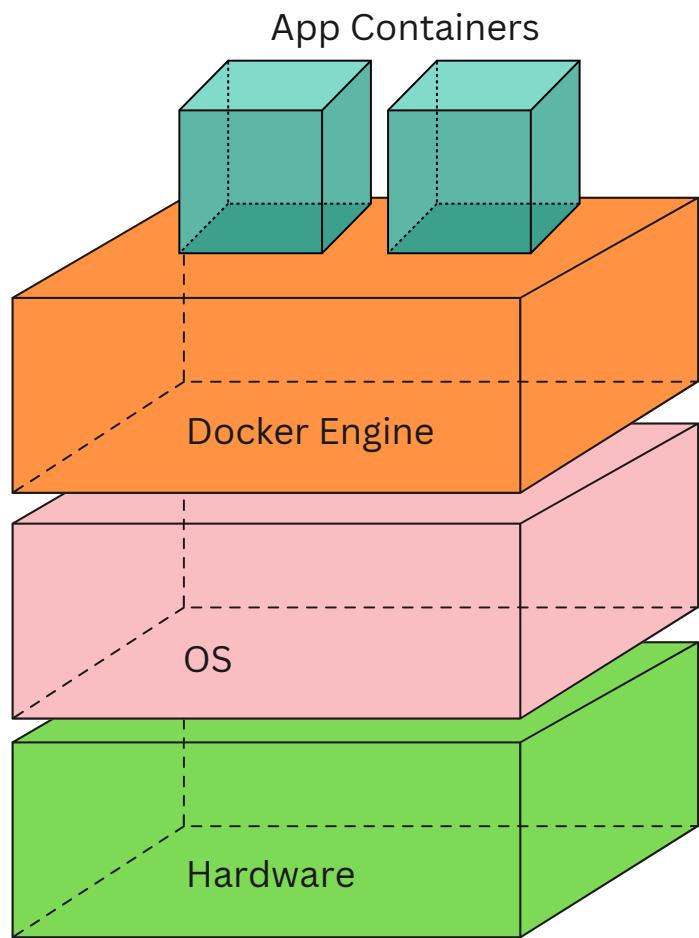
# Docker Challenges →



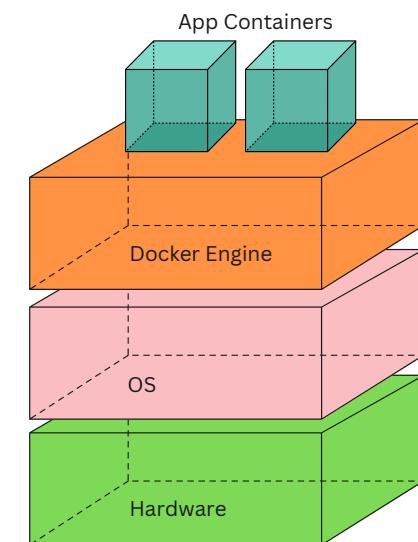
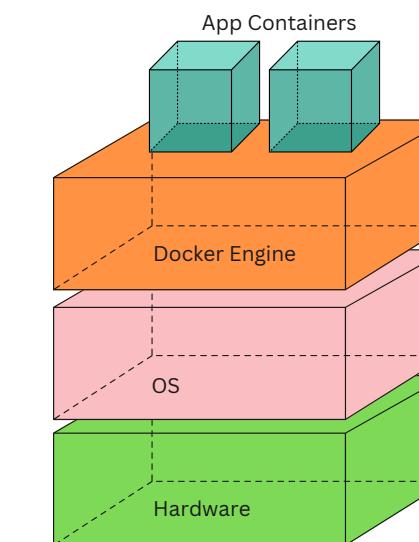
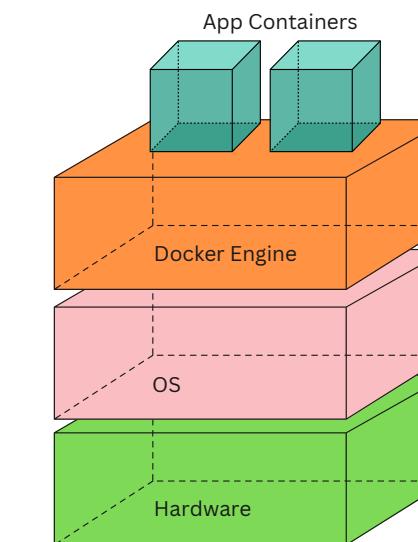
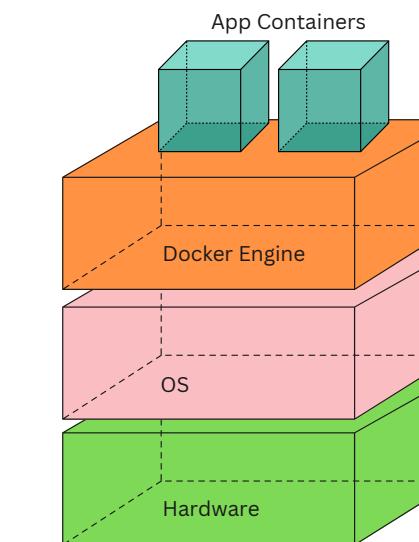
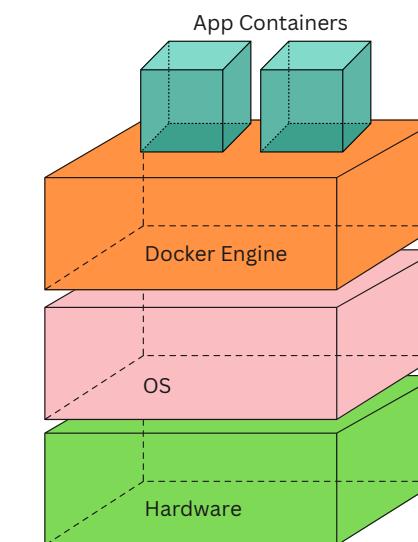
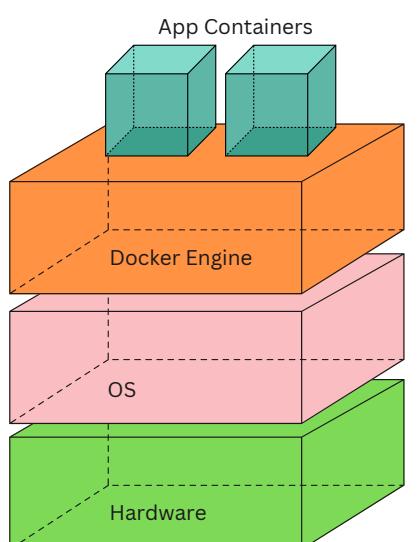
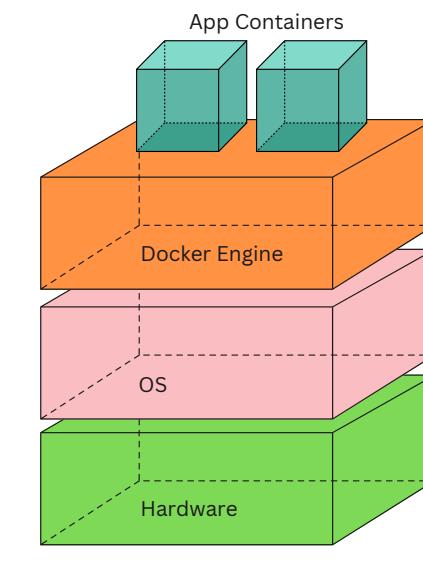
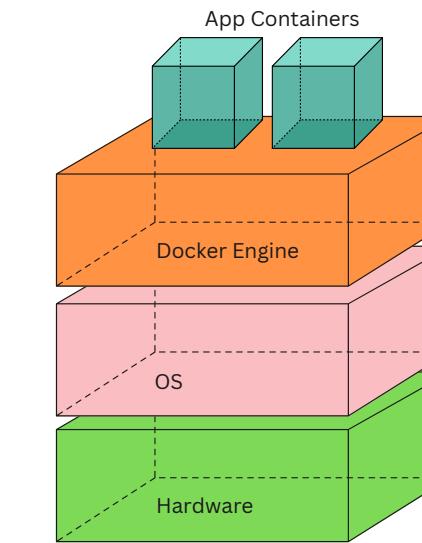
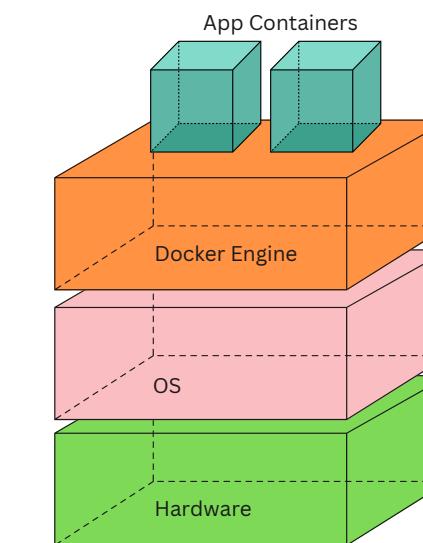
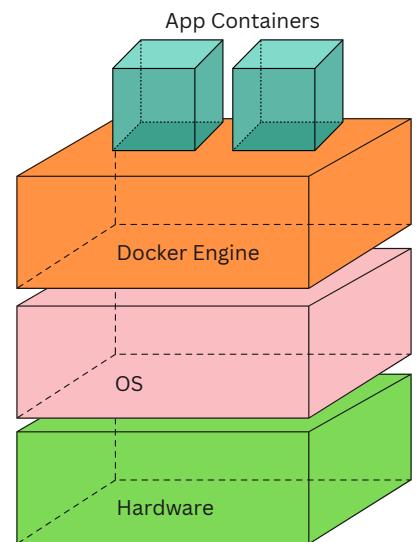
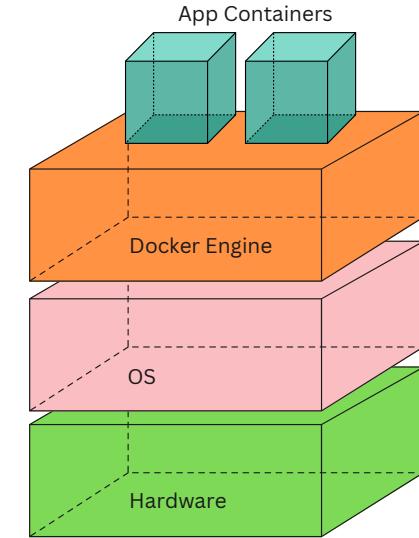
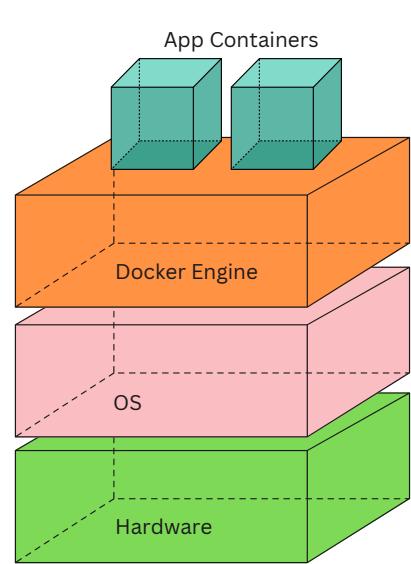
# Docker Challenges →



# Docker Challenges →



# Docker Challenges →



# Docker Challenges

As the number of containers and machines increases, managing Docker environments becomes increasingly complex.

Manual processes for scaling, updating, and configuration lead to inefficiencies and potential inconsistencies, highlighting the need for better orchestration solutions.

---

## Managing Multiple Containers

Scaling and resource allocation become complex as the number of containers grows.

---

## Manual Scaling

Adding new machines and containers manually is time-consuming and inefficient.

---

## Update and Deployment Issues

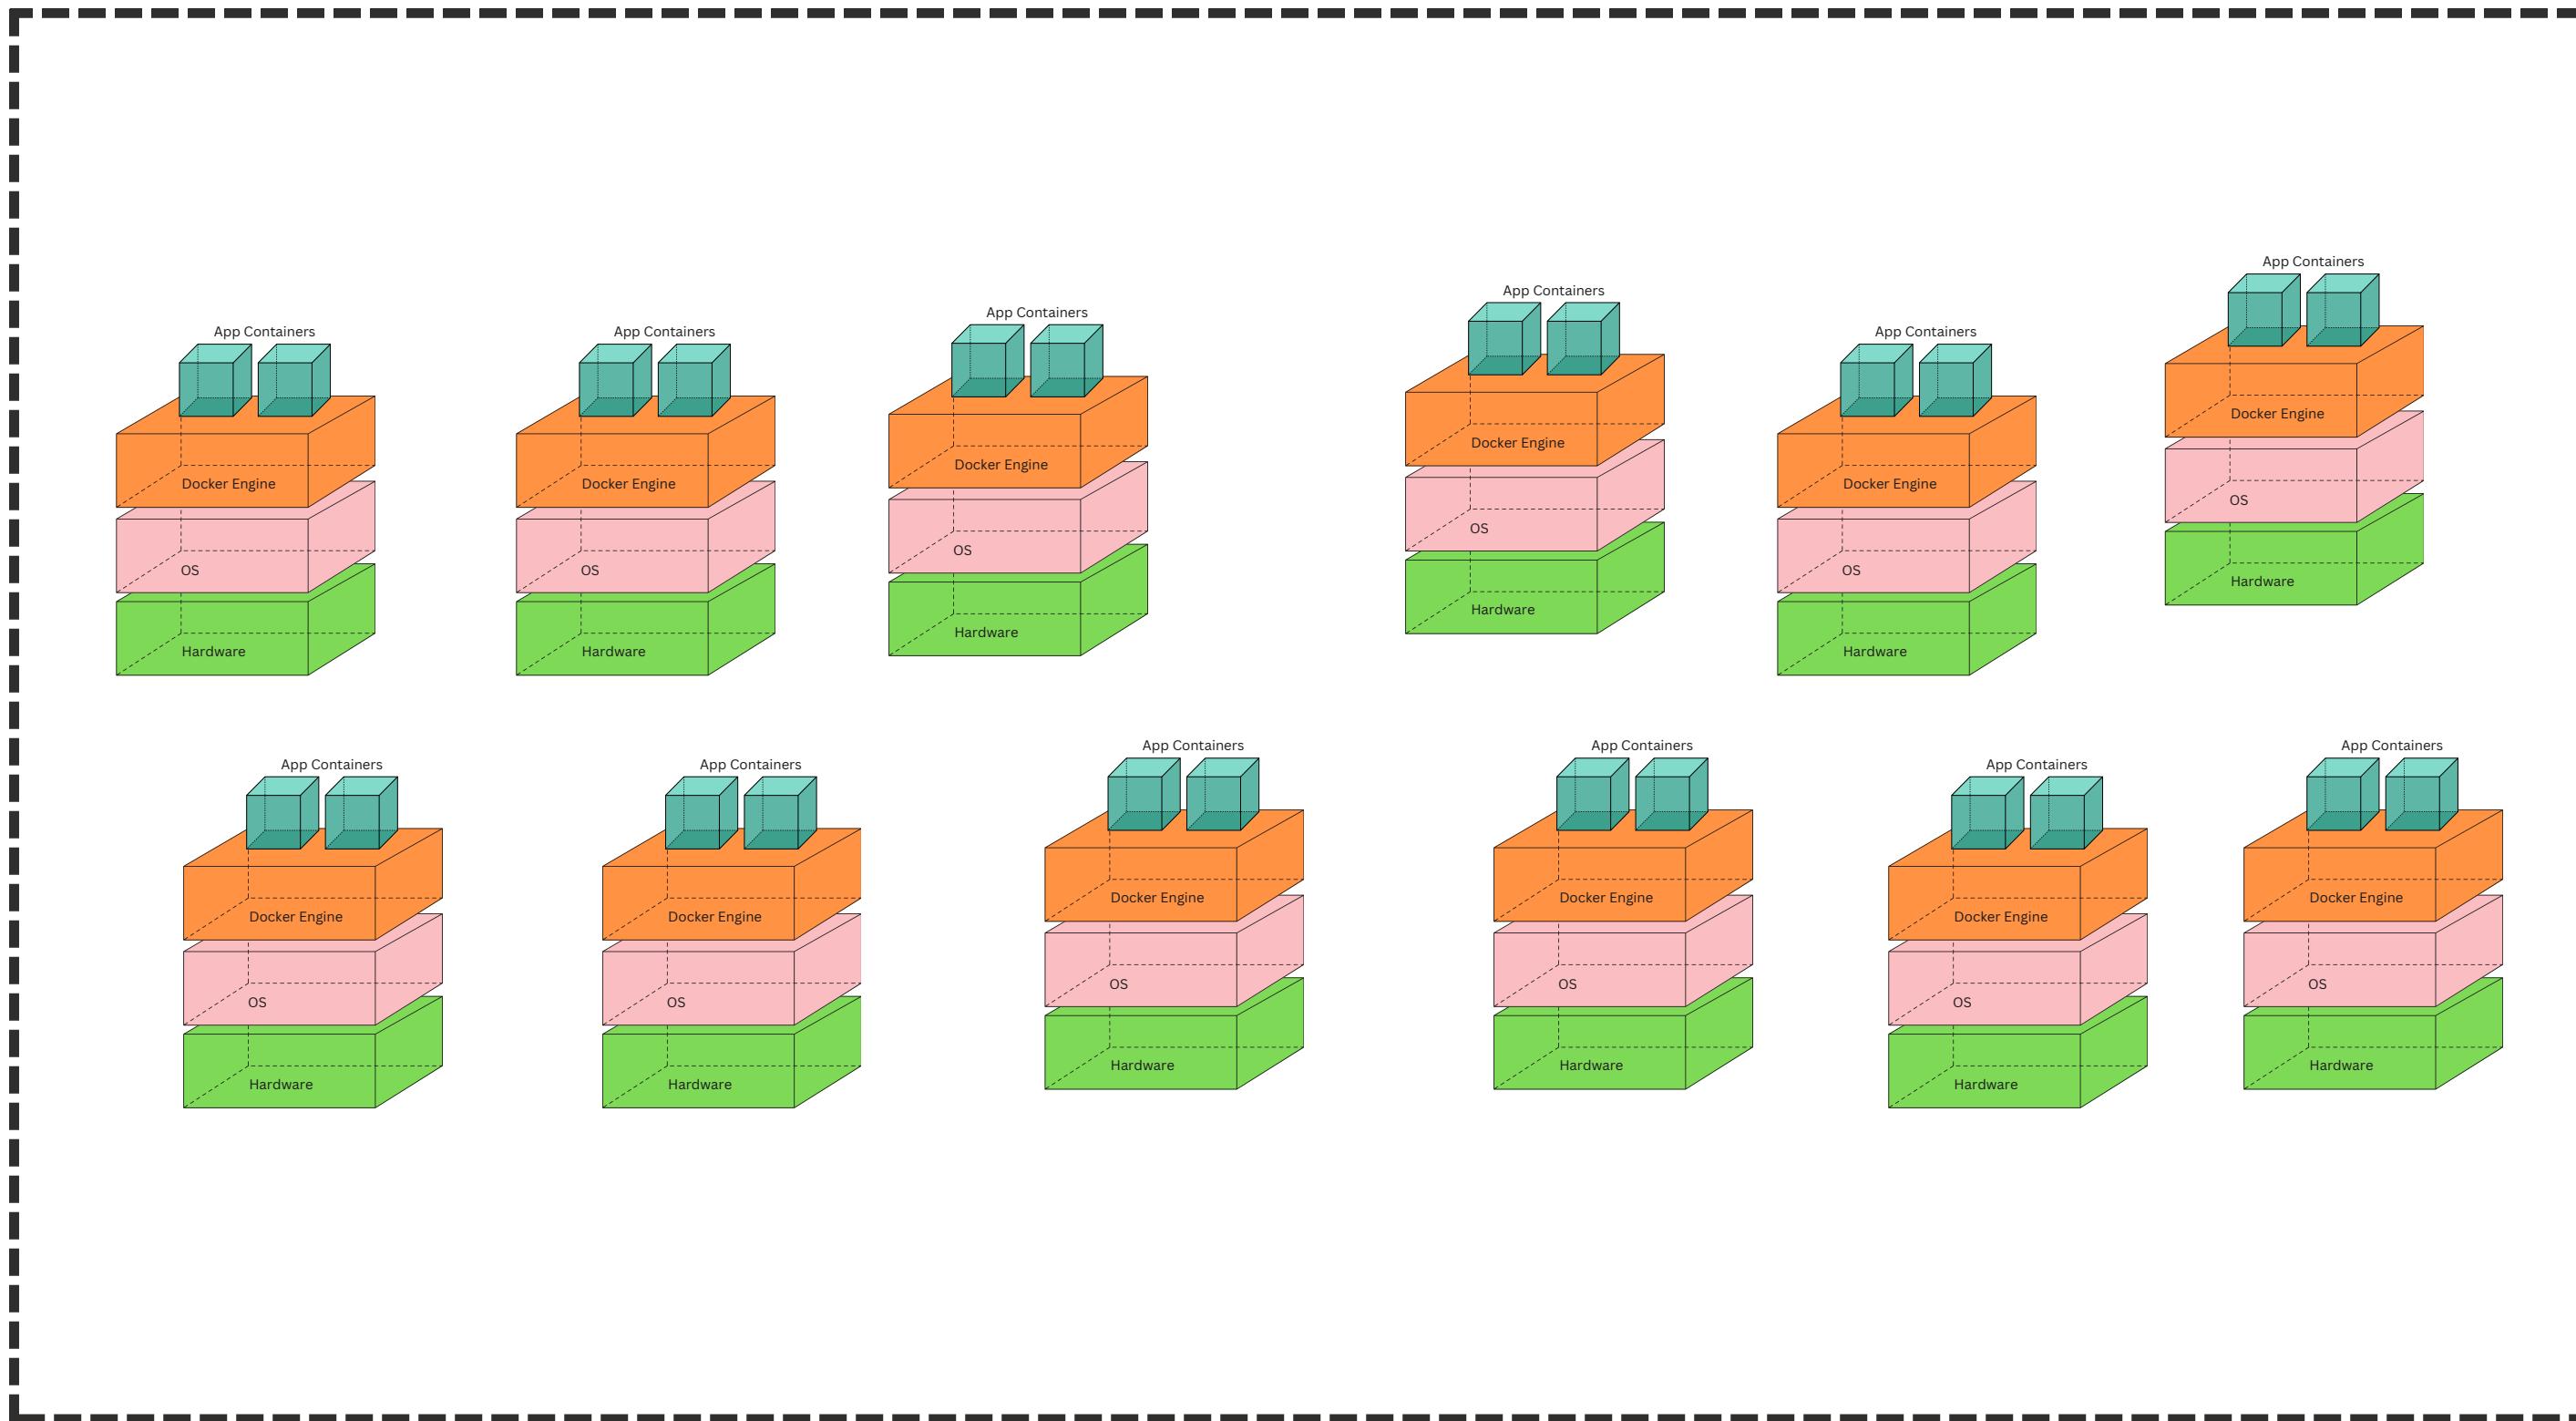
Updates require manual changes on each machine, leading to potential inconsistencies.

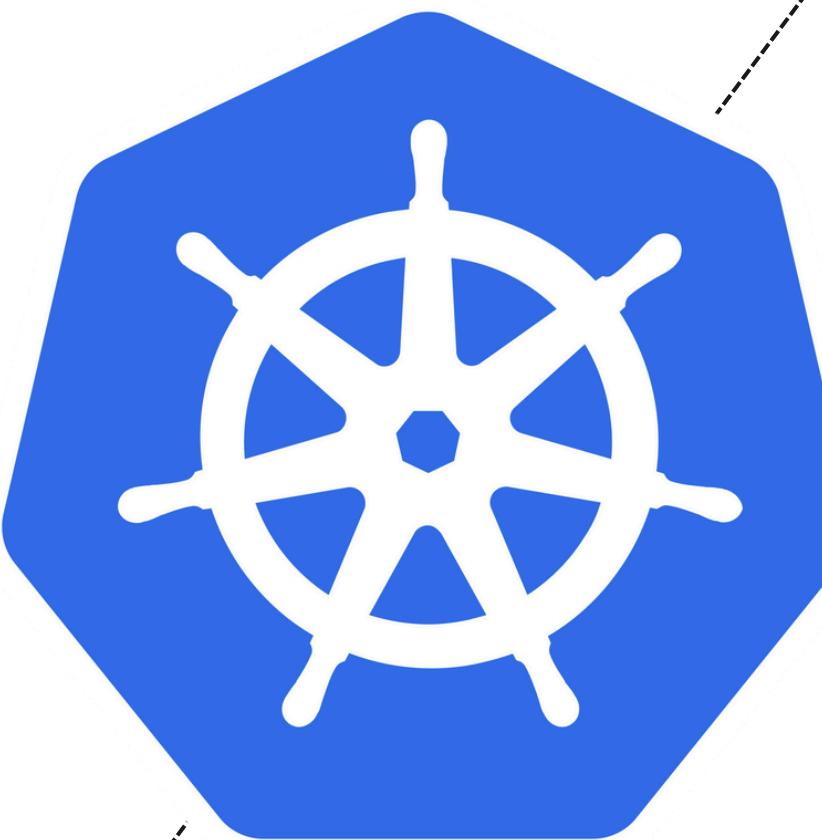
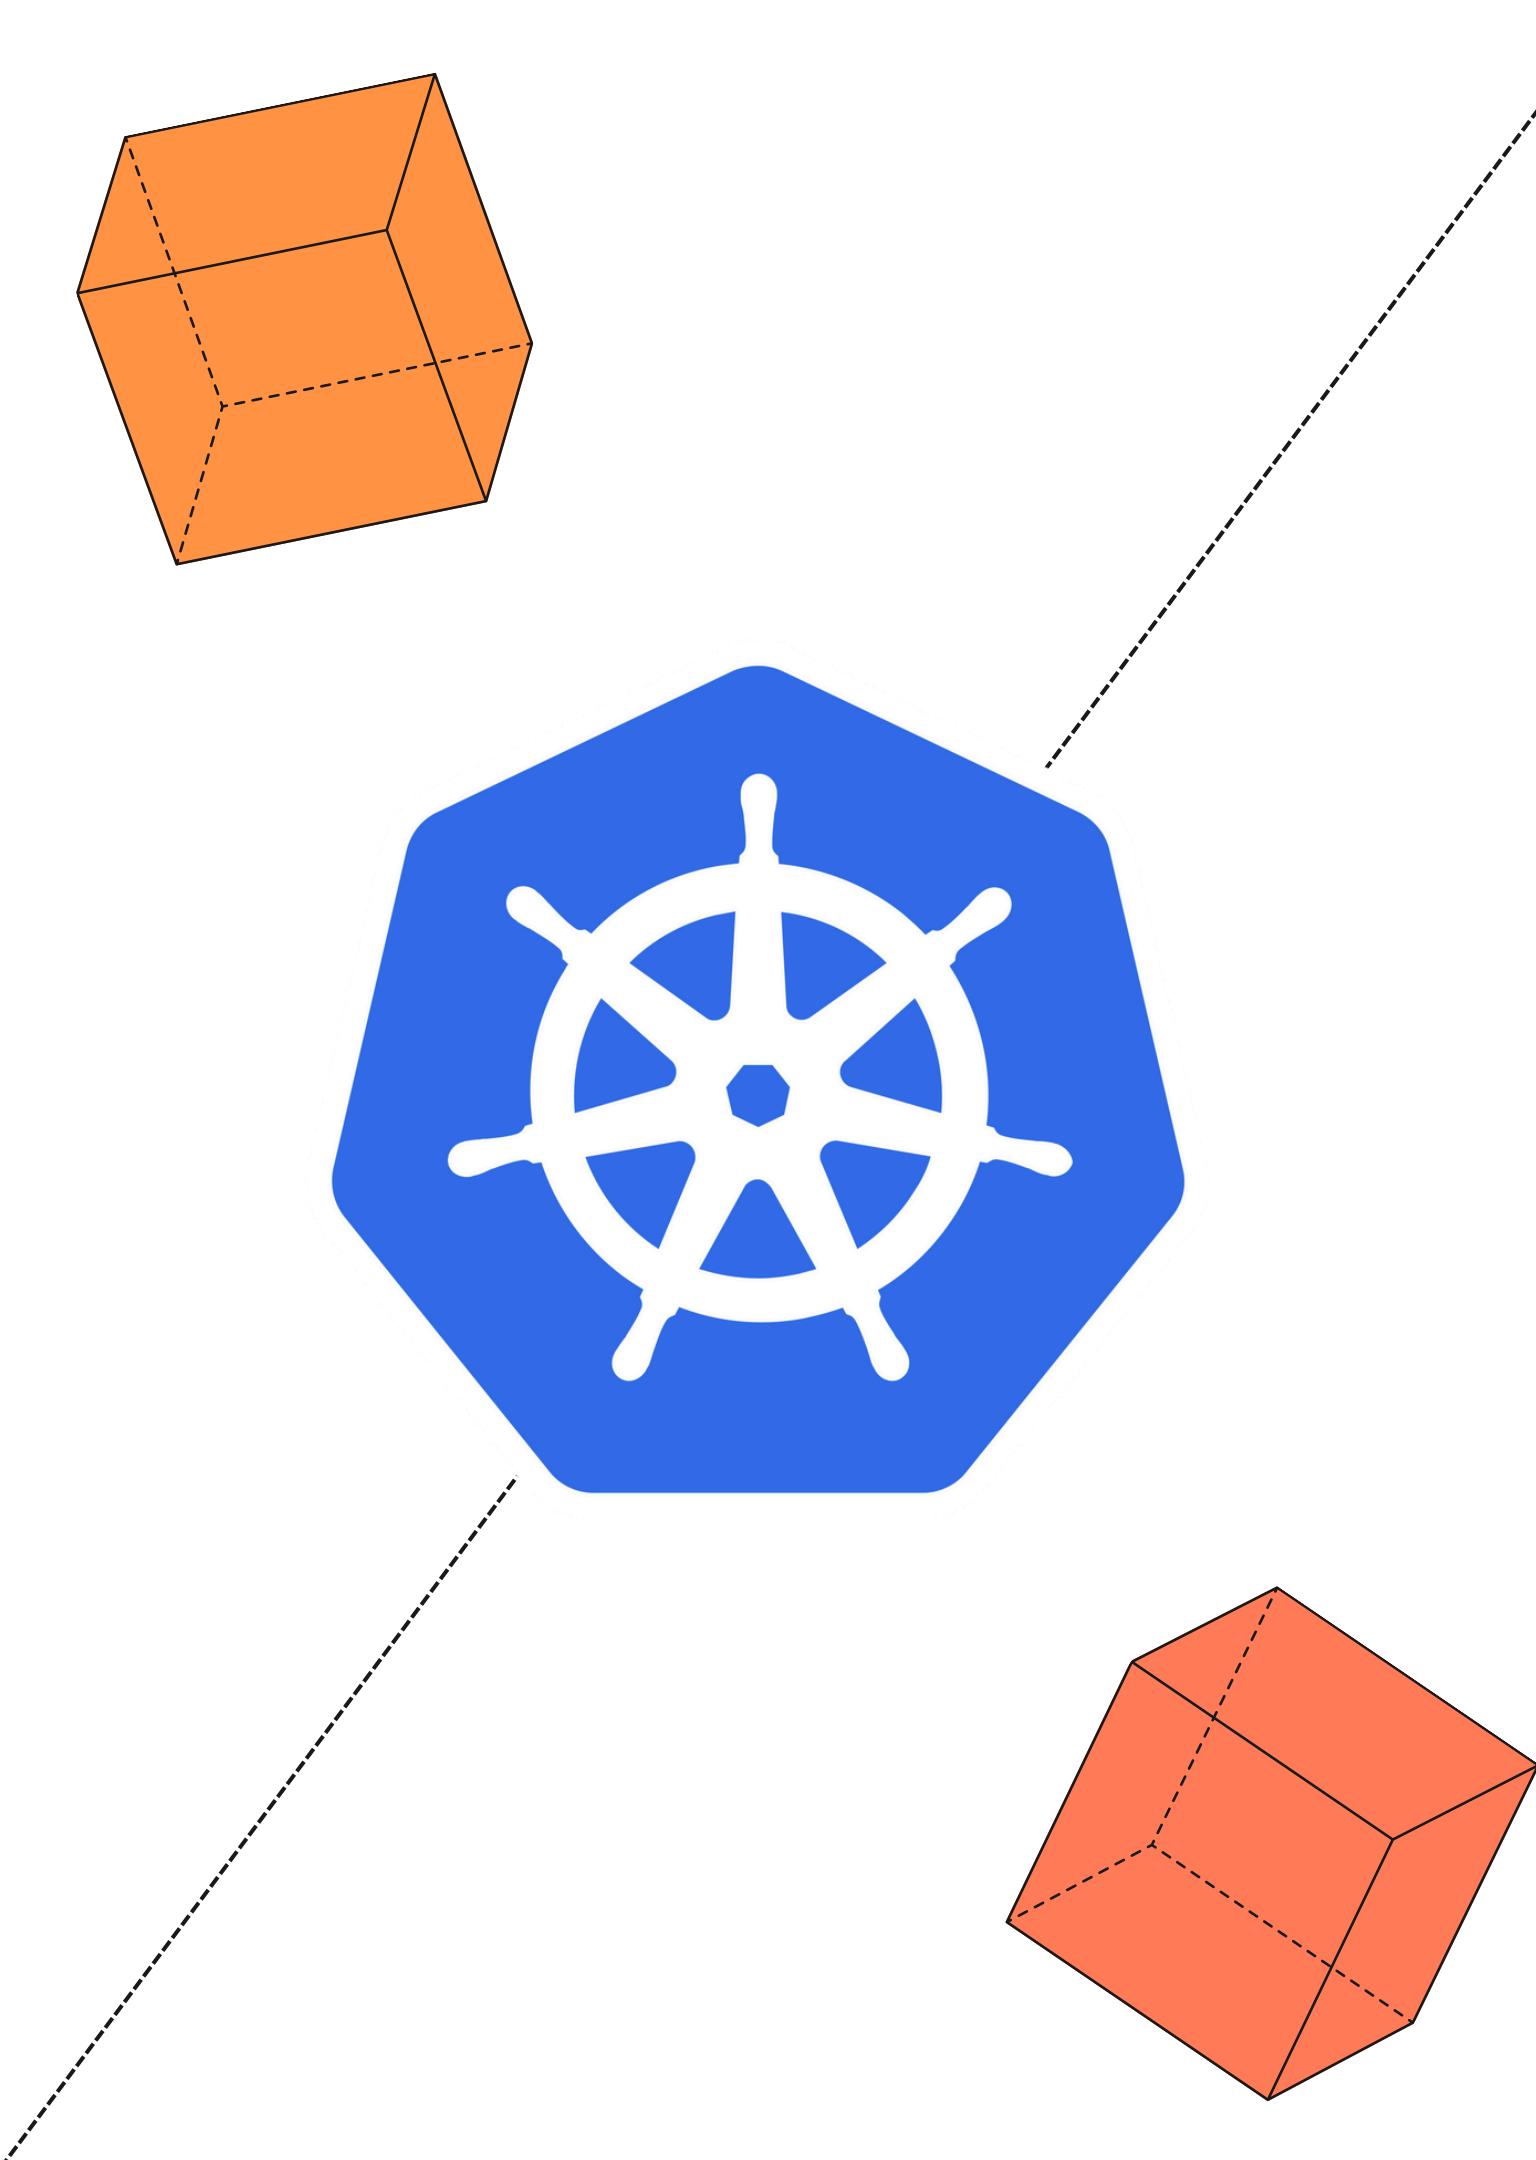
---

## Configuration Management

Ensuring consistent configurations across all resources is error-prone and difficult.

# Orchestrating the chaos →





# Kubernetes

Container Orchestration



# What is Kubernetes?

---

## Automated Deployments

Kubernetes automates the process of deploying containerized applications across a cluster of machines.

Developers can define the desired state of their application, and Kubernetes takes care of the deployment process.

## Self-Healing Capabilities

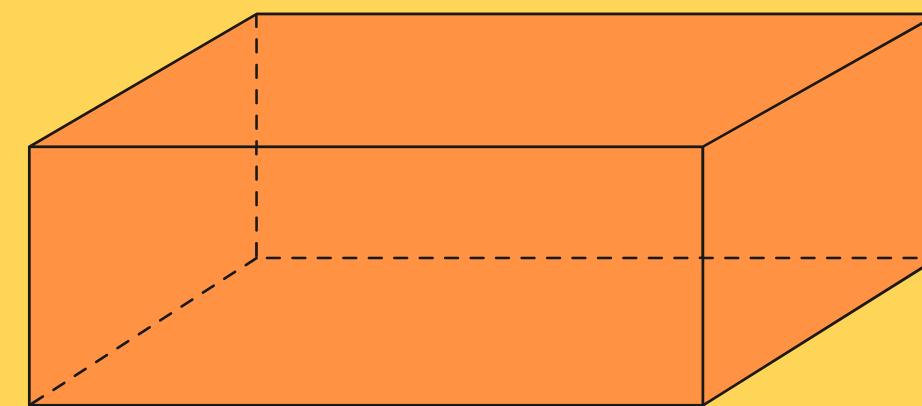
Kubernetes can automatically restart failed containers and reschedule them on healthy nodes (machines) within the cluster, ensuring high availability of your application.

## Service Discovery

Kubernetes helps applications discover and communicate with each other within the cluster.

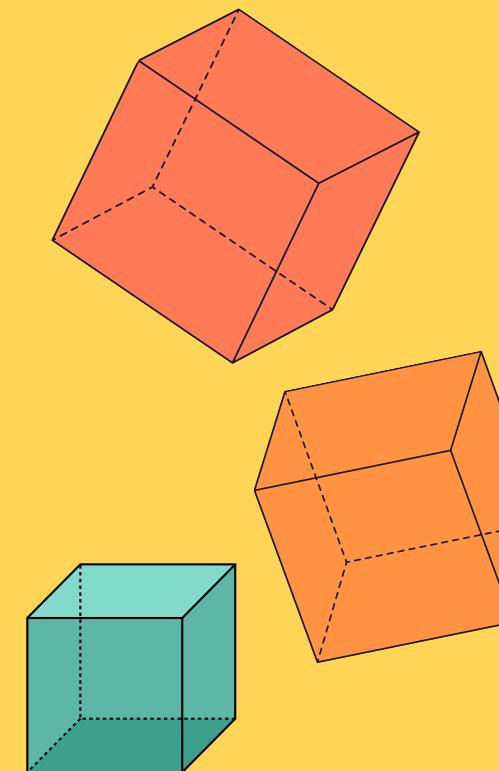
It also provides built-in load balancing to distribute traffic across multiple container instances of a microservice.

# ARCHITECTURE



## MASTER NODE

- Central control unit managing the cluster.
- Controls cluster , Manages resources



## WORKER NODES

- Nodes that run application workloads.
- Runs applications , Manages Pods

# Worker Node

**Kubelet** (Agent ensuring containers are running in a Pod)

Manages containerized applications on the node, ensuring they're running as expected.

**Kube-Proxy** (Handles network proxying and load balancing.)

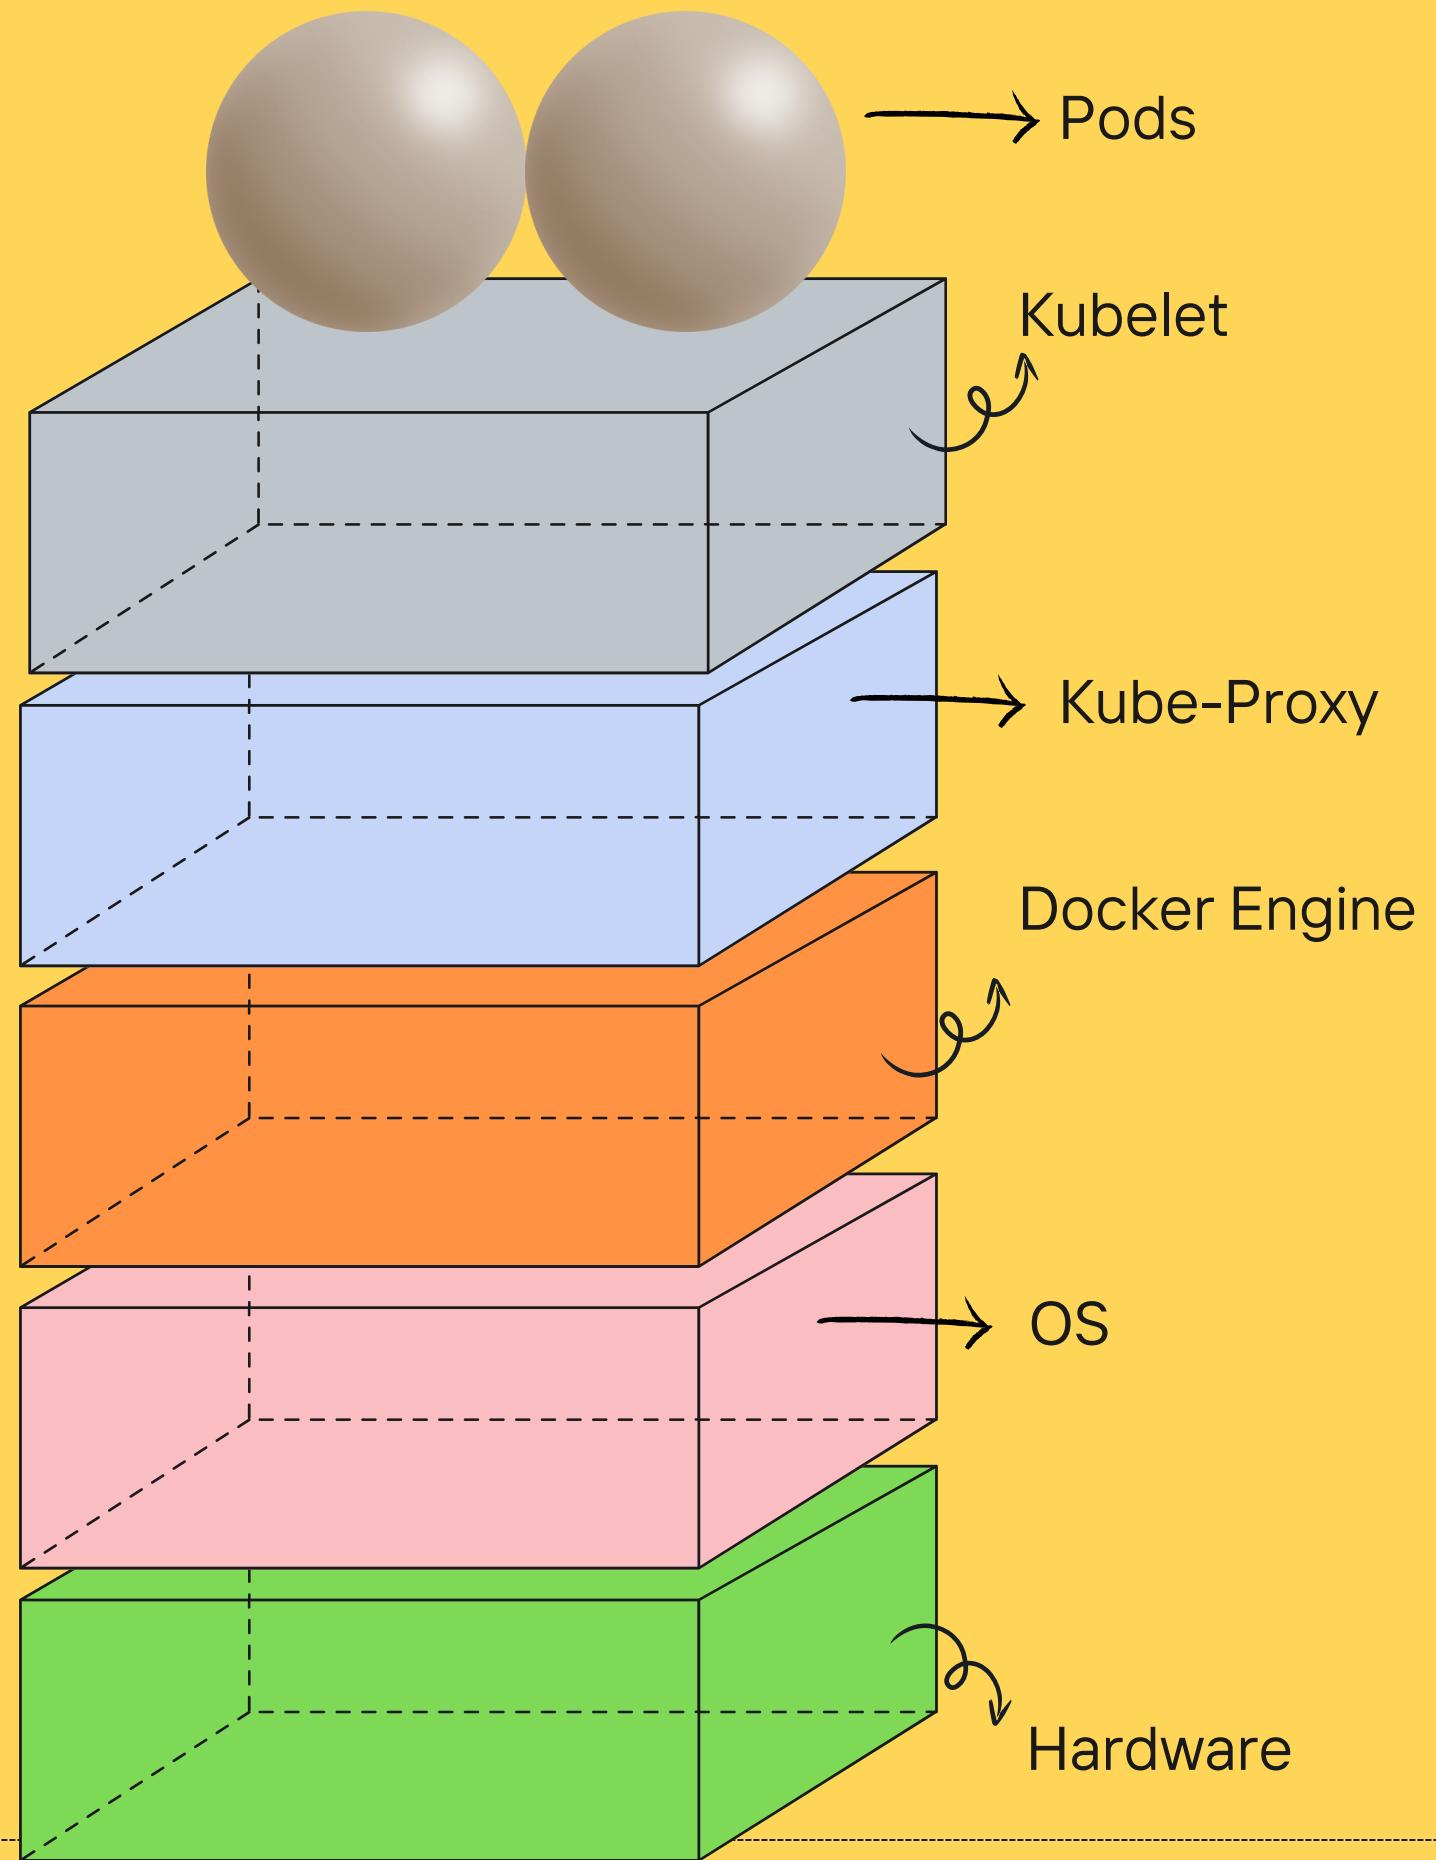
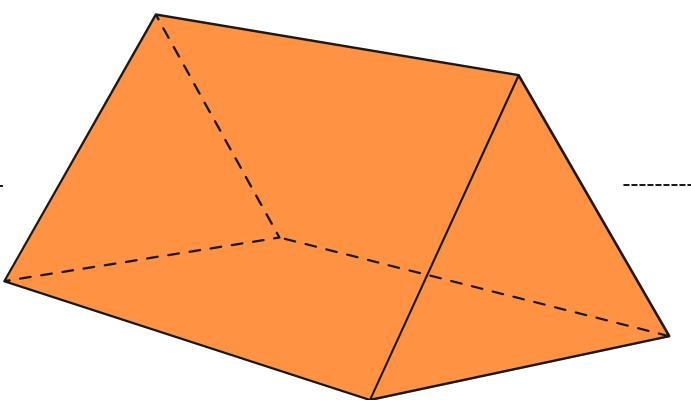
Manages network rules and traffic routing for services within the node.

**Container Runtime** (Software for running containers.)

Executes and manages container life cycles, such as Docker or containerd.

**Pods** (Smallest deployable units in Kubernetes.)

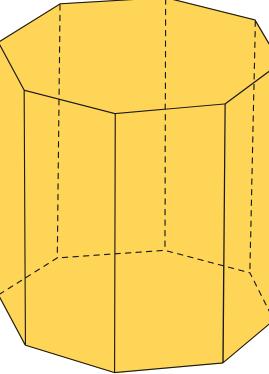
Encapsulates one or more containers, sharing storage and network resources.



# Master Node

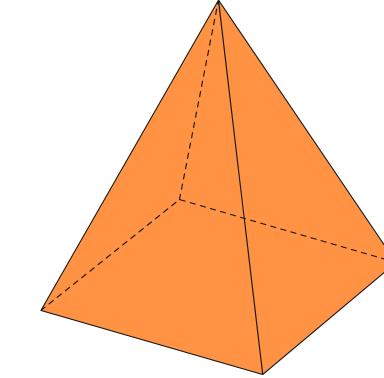
The master node is the central control unit of a Kubernetes cluster.

It manages the cluster's state, schedules workloads, and ensures the desired state matches the actual state through various components.



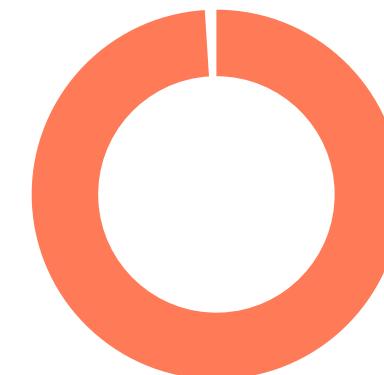
## API Server (Frontend to the Kubernetes control plane)

Updates require manual changes on each machine, leading to potential inconsistencies.



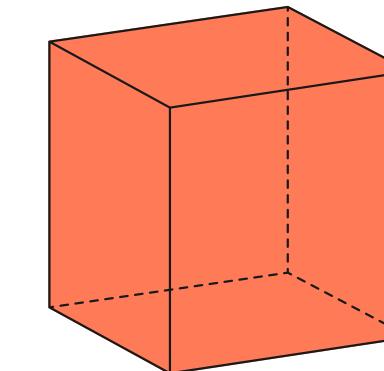
## Scheduler (Assigns workloads to specific nodes)

Determines which nodes are suitable for scheduling new Pods based on resource availability.



## Controller Manager (Ensures desired state of the cluster)

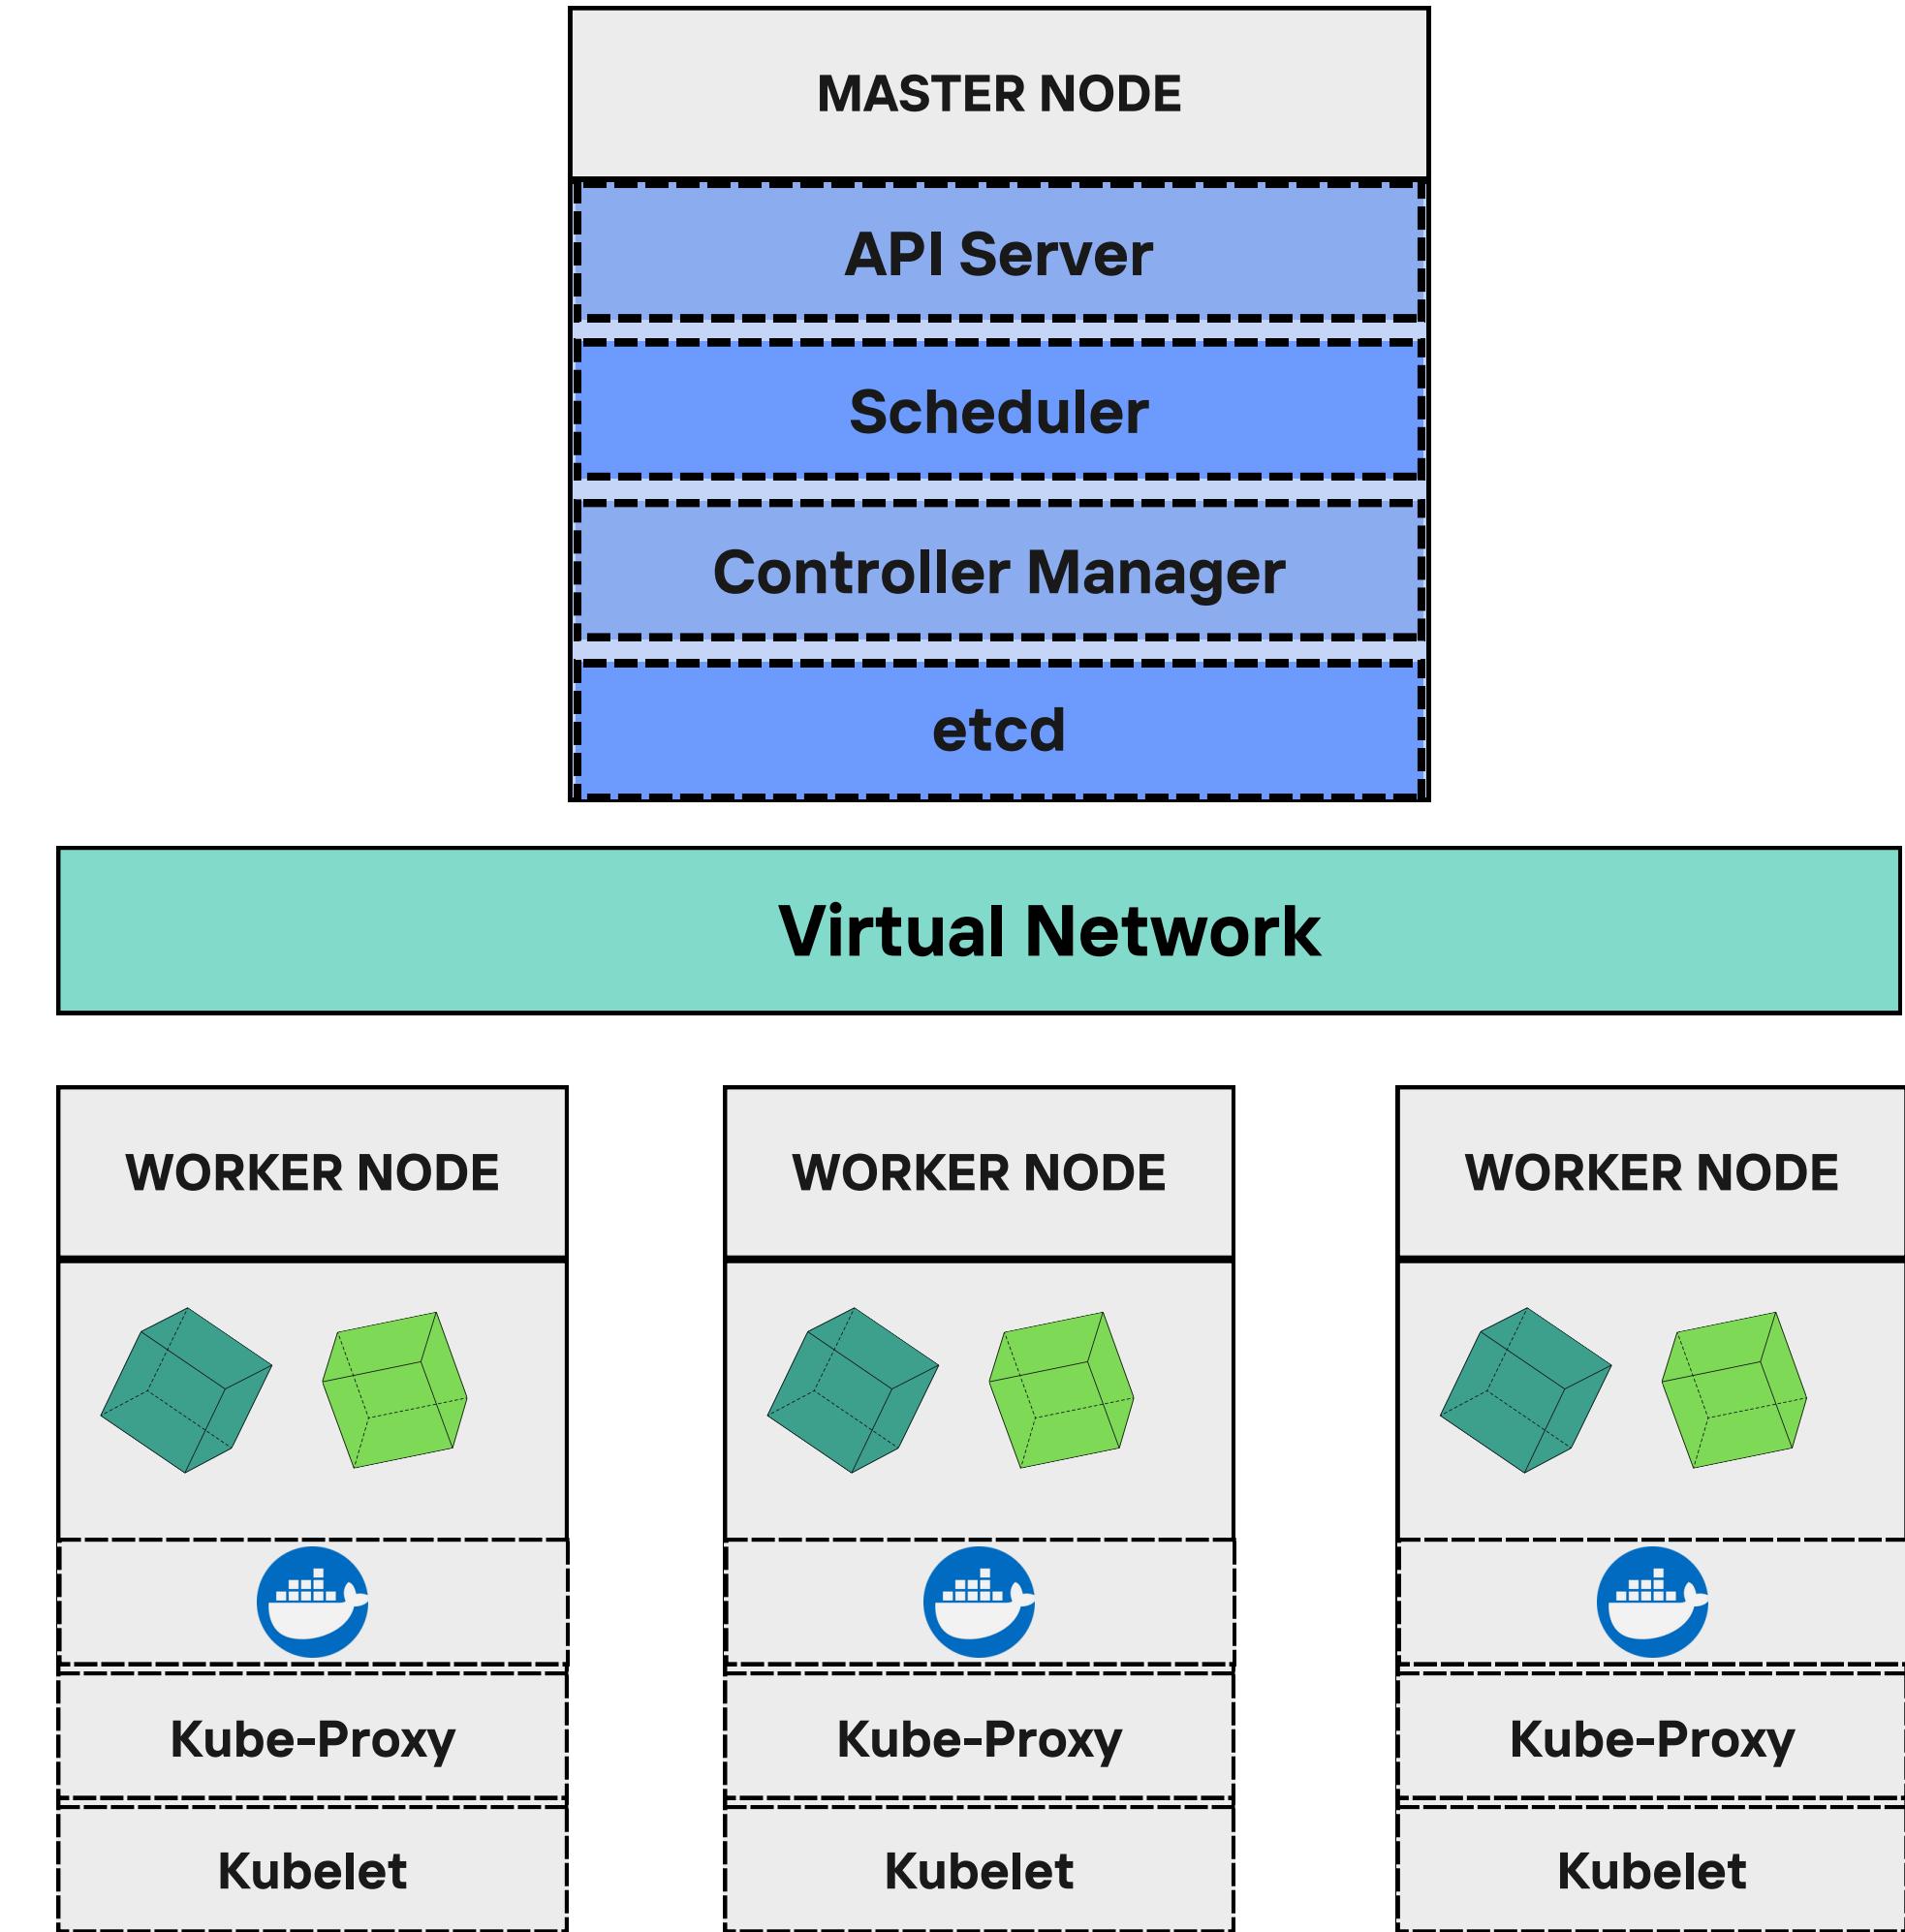
Runs various controllers that handle routine tasks, such as replication and endpoint management.



## etcd: (Key-value store for all cluster data.)

Stores configuration data, state, and metadata for the Kubernetes cluster.

# K8 ARCHITECTURE



# Kubernetes Benefits

Kubernetes offers numerous benefits that make it a powerful and essential tool for managing containerized applications in production environments.

## **Efficient Resource Utilization**

Kubernetes schedules containers based on resource requirements and available capacity, ensuring efficient use of resources.

## **Automated Deployment and Scaling**

Kubernetes automates the deployment and scaling of applications, allowing for continuous delivery and integration.

## **High Availability and Resiliency**

Kubernetes provides self-healing mechanisms, such as auto-restarting failed containers and rescheduling them on healthy nodes.

## **Portability and Flexibility**

Kubernetes is platform-agnostic and can run on various environments, including on-premises, cloud, and hybrid setups.

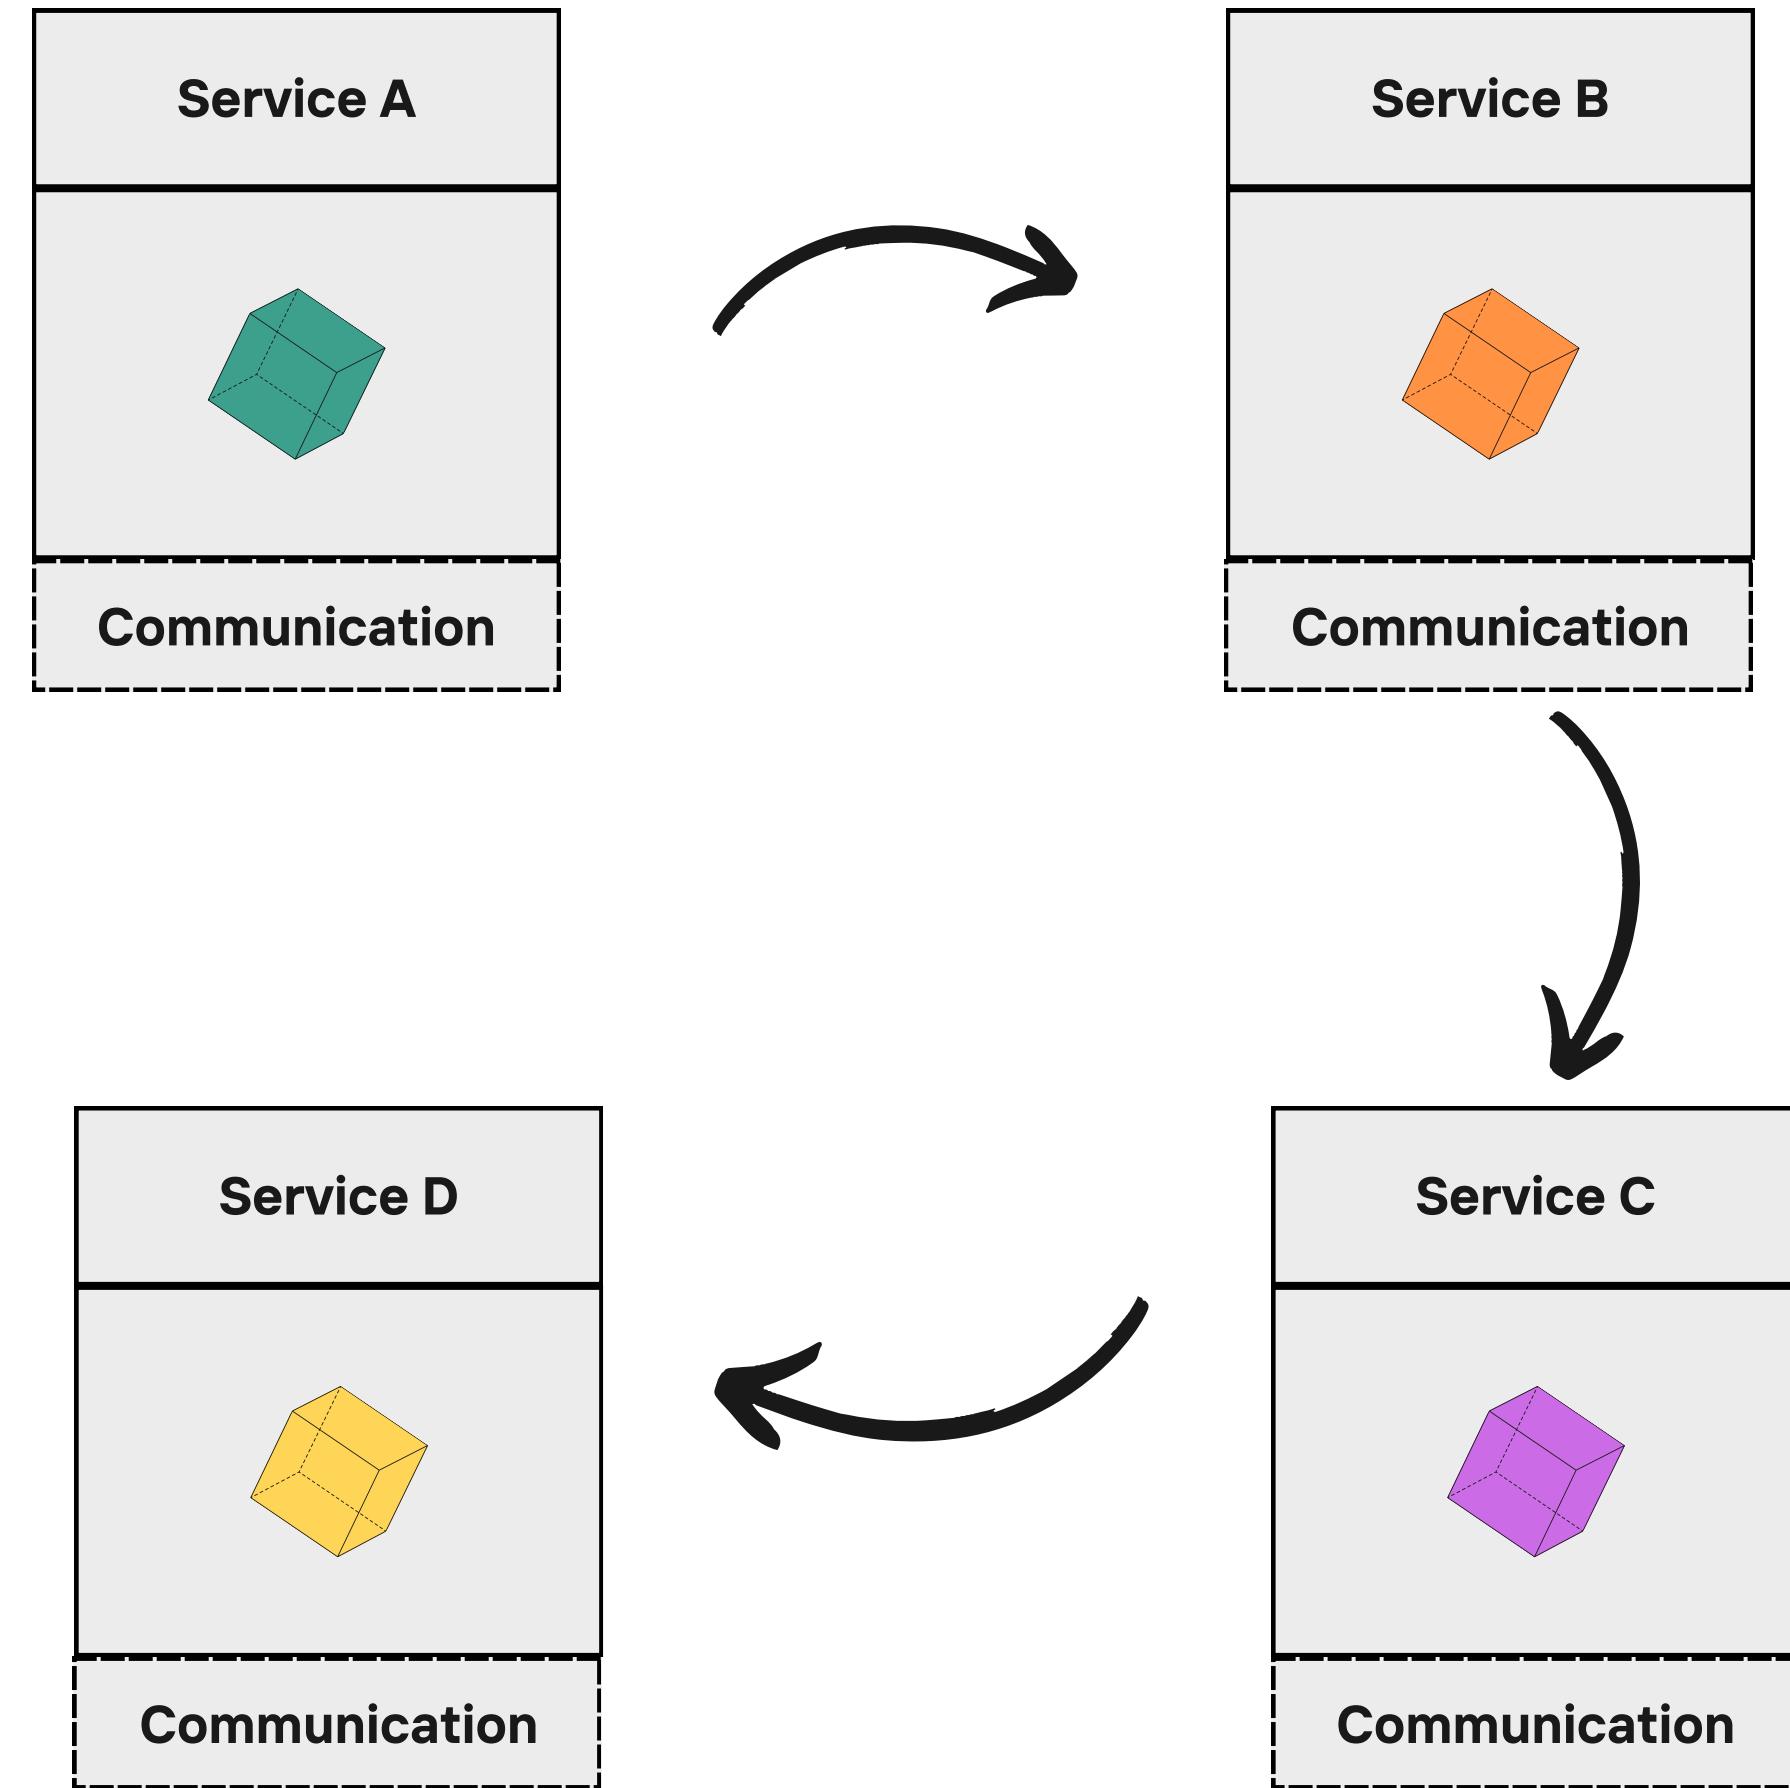


# Limitations of Kubernetes



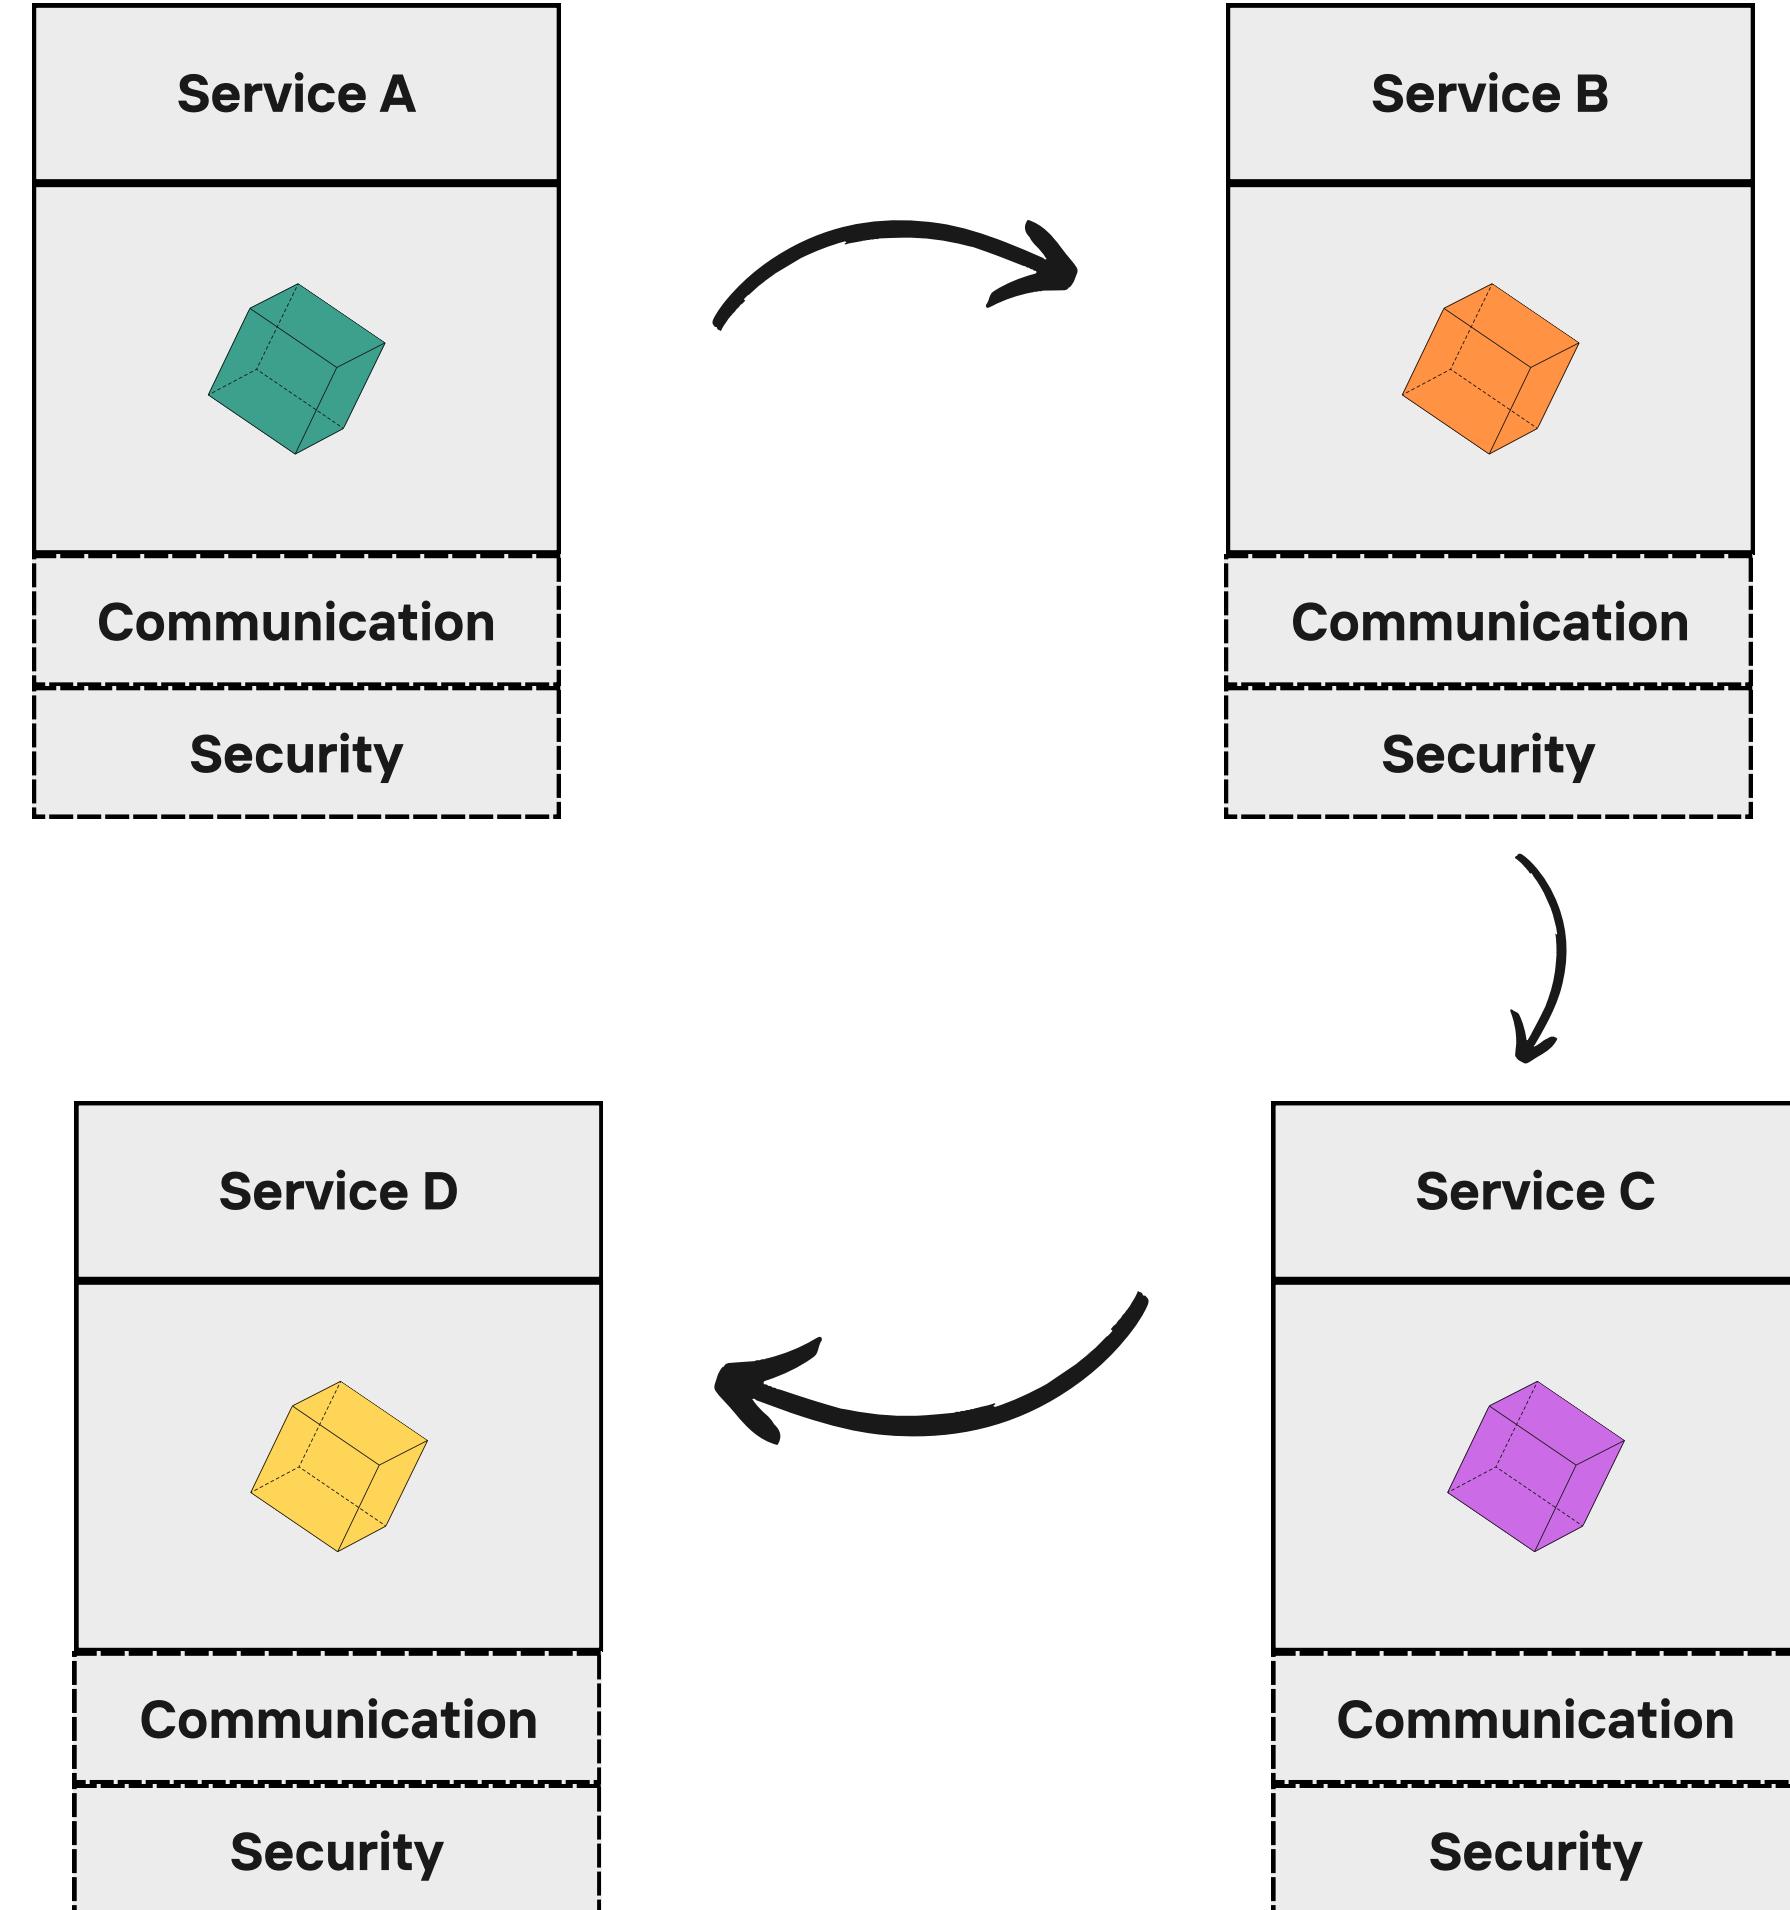
# Communication Logic

- **Service Discovery:** Automatically detecting and connecting services.
- **Retry Logic:** Implementing retry mechanisms for failed requests.
- **Load Balancing:** Distributing network traffic across multiple services.



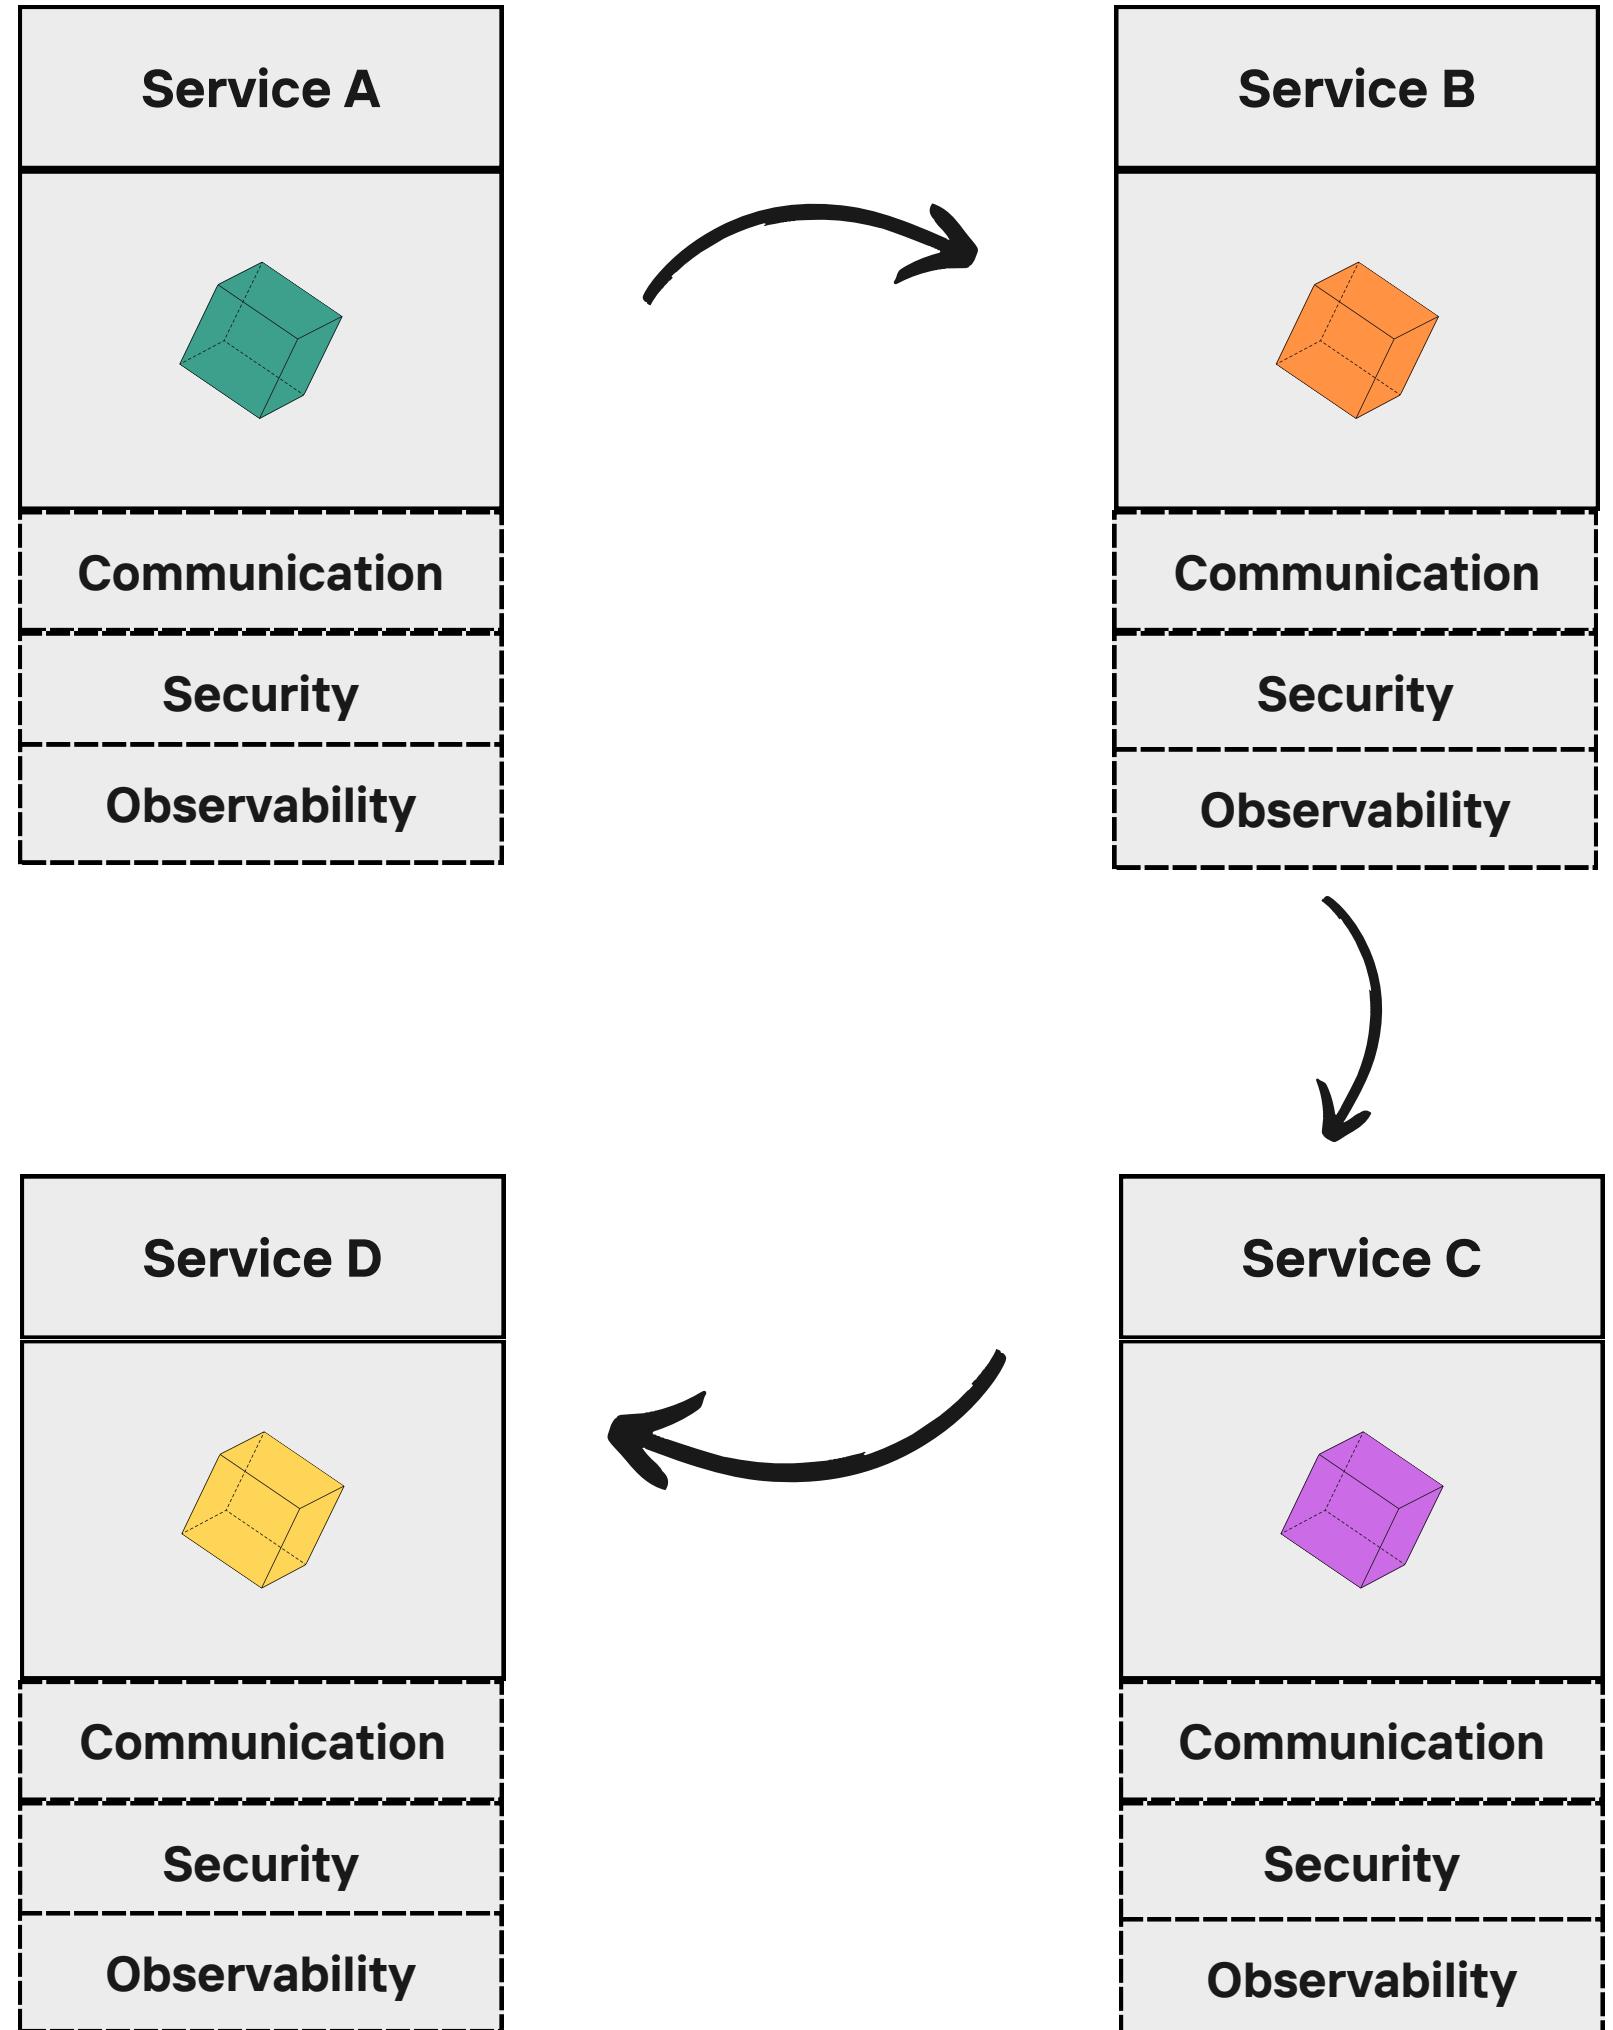
# Security Concerns

- **Authorization:** Ensuring only authorized services can access specific resources.
- **Authentication:** Verifying the identity of services communicating with each other.
- **Encryption:** Securing data in transit between services to prevent interception.



# Limited Observability

- **Health Monitoring:** Tracking the status and health of individual services.
- **Performance Metrics:** Measuring response times, resource usage, and throughput.
- **Communication Patterns:** Visualizing and tracing interactions between microservices.



# DIY Service Management Challenges



Building your own solutions to manage inter-service communication, observability, and security can lead to increased complexity and maintenance overhead.

Each new service must incorporate logic for retries, load balancing, monitoring, and security, which can be time-consuming and error-prone.

---

**Increased Complexity:** Each service requires custom logic for critical functions.

Adding retry mechanisms, load balancing, health monitoring, and security policies to each service increases complexity and potential for errors.

---

**Maintenance Overhead:** Ongoing updates and troubleshooting are more challenging.

Keeping custom implementations up to date with best practices and security standards is labor-intensive.

---

**Scalability Issues:** Custom solutions can hinder scalability.

Managing these functions manually can limit the ability to scale services efficiently.

# Introducing Service Mesh

**A Communication Layer for Microservices**



# Introducing Service Mesh

A **service mesh** is a dedicated infrastructure layer that manages service-to-service communication within a microservices architecture, providing

---

## Service Discovery and Routing

The service mesh helps microservices discover each other and efficiently route traffic to the appropriate service instances.

---

## Security Enforcement

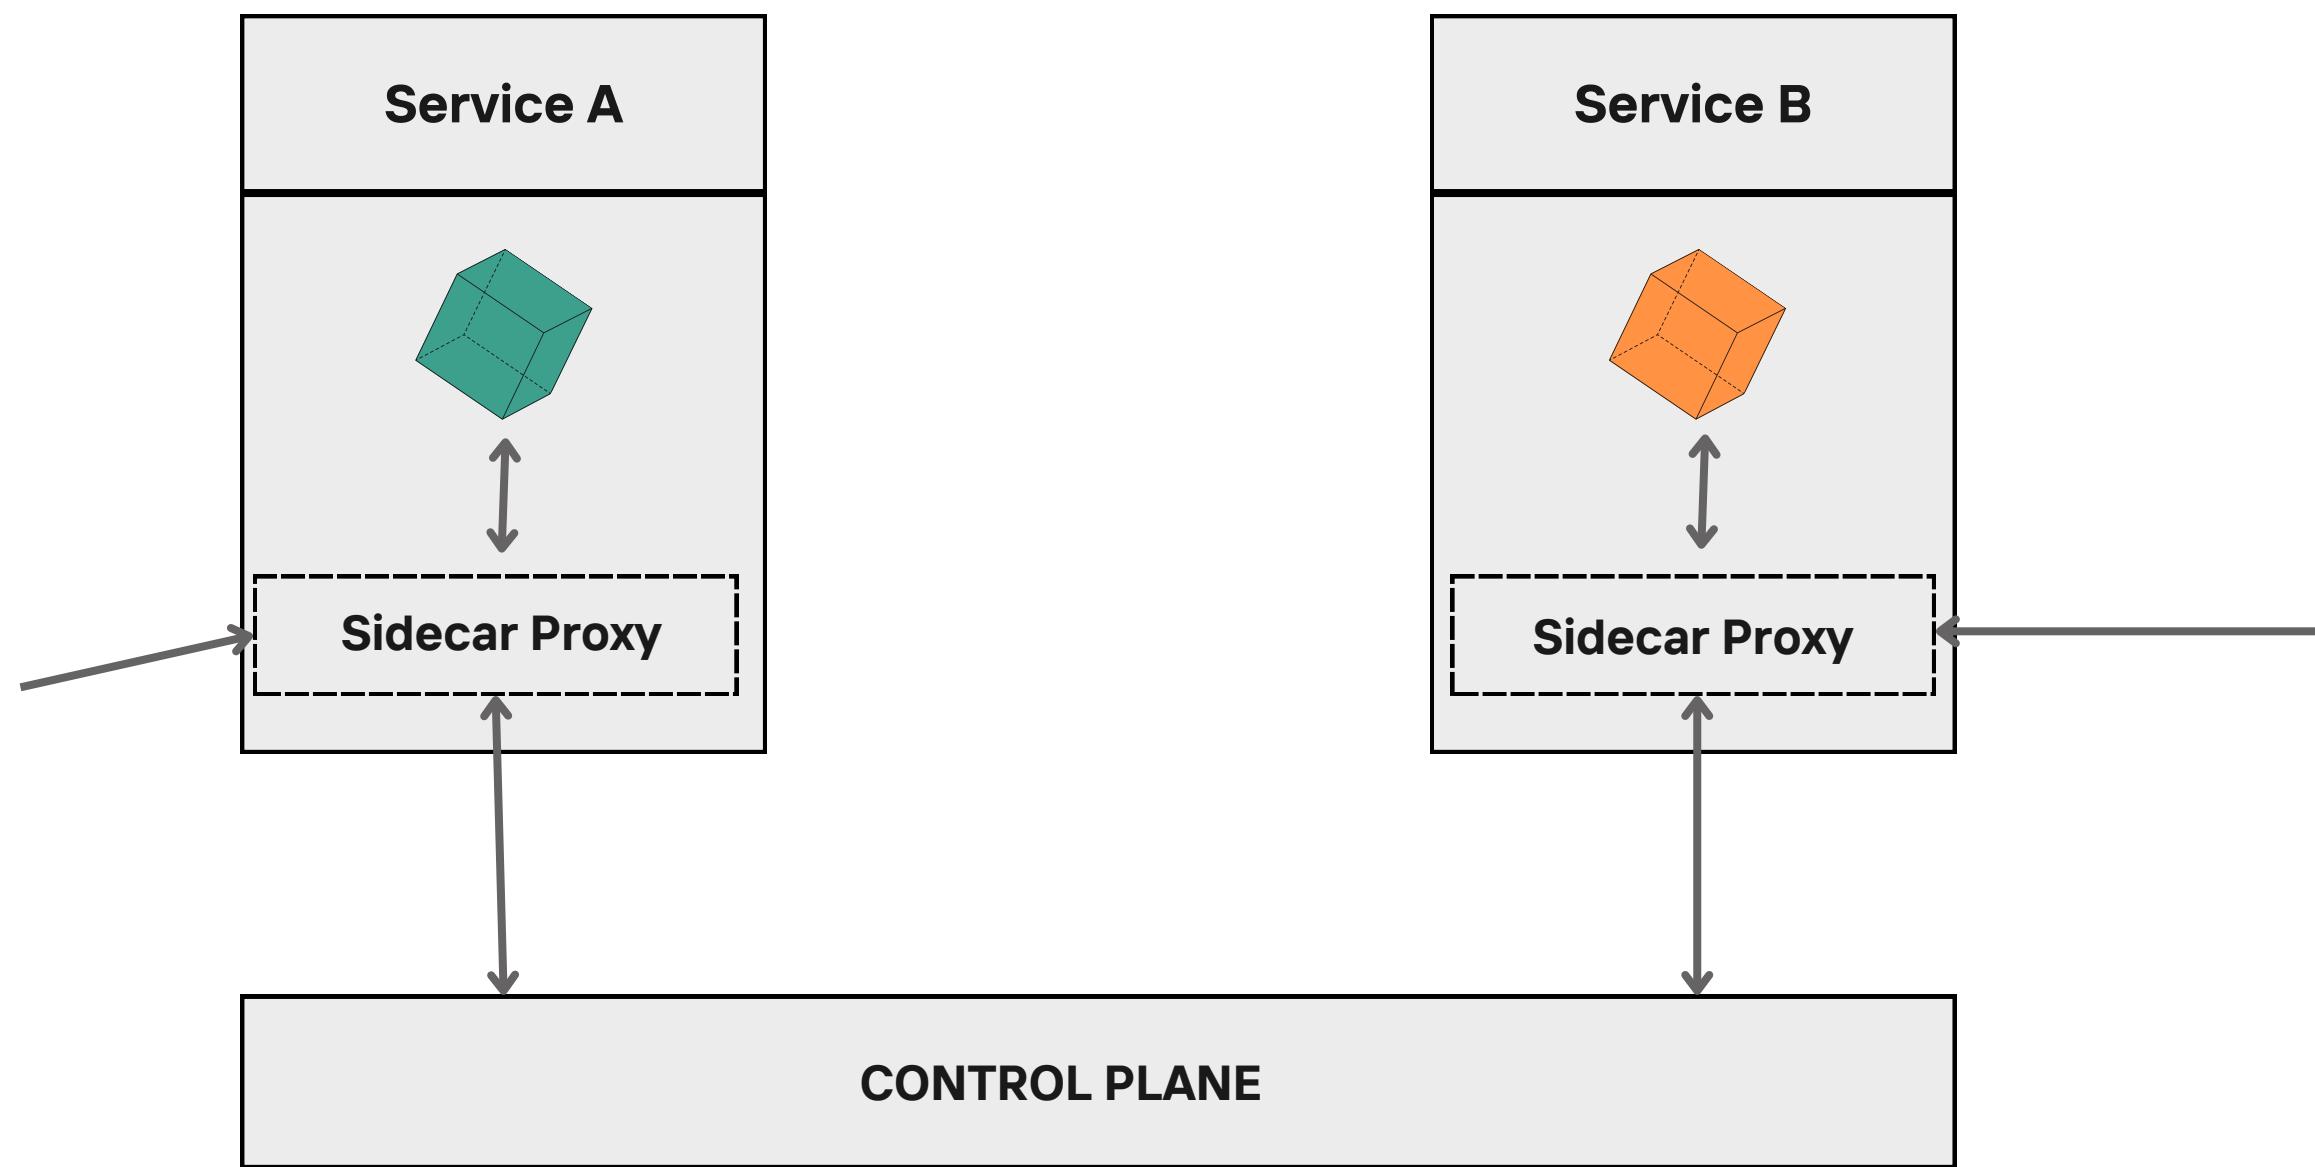
The service mesh can enforce security policies like authentication, authorization, and encryption for communication between microservices.

---

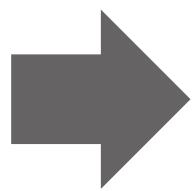
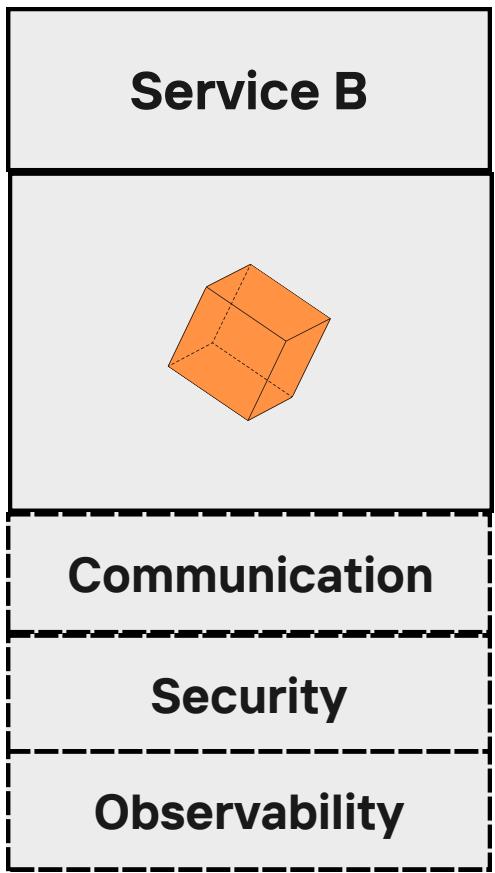
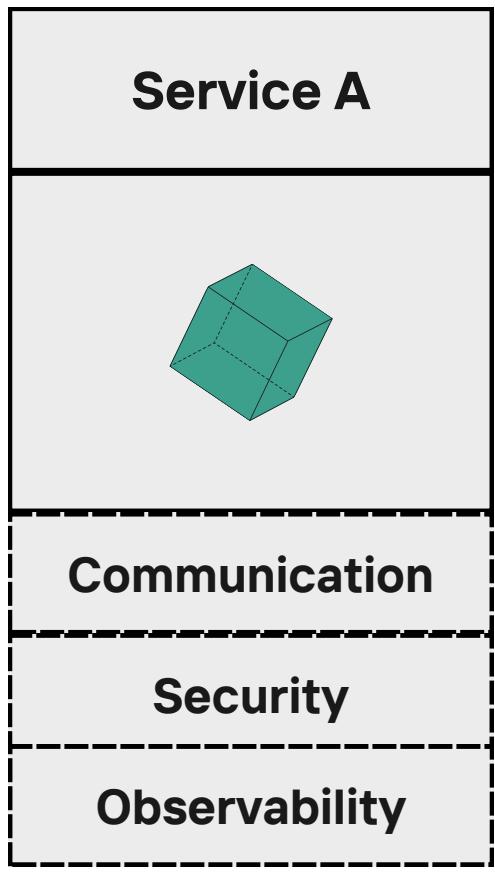
## Enhanced Observability

The service mesh provides valuable insights into service-to-service communication patterns, aiding in troubleshooting and performance optimization.

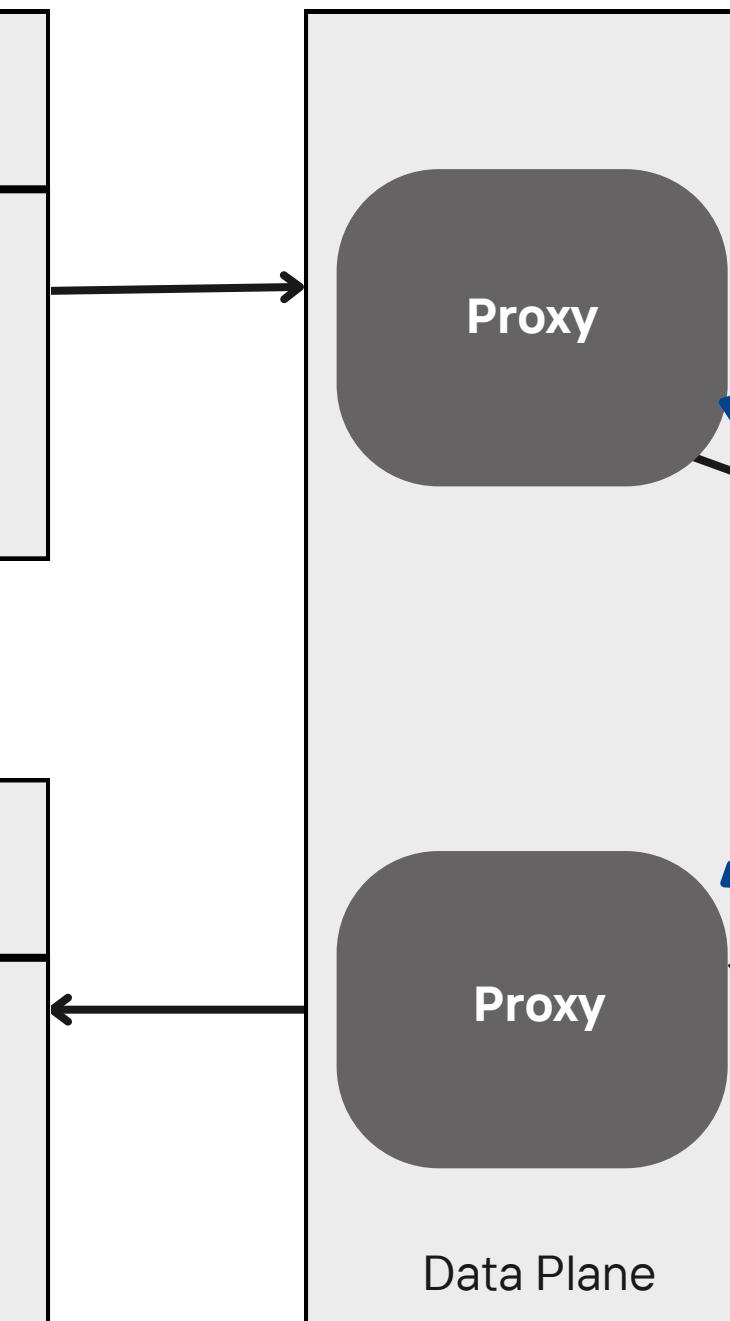
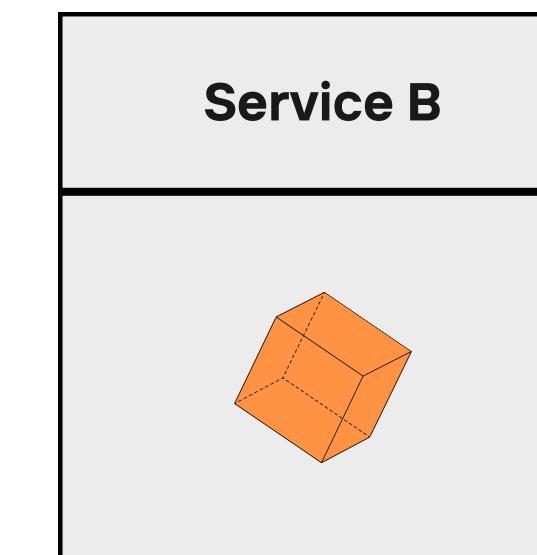
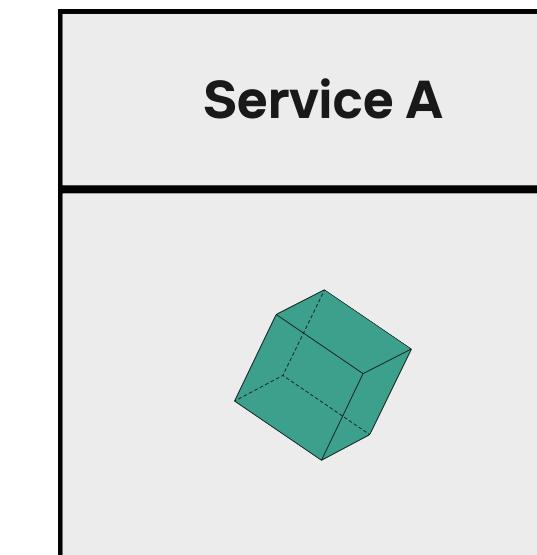
# Service Mesh Architecture



## Without Service Mesh



## With Service Mesh



Control Plane

# Service Mesh Benefits

A service mesh provides a range of core features that enhance the management, observability, and security of microservices within a Kubernetes environment.

---

## **Service Discovery and Load Balancing**

The service mesh automatically discovers and registers microservices within the cluster.

---

## **Manual Scaling**

Adding new machines and containers manually is time-consuming and inefficient.

---

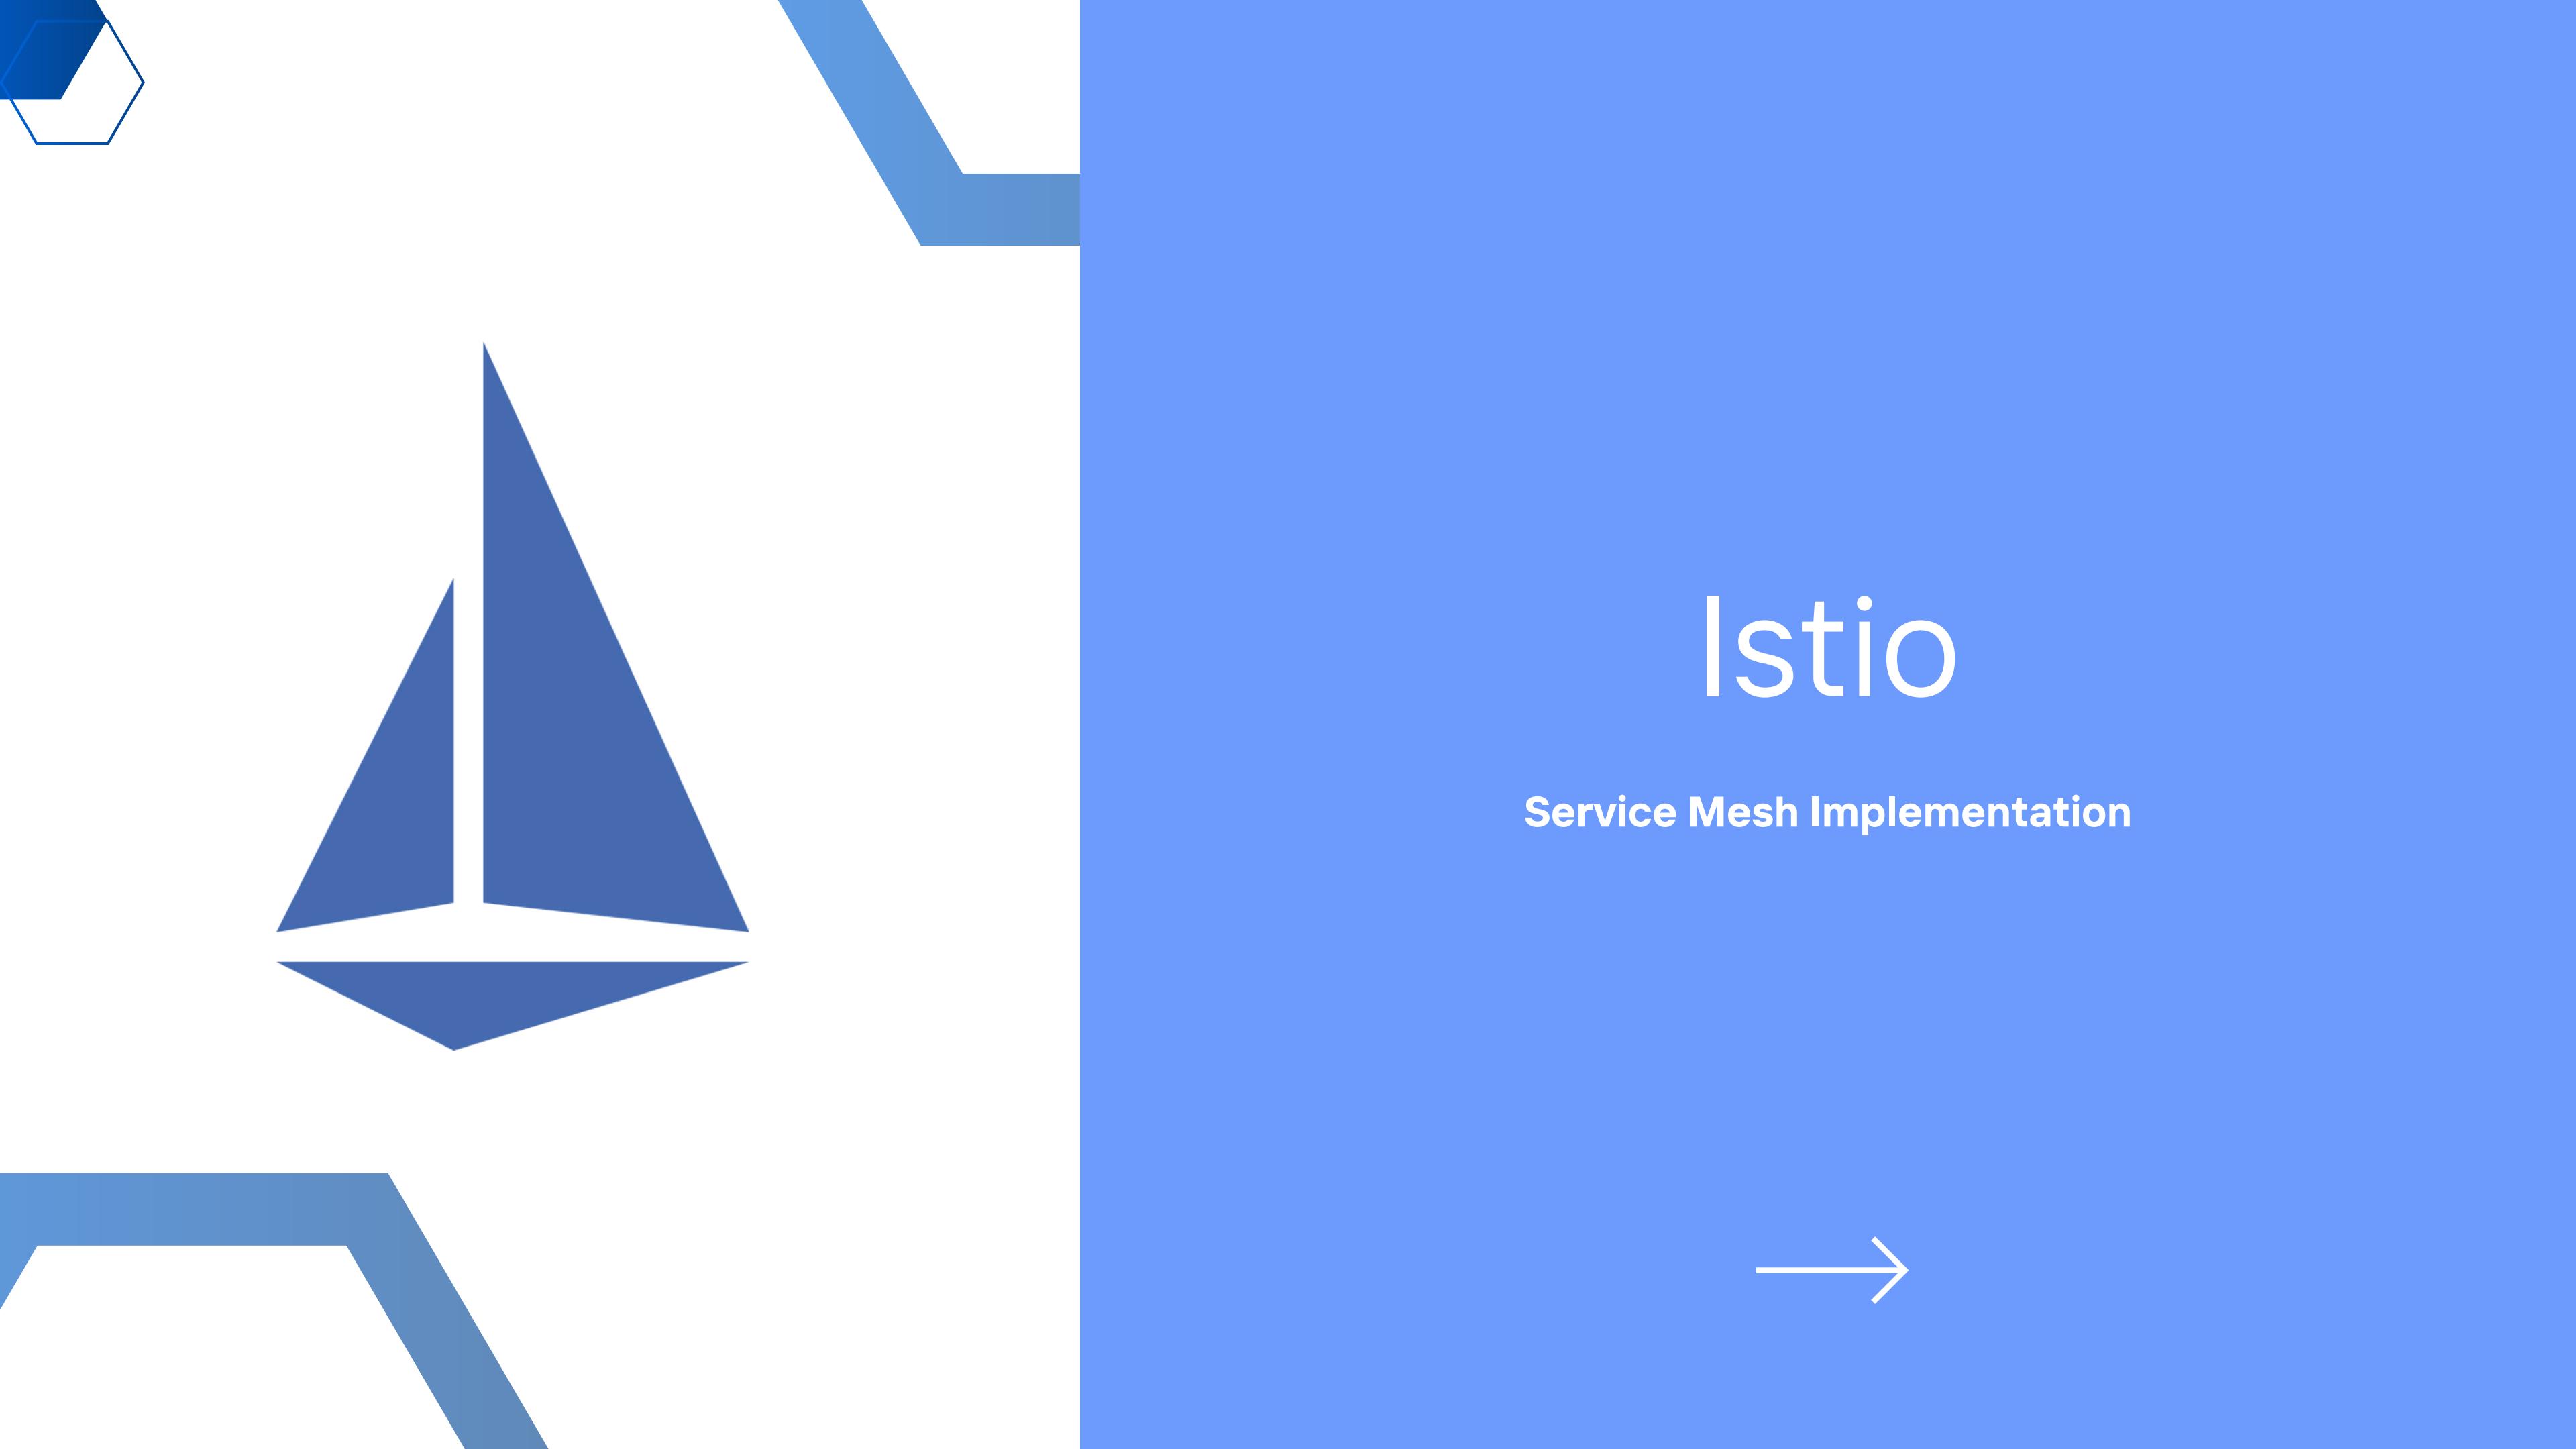
## **Update and Deployment Issues**

Updates require manual changes on each machine, leading to potential inconsistencies.

---

## **Configuration Management**

Ensuring consistent configurations across all resources is error-prone and difficult.

The background features abstract blue geometric shapes: a hexagon at the top left, a large triangle in the center-left, and a trapezoid at the bottom left. A thick blue arrow points from the right towards the central text.

# Istio

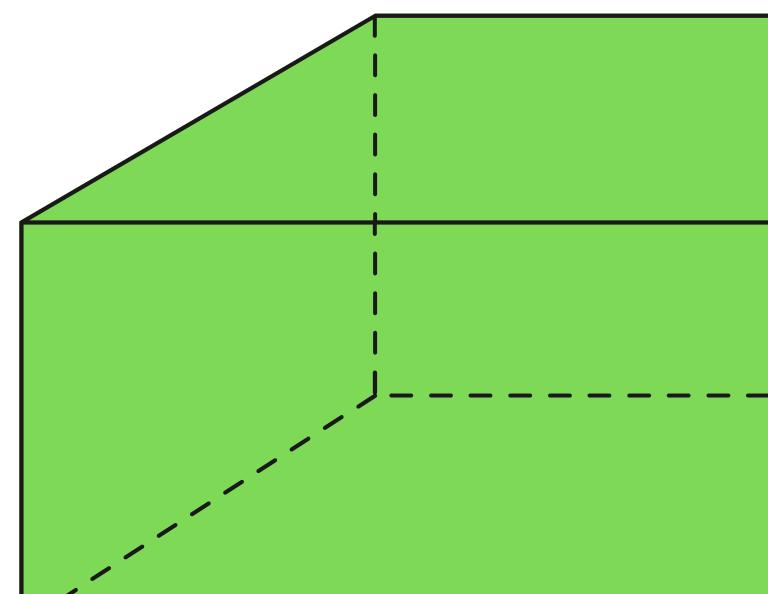
**Service Mesh Implementation**

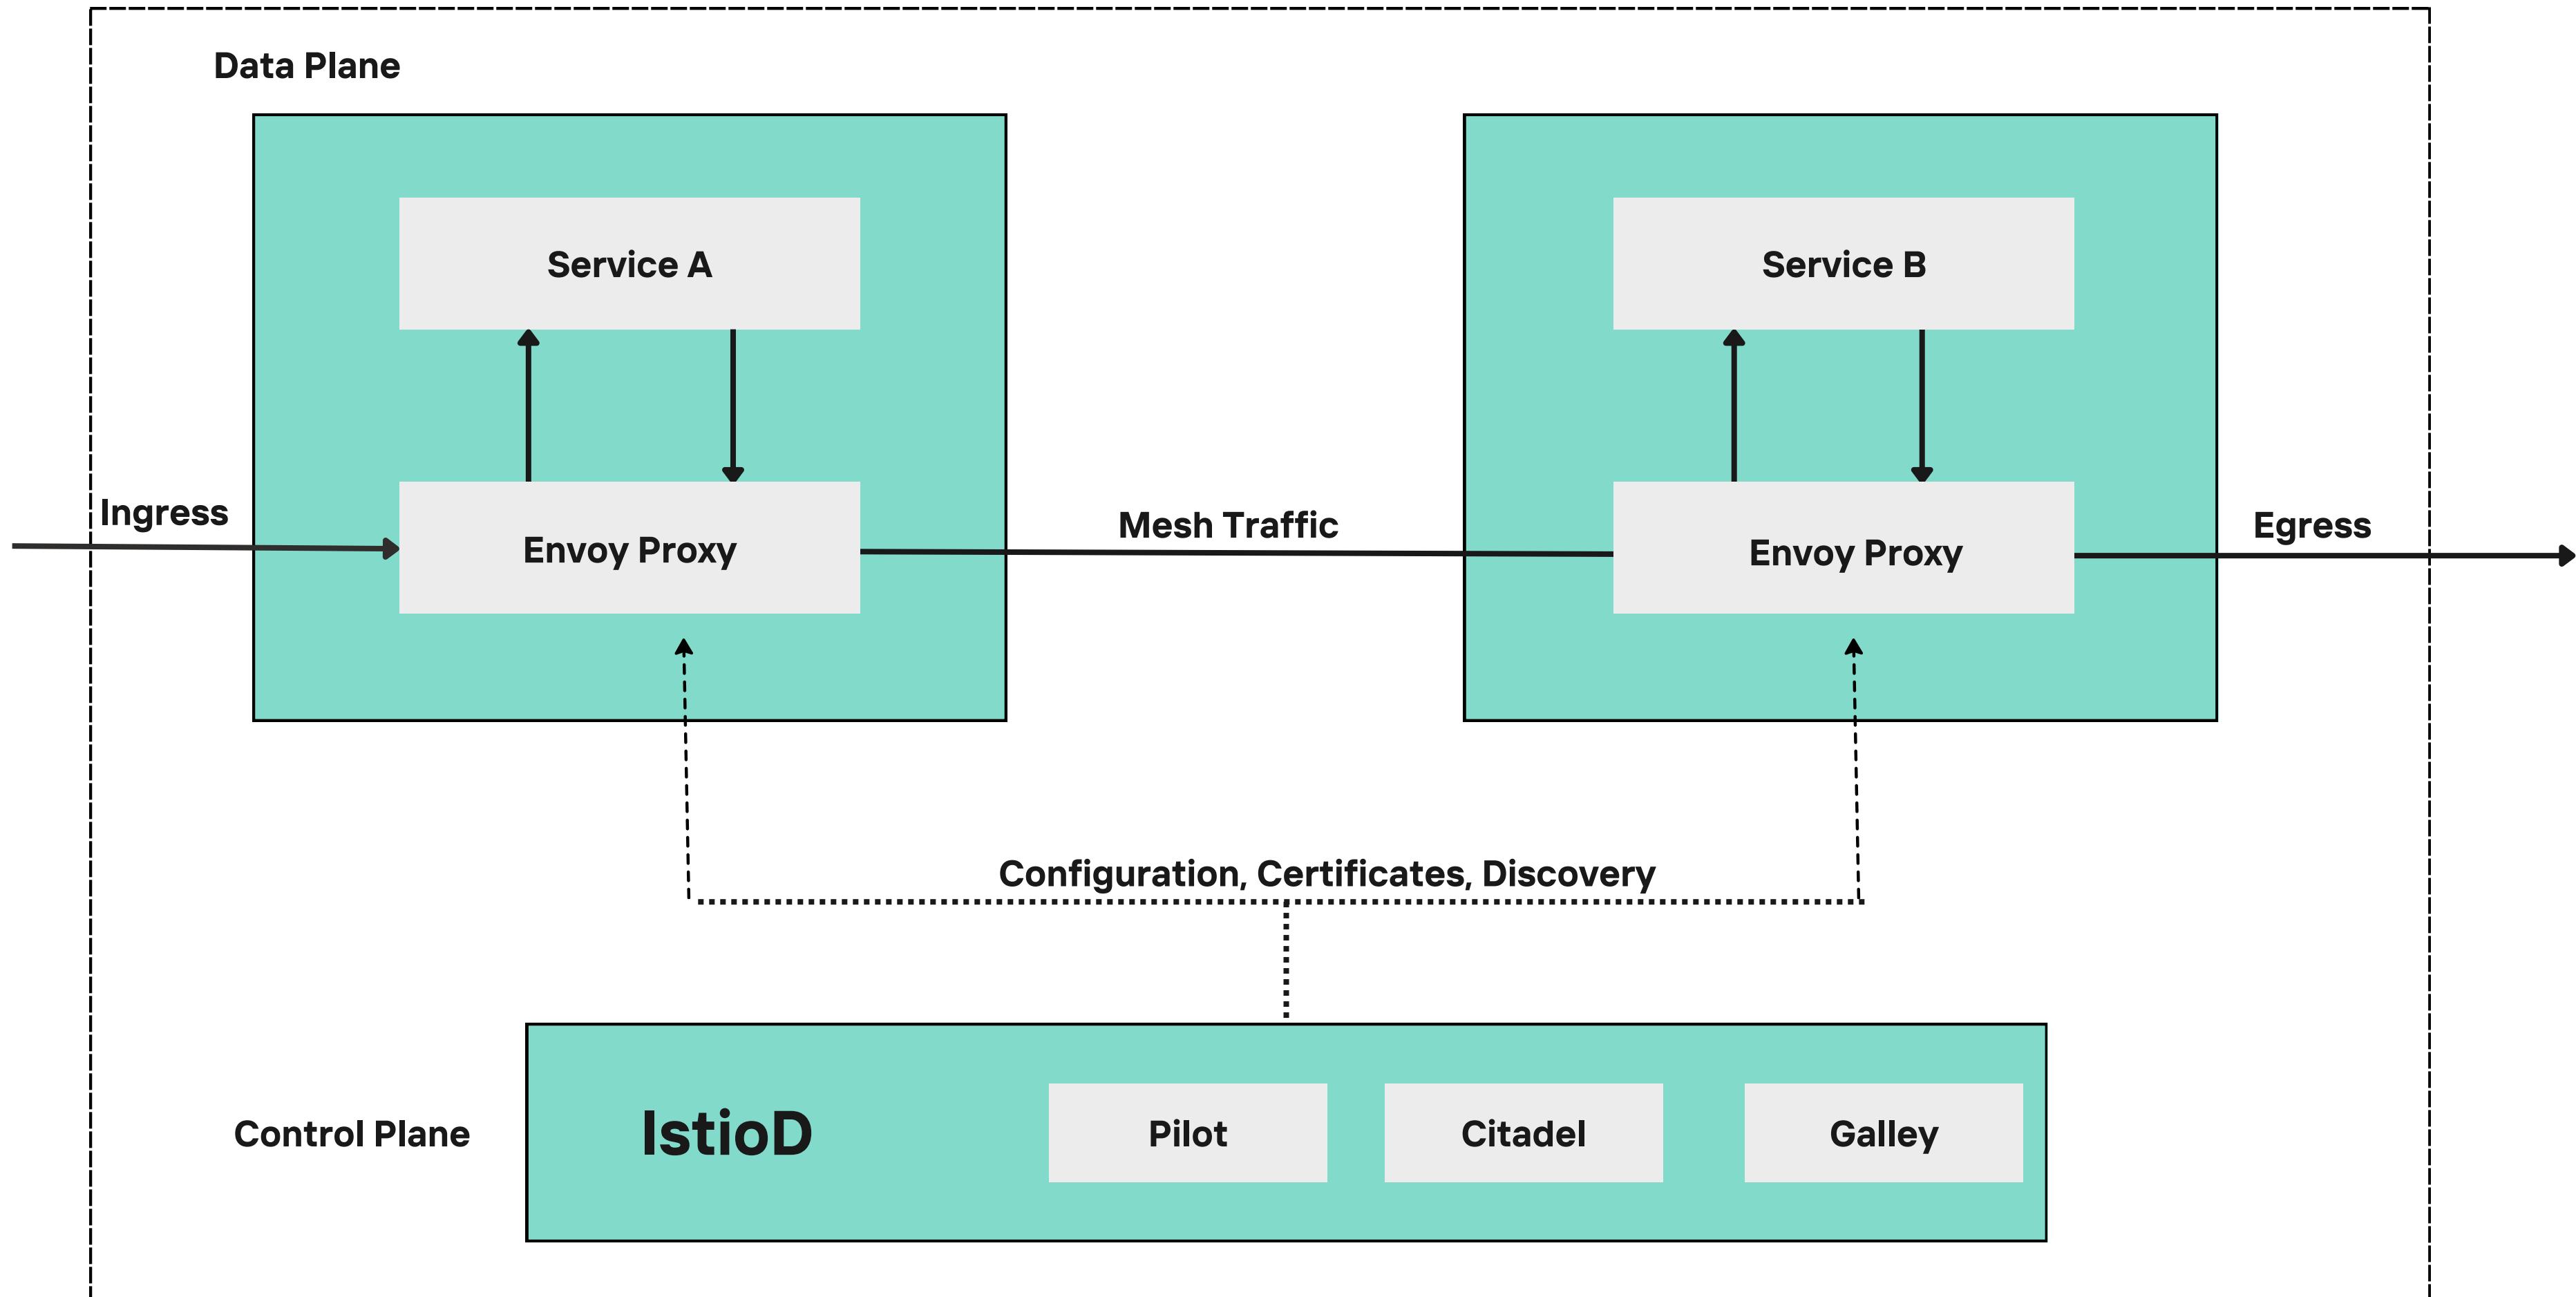


# What is ISTIO?

Istio is an open-source service mesh that simplifies managing microservices by providing tools for traffic management, observability, and security.

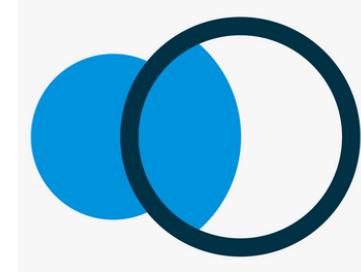
- **Service Mesh:** A layer for managing service-to-service communication.
- **Open Source:** Community-driven and widely supported.





# Observability in Istio

Istio integrates seamlessly with powerful observability tools like Jaeger, Kiali, and Grafana, enhancing the visibility and monitoring of your microservices.



## Kiali

Provides a graphical interface to observe the service mesh, view metrics, and manage microservices



## Jaeger

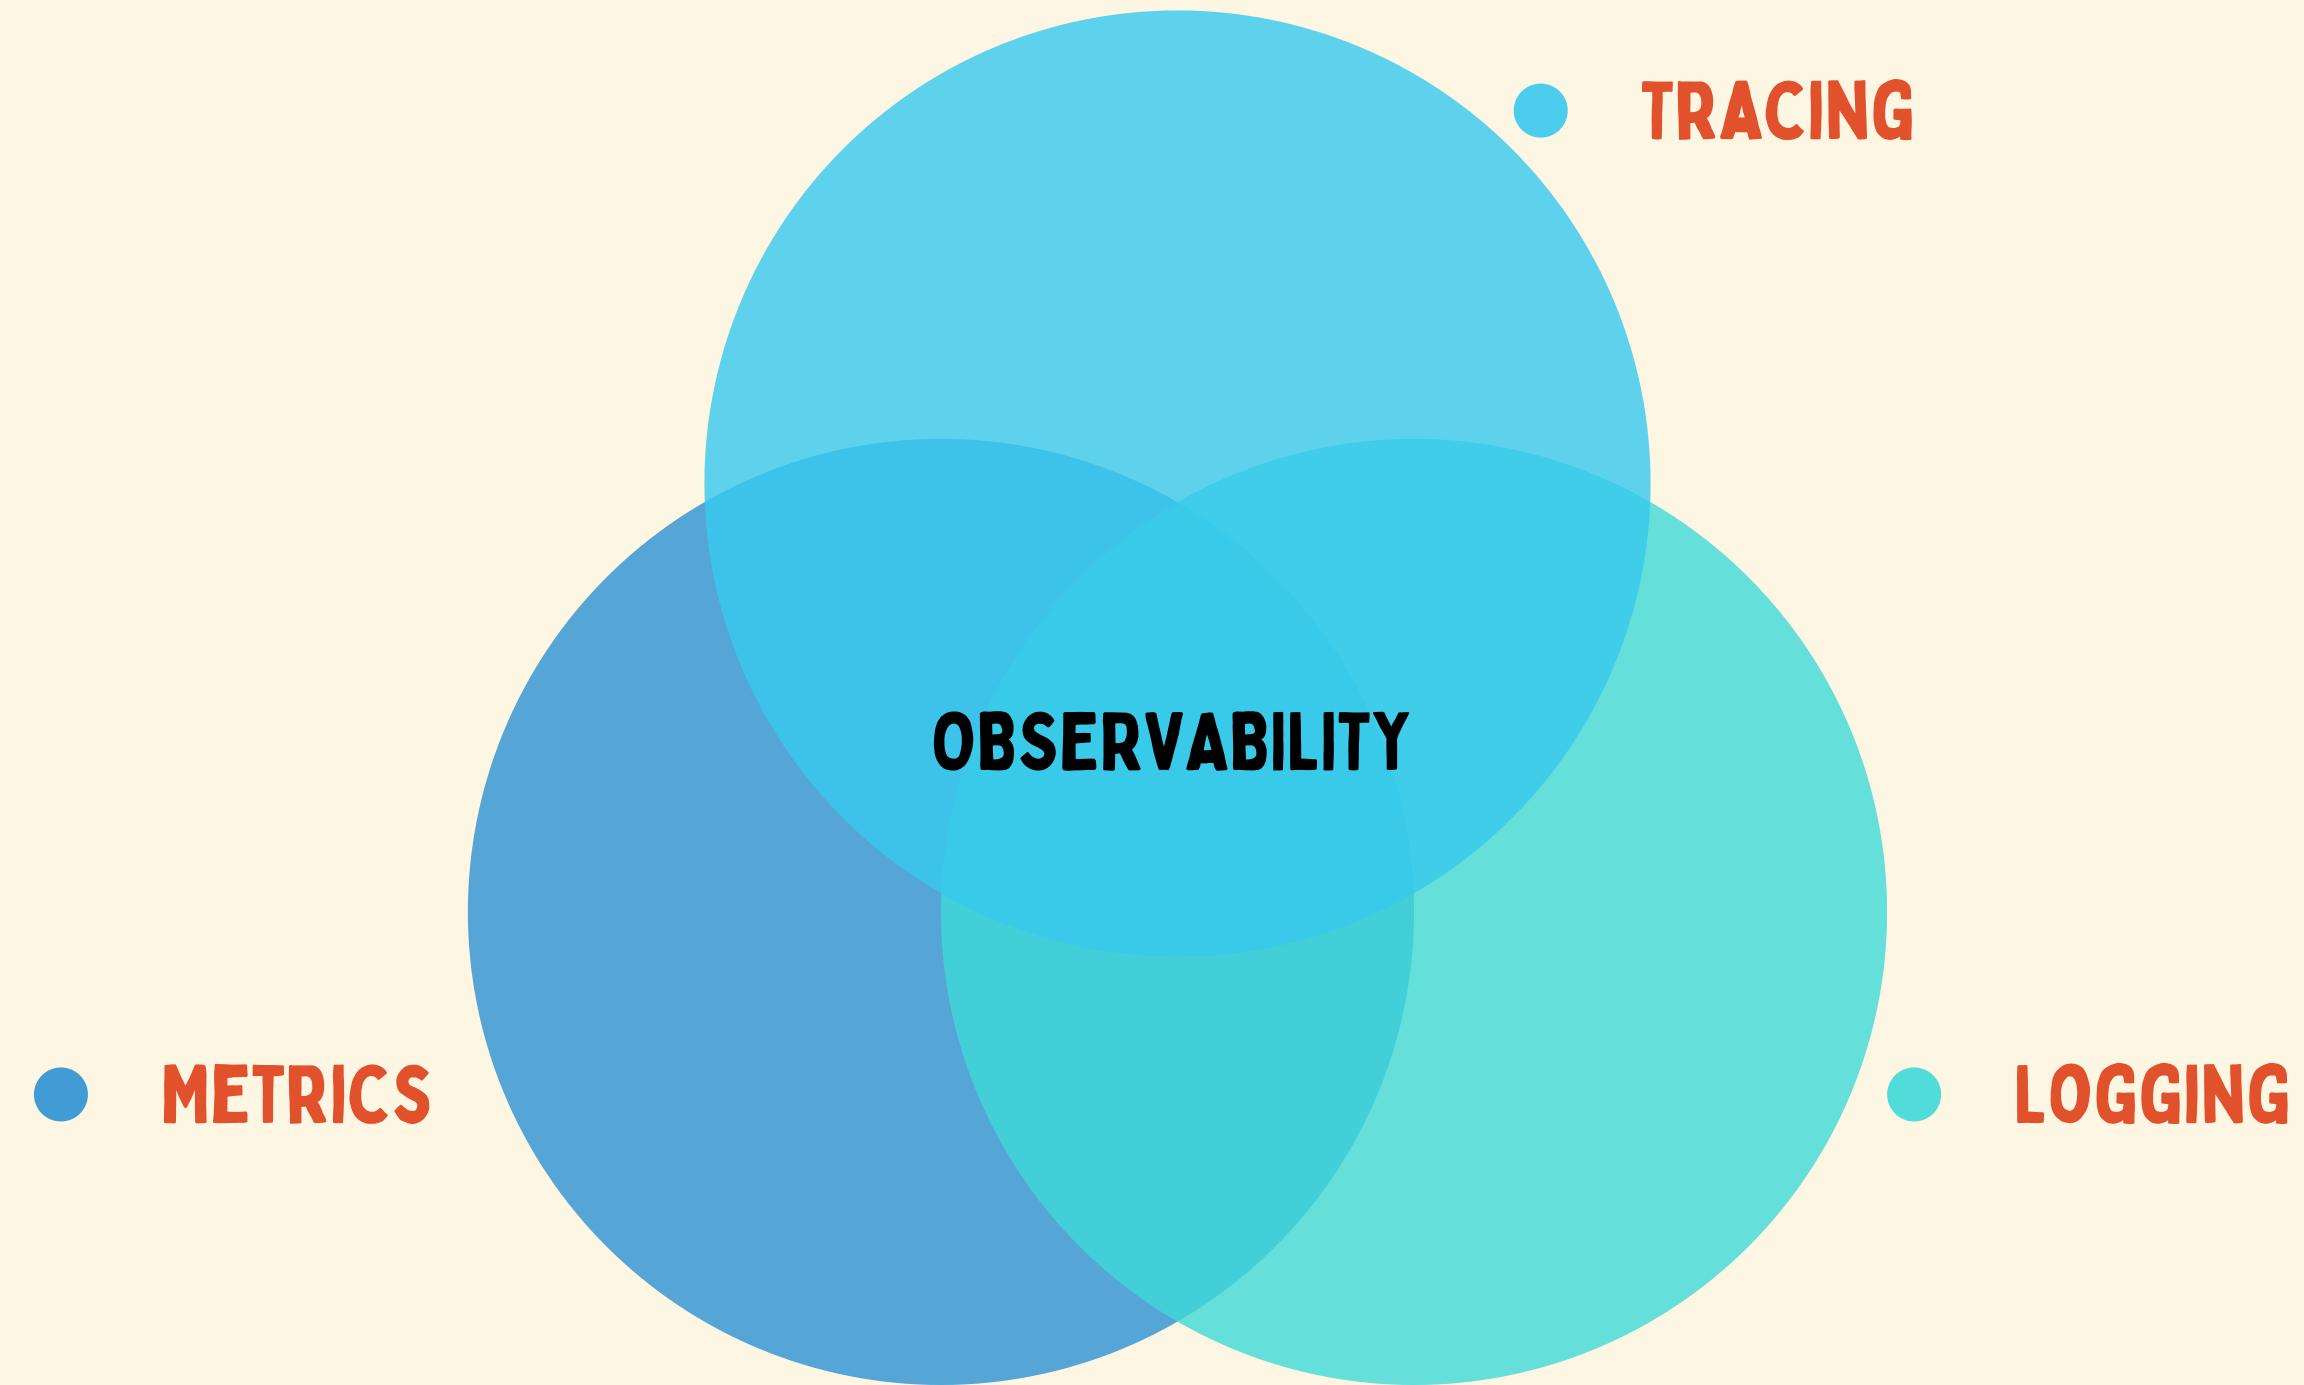
Tracks and visualizes the flow of requests through services, helping identify performance bottlenecks and failures



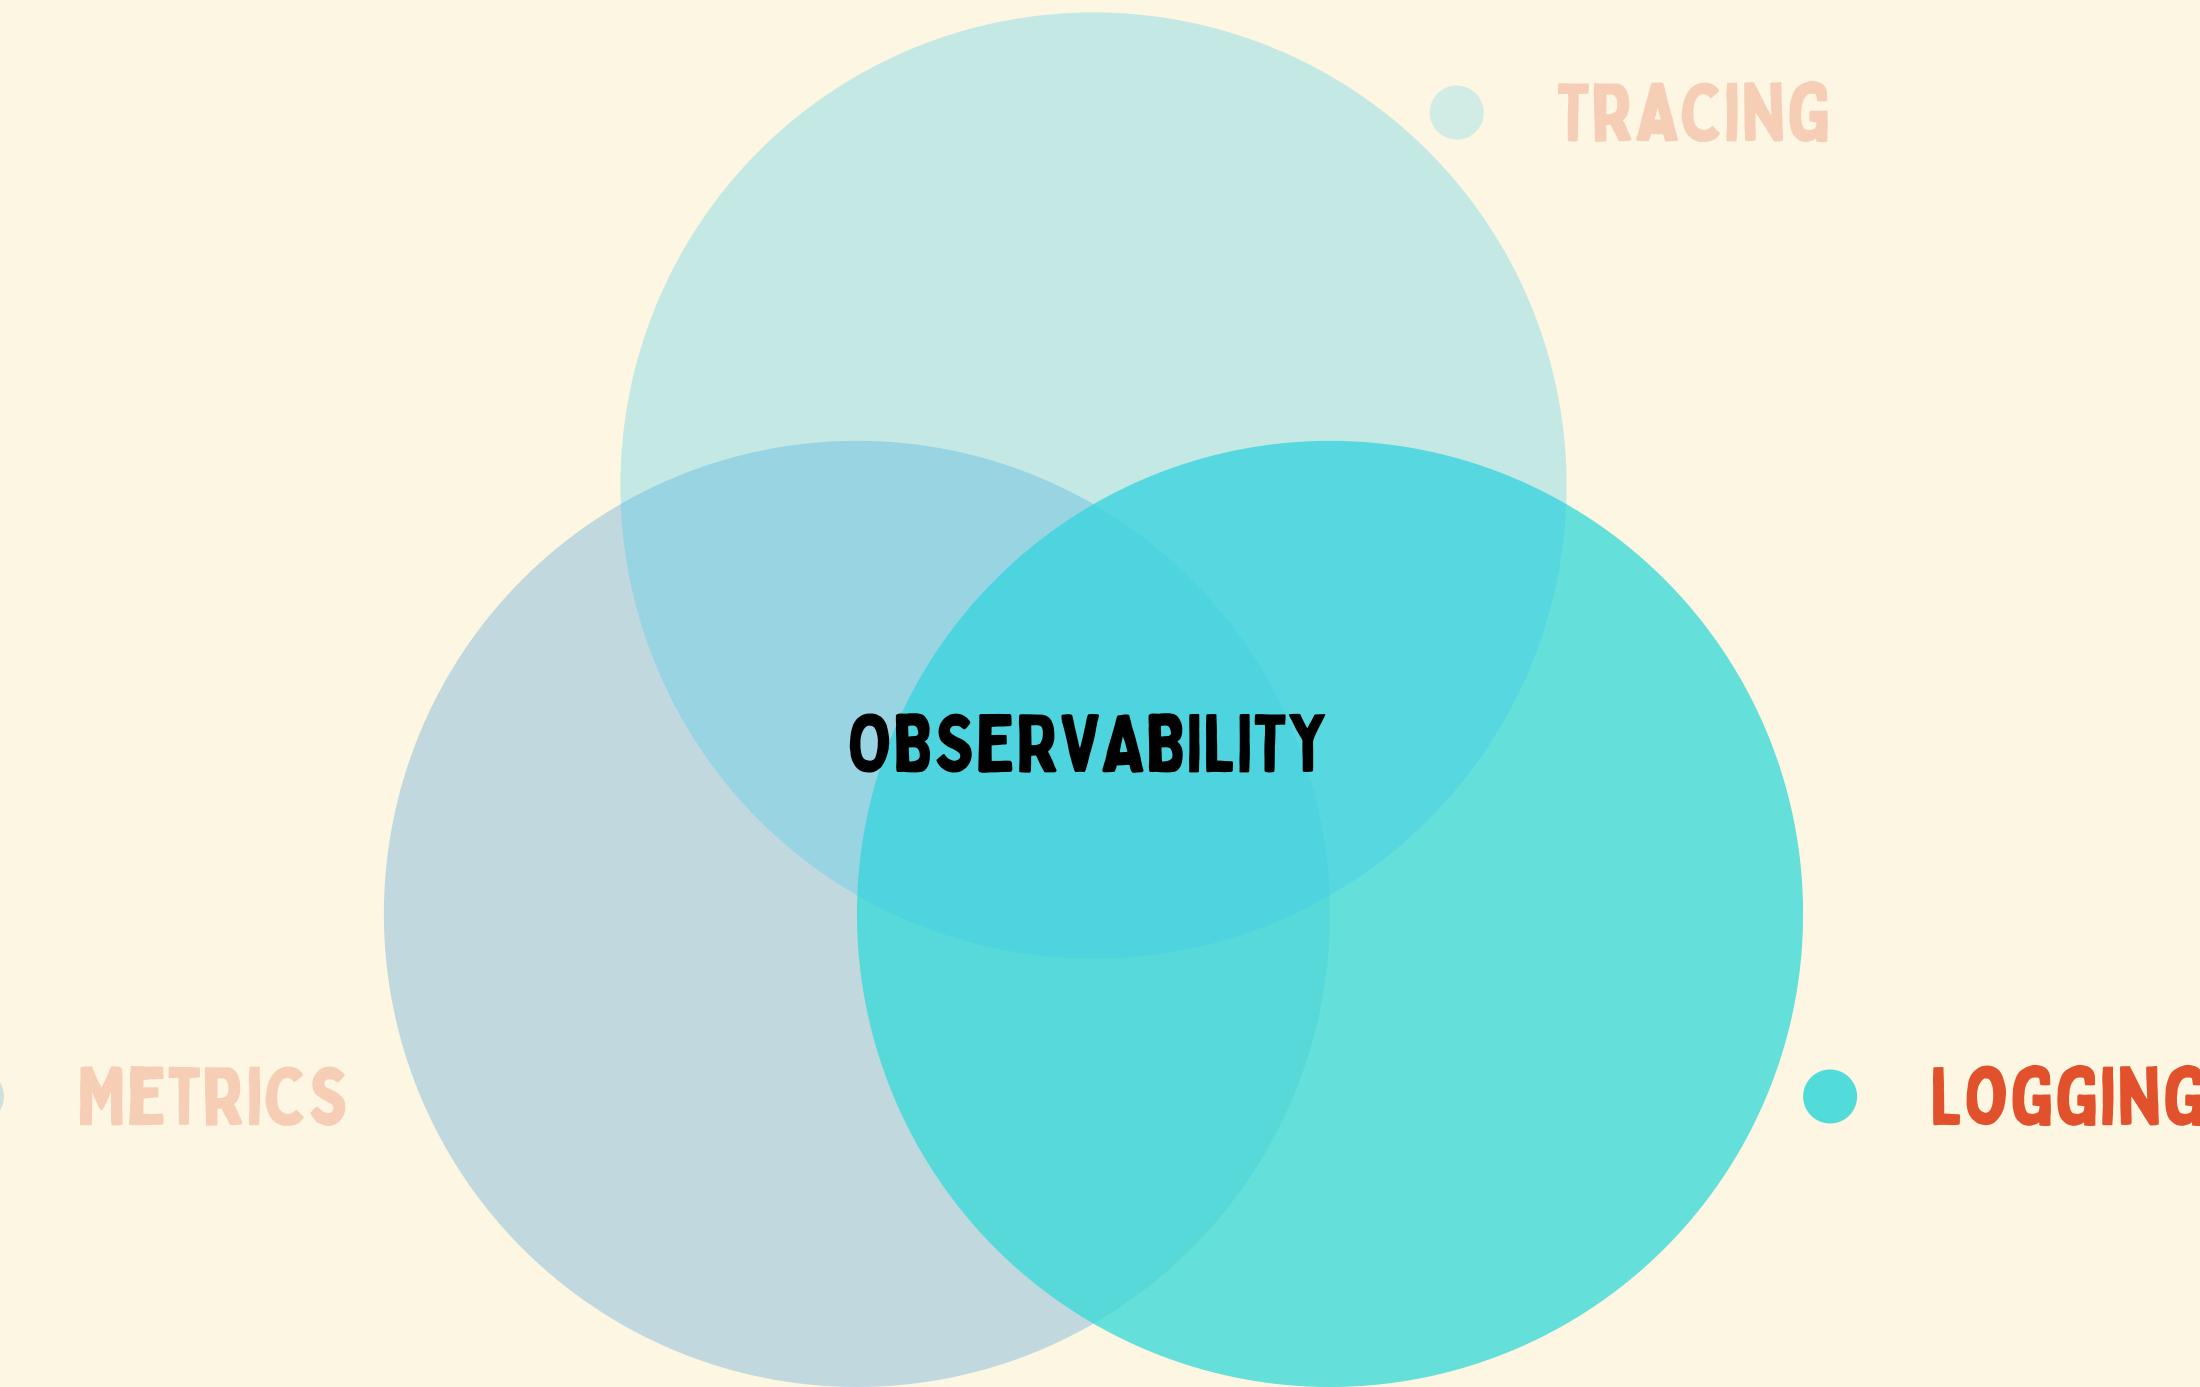
## Kiali

Visualizes metrics collected from Prometheus, providing dashboards for real-time monitoring of system health and performance.

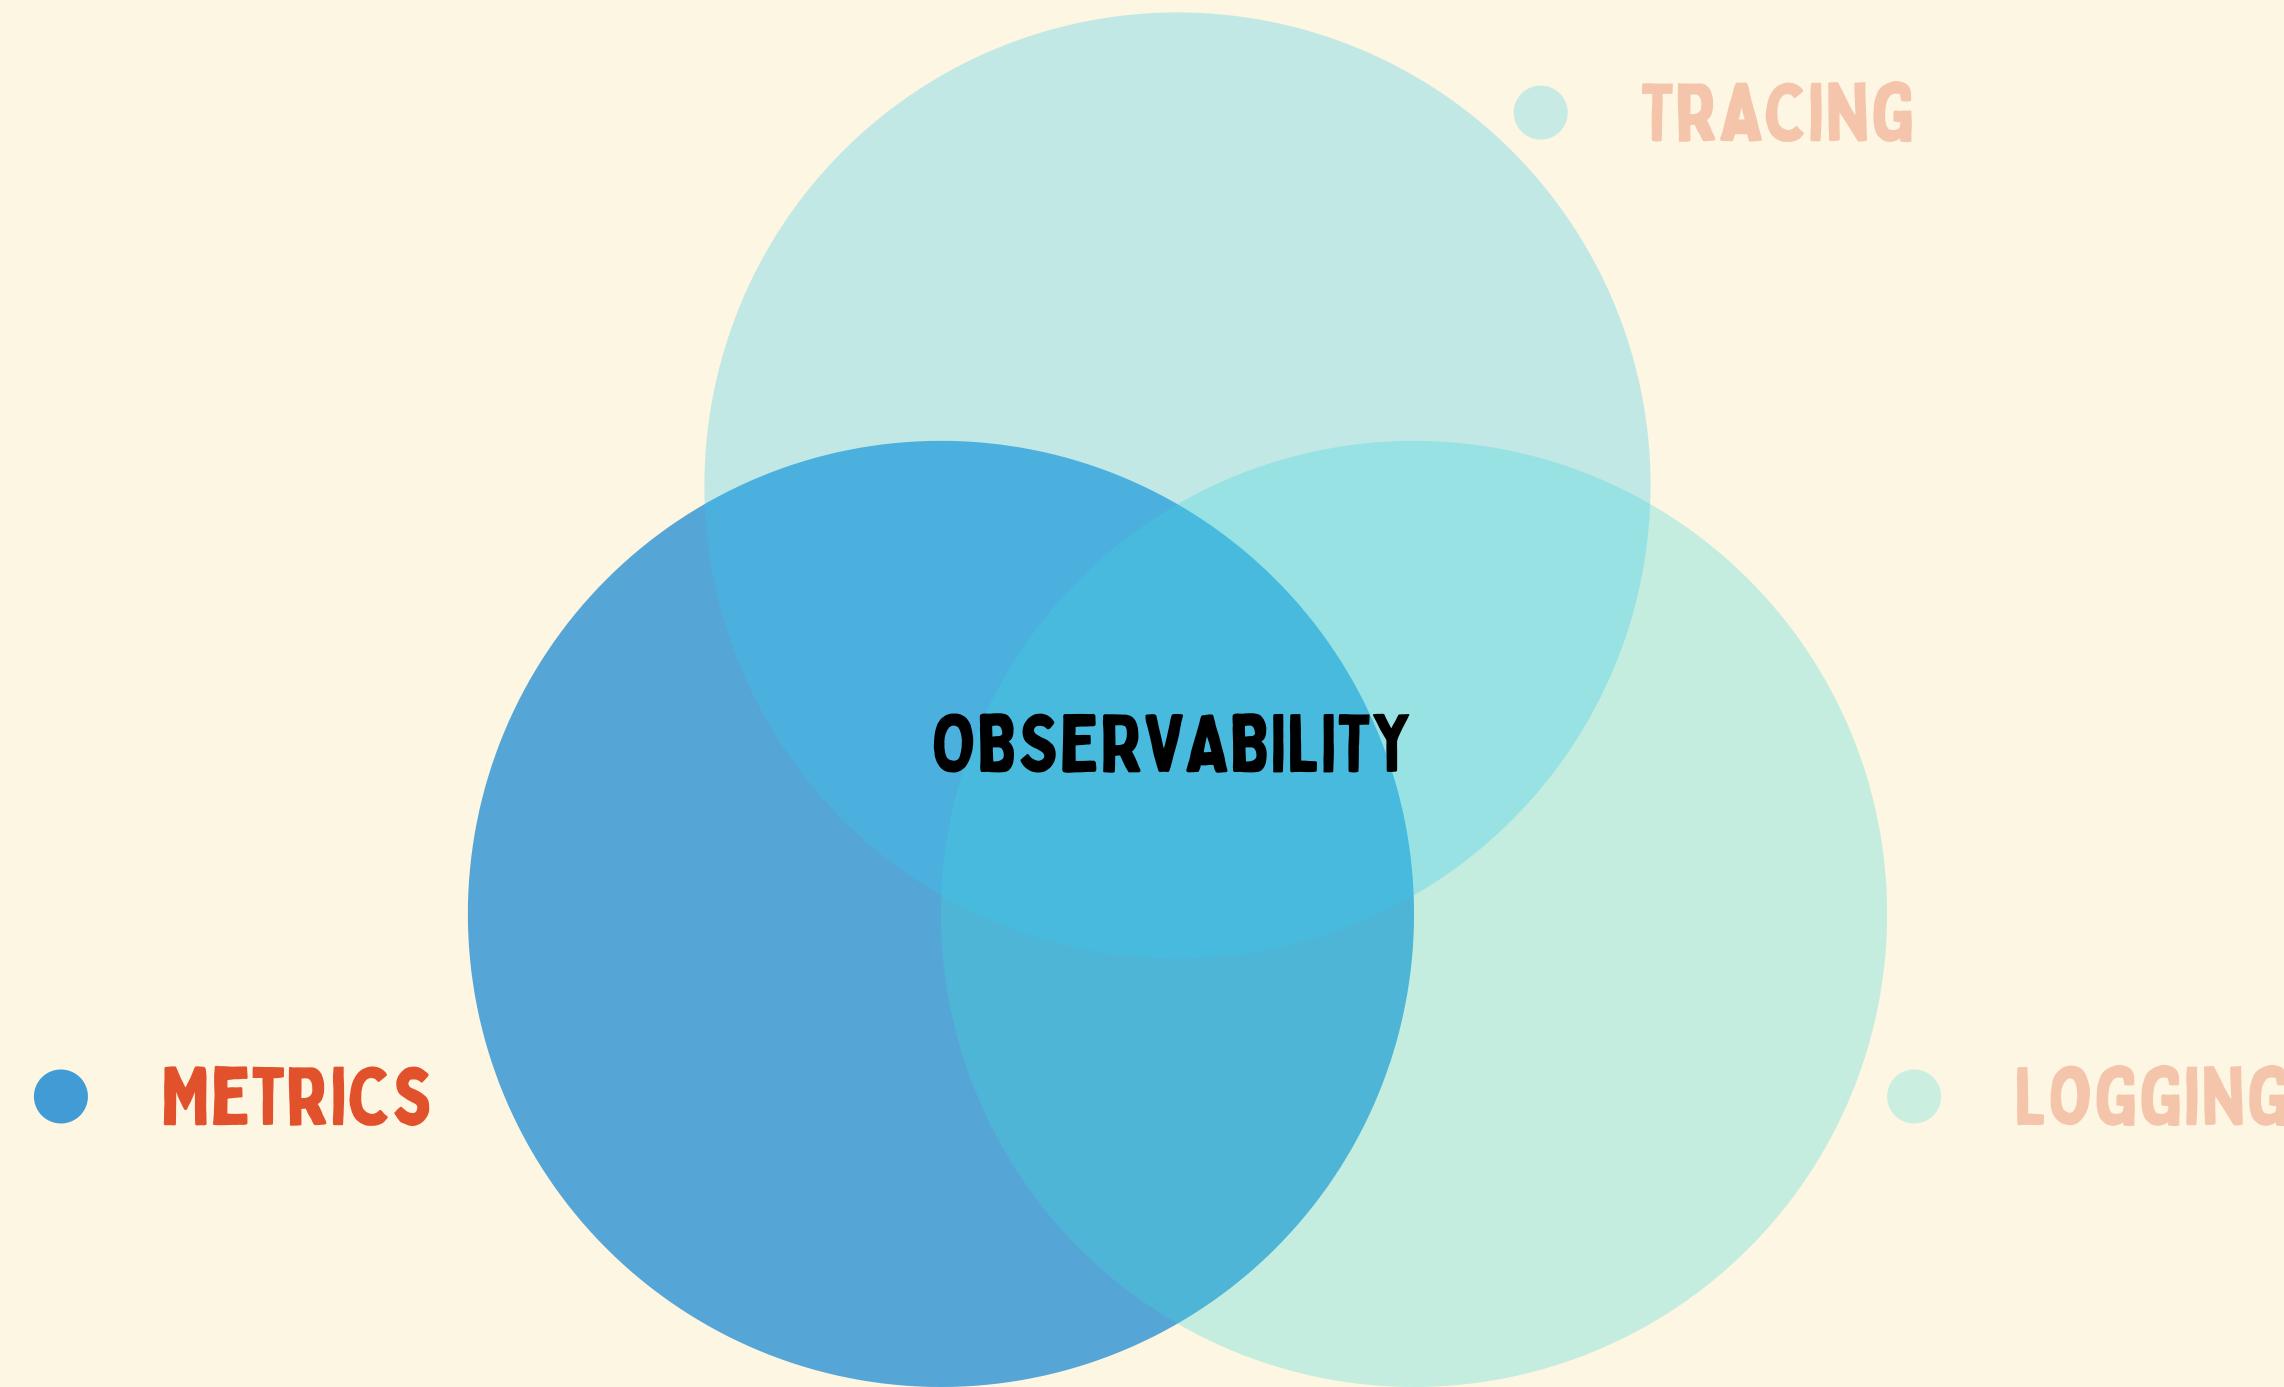
# THREE PILLARS



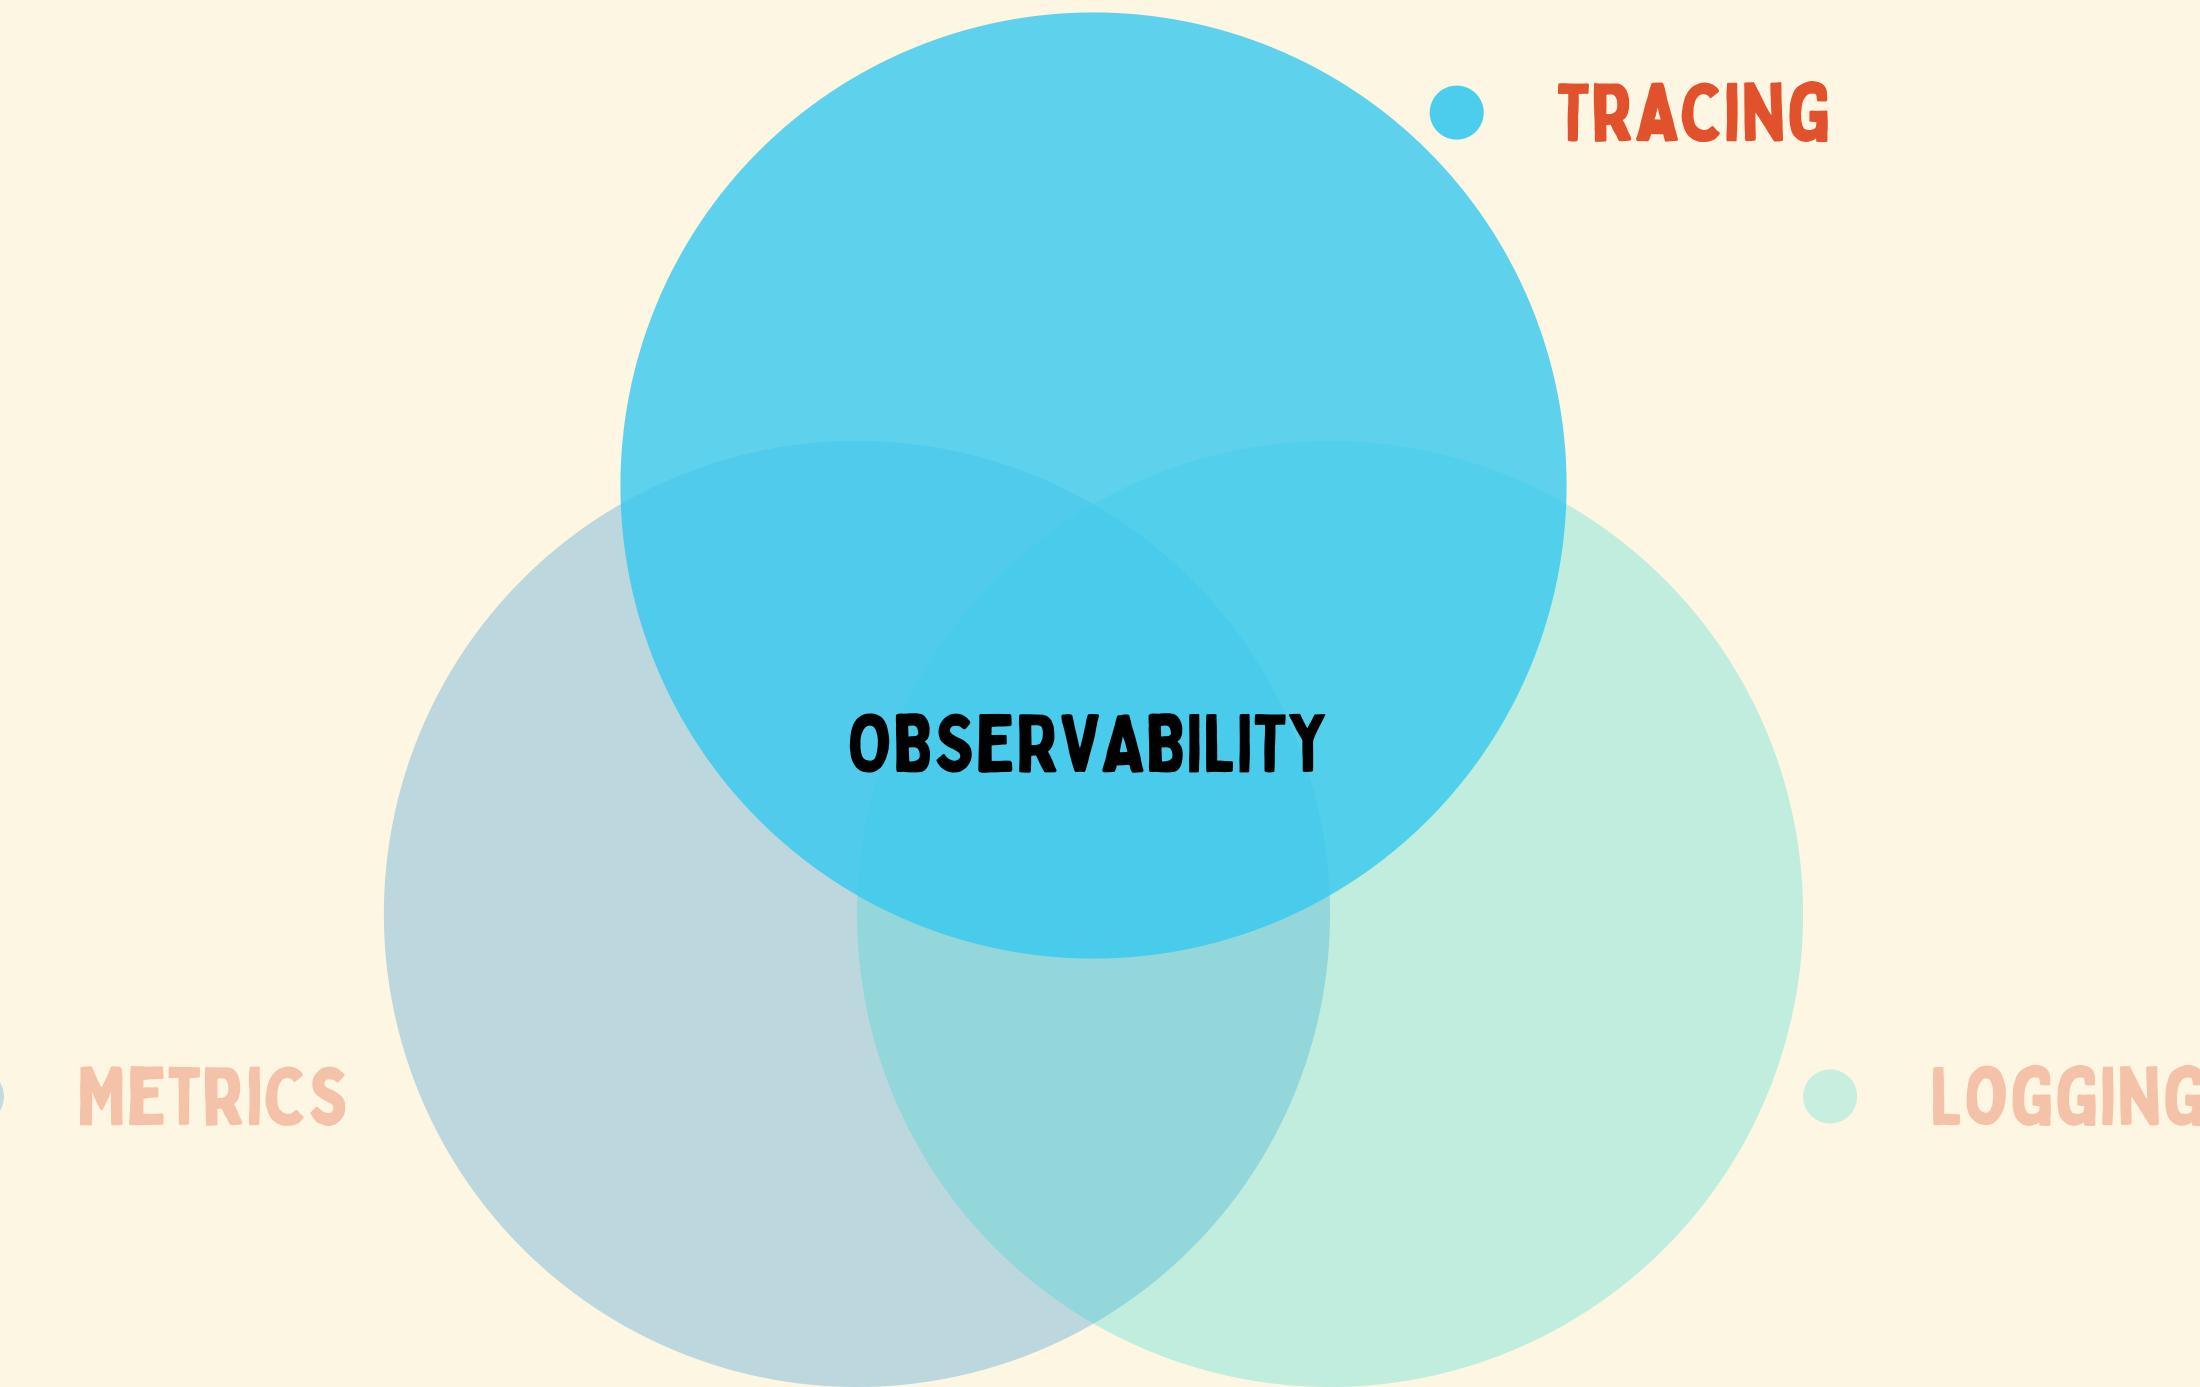
# THREE PILLARS



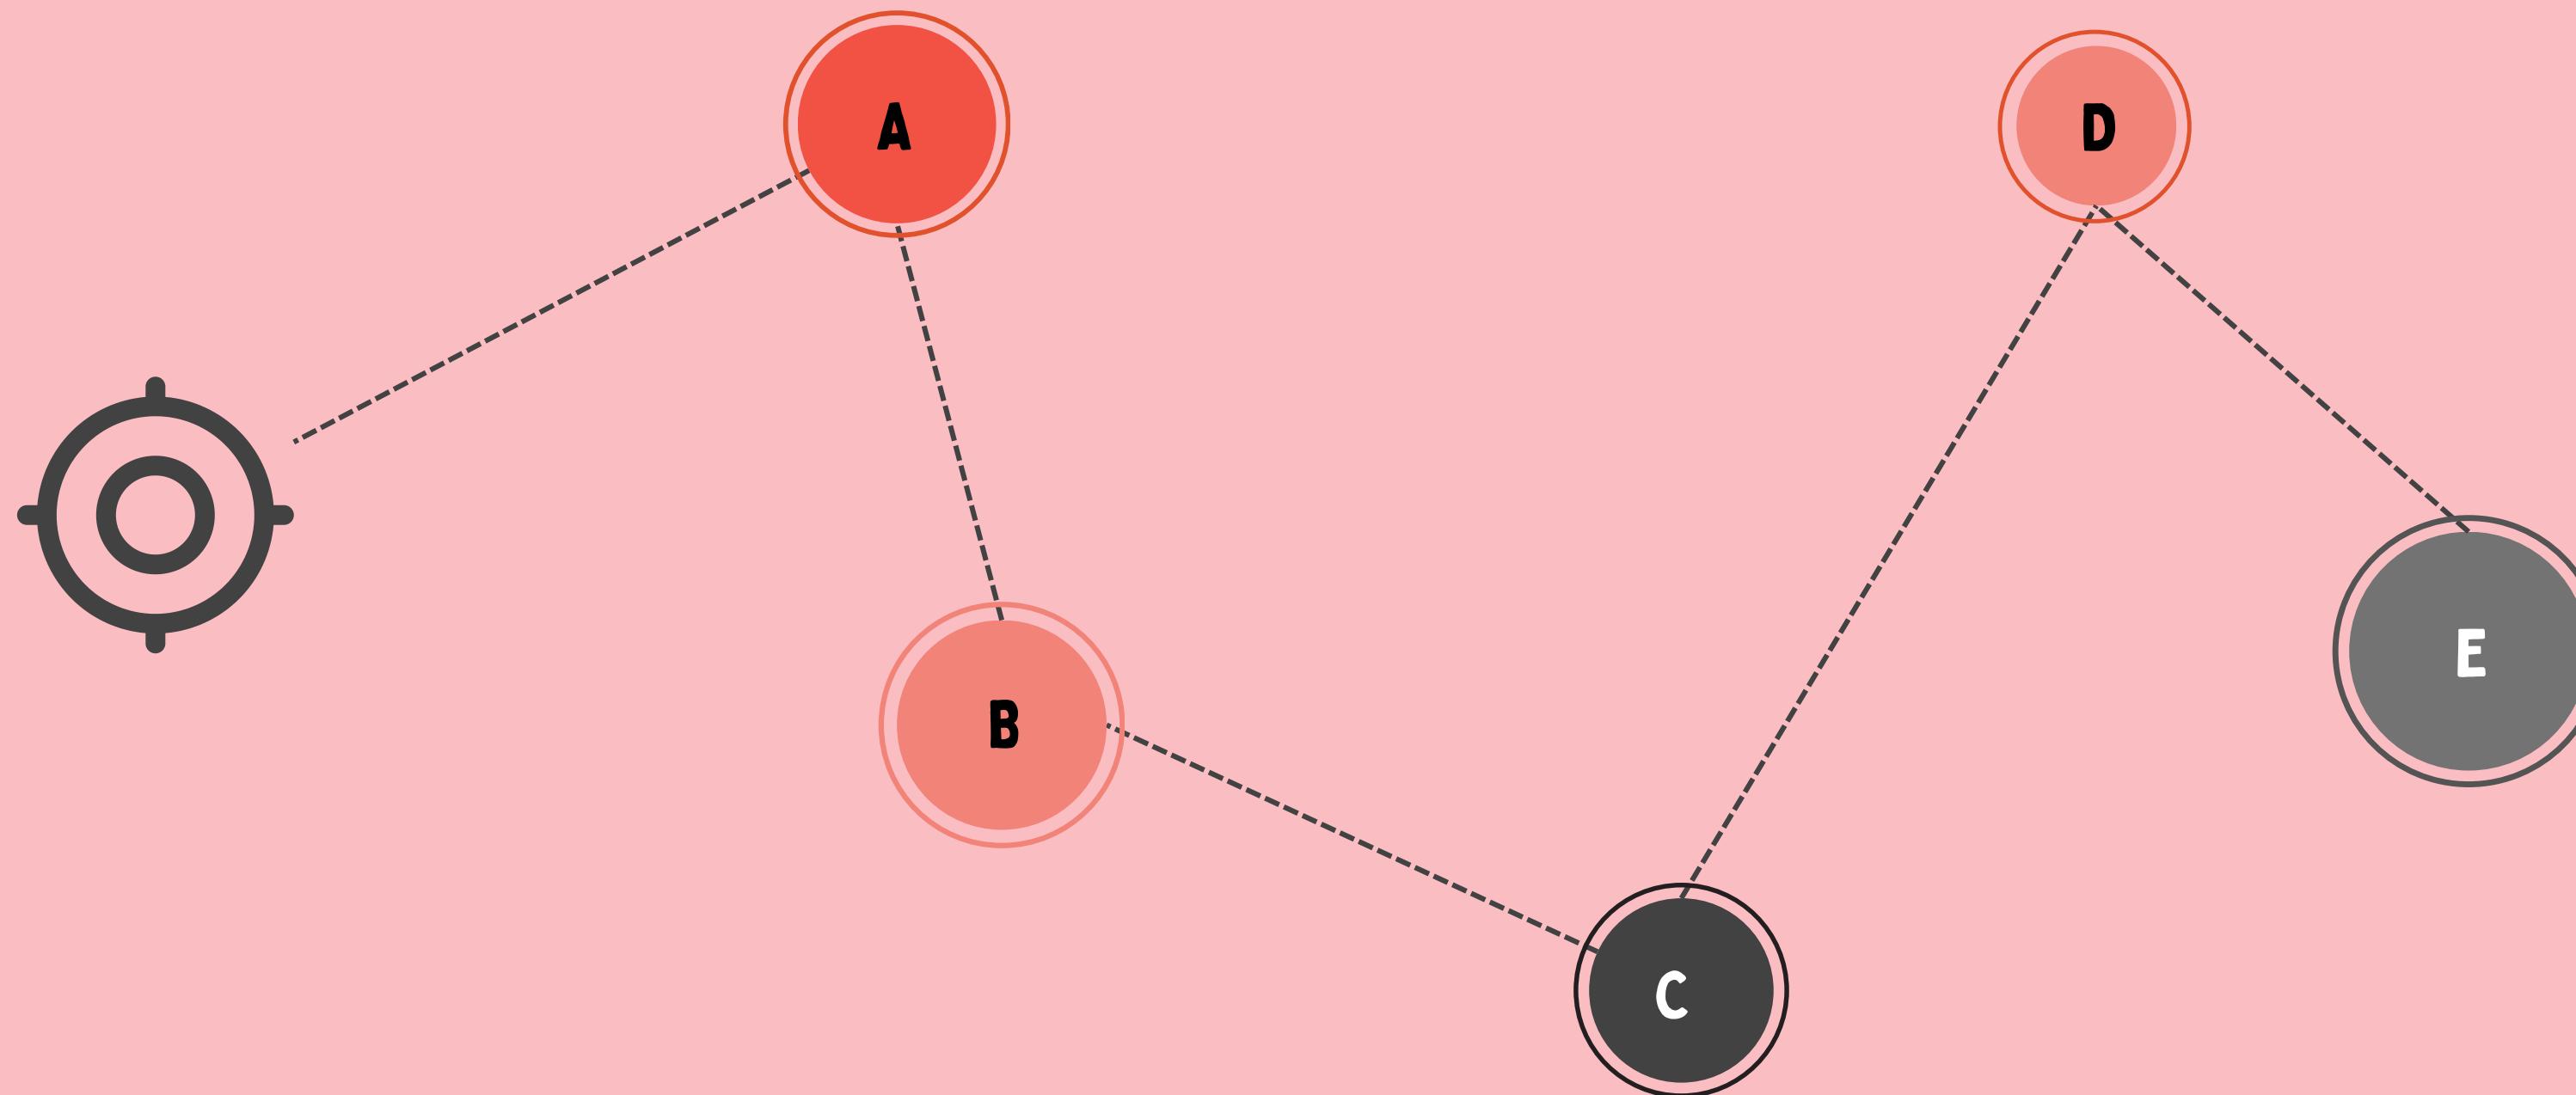
# THREE PILLARS



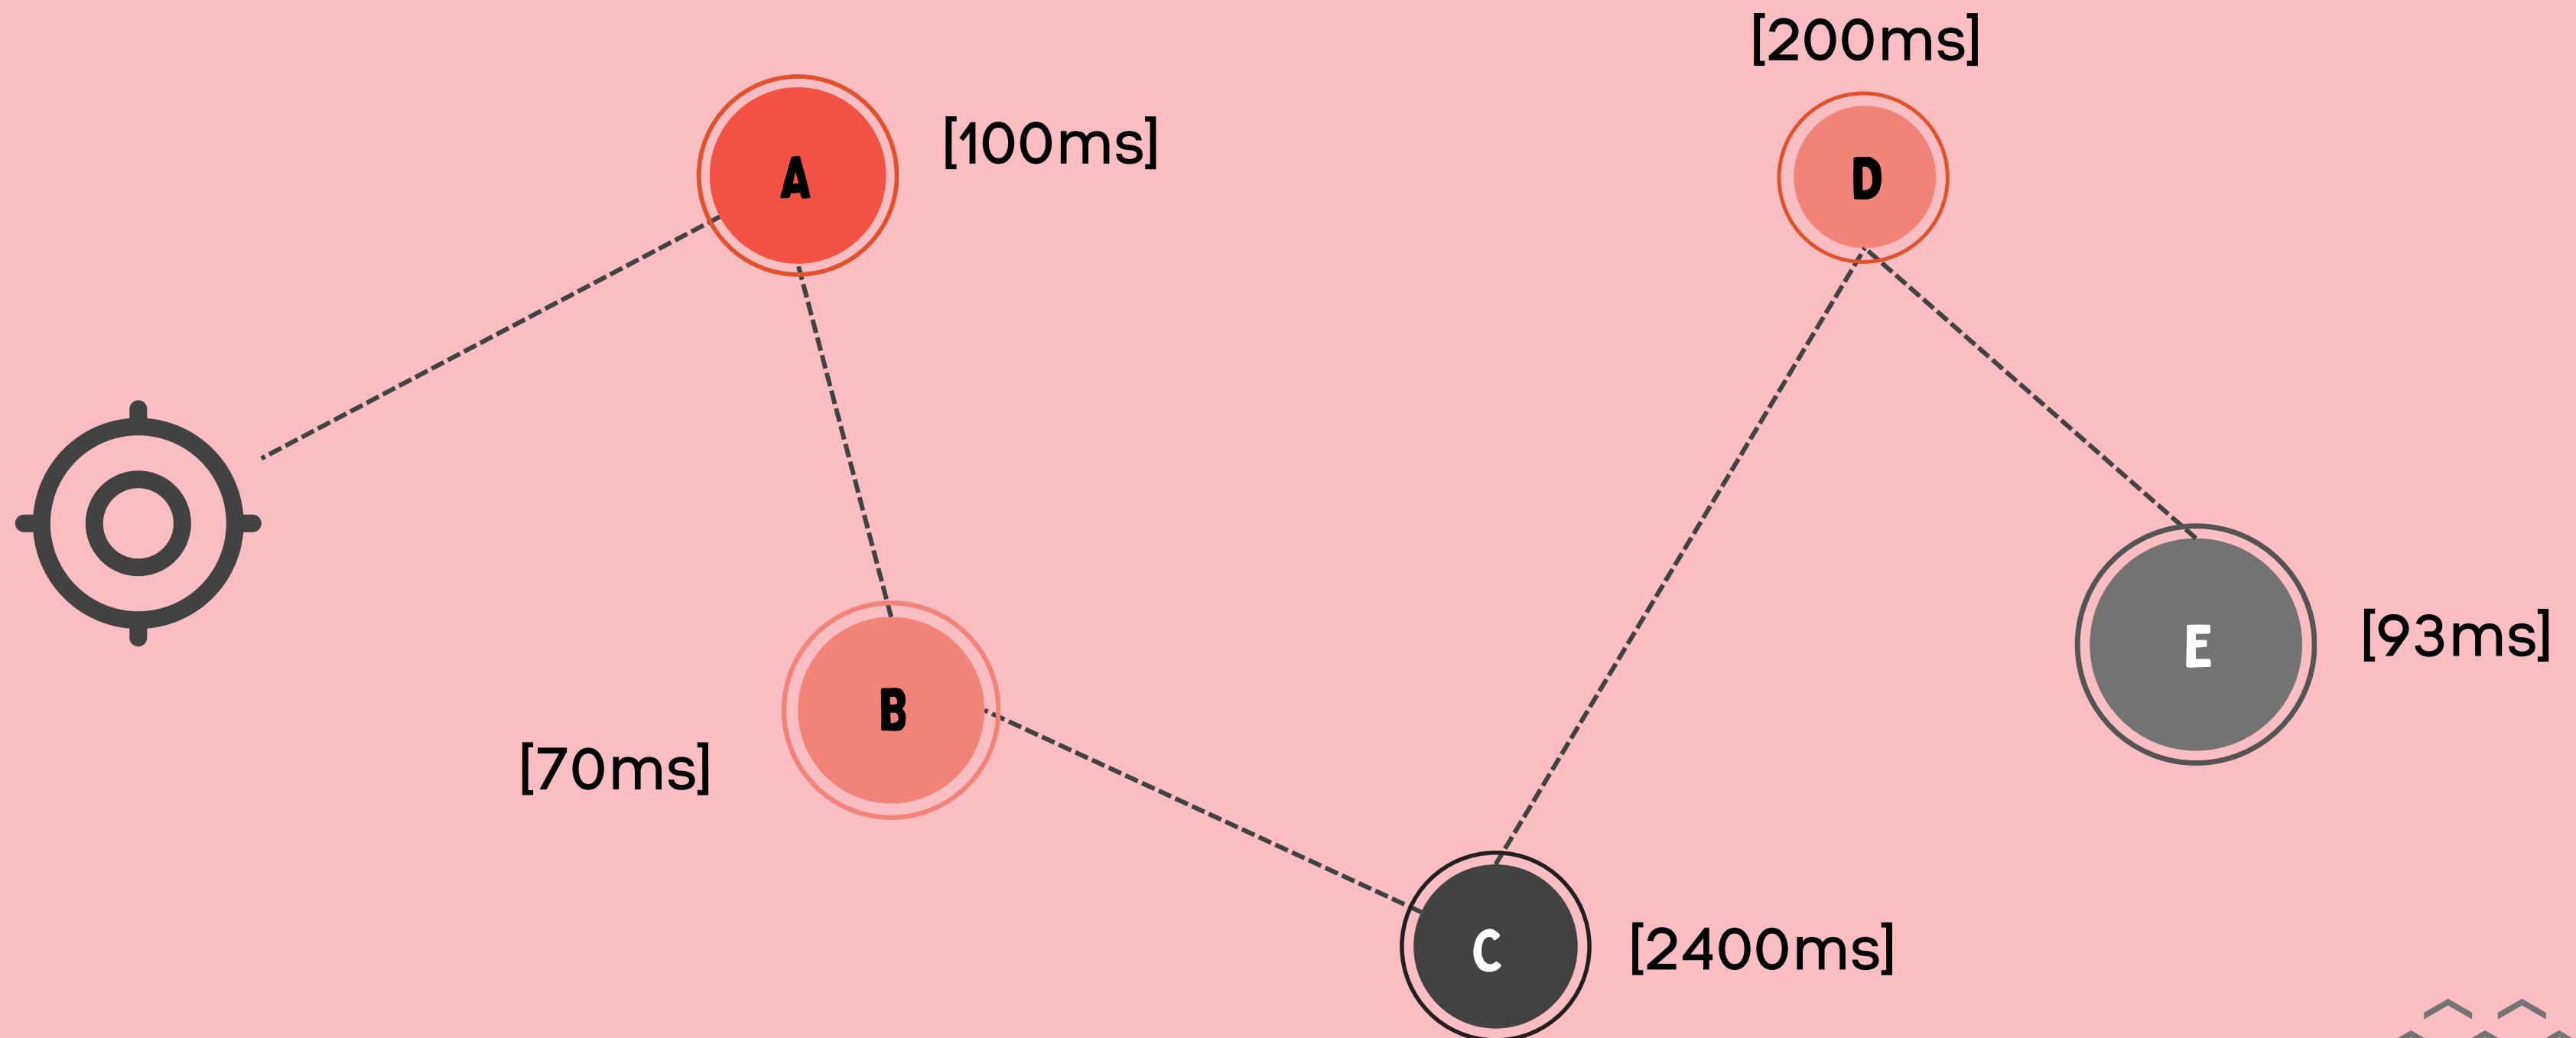
# THREE PILLARS



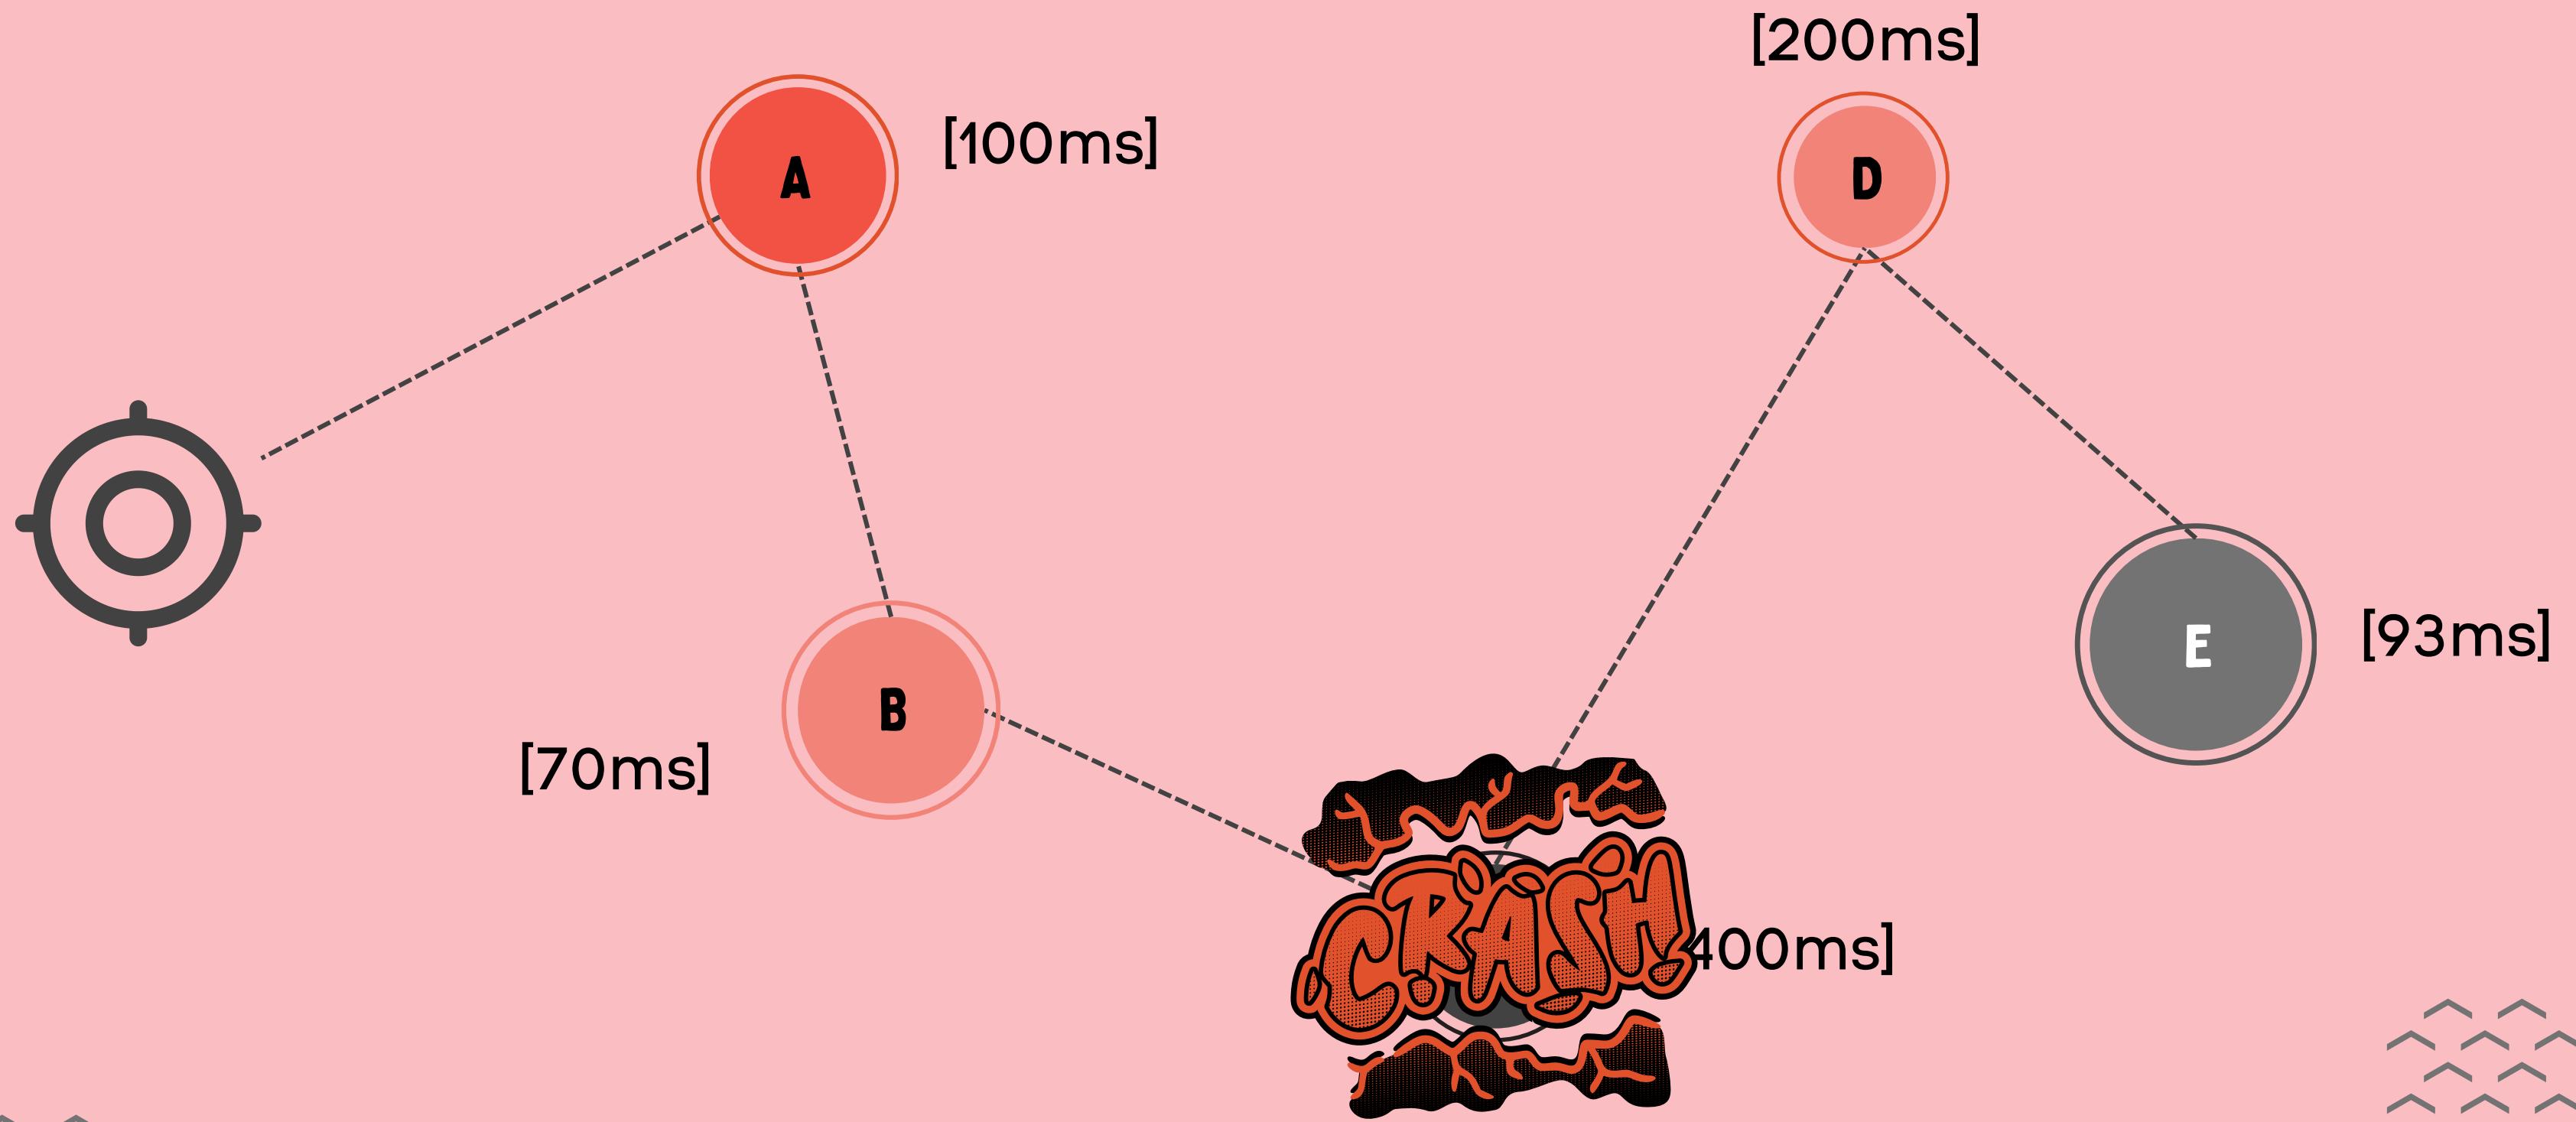
# REQUEST



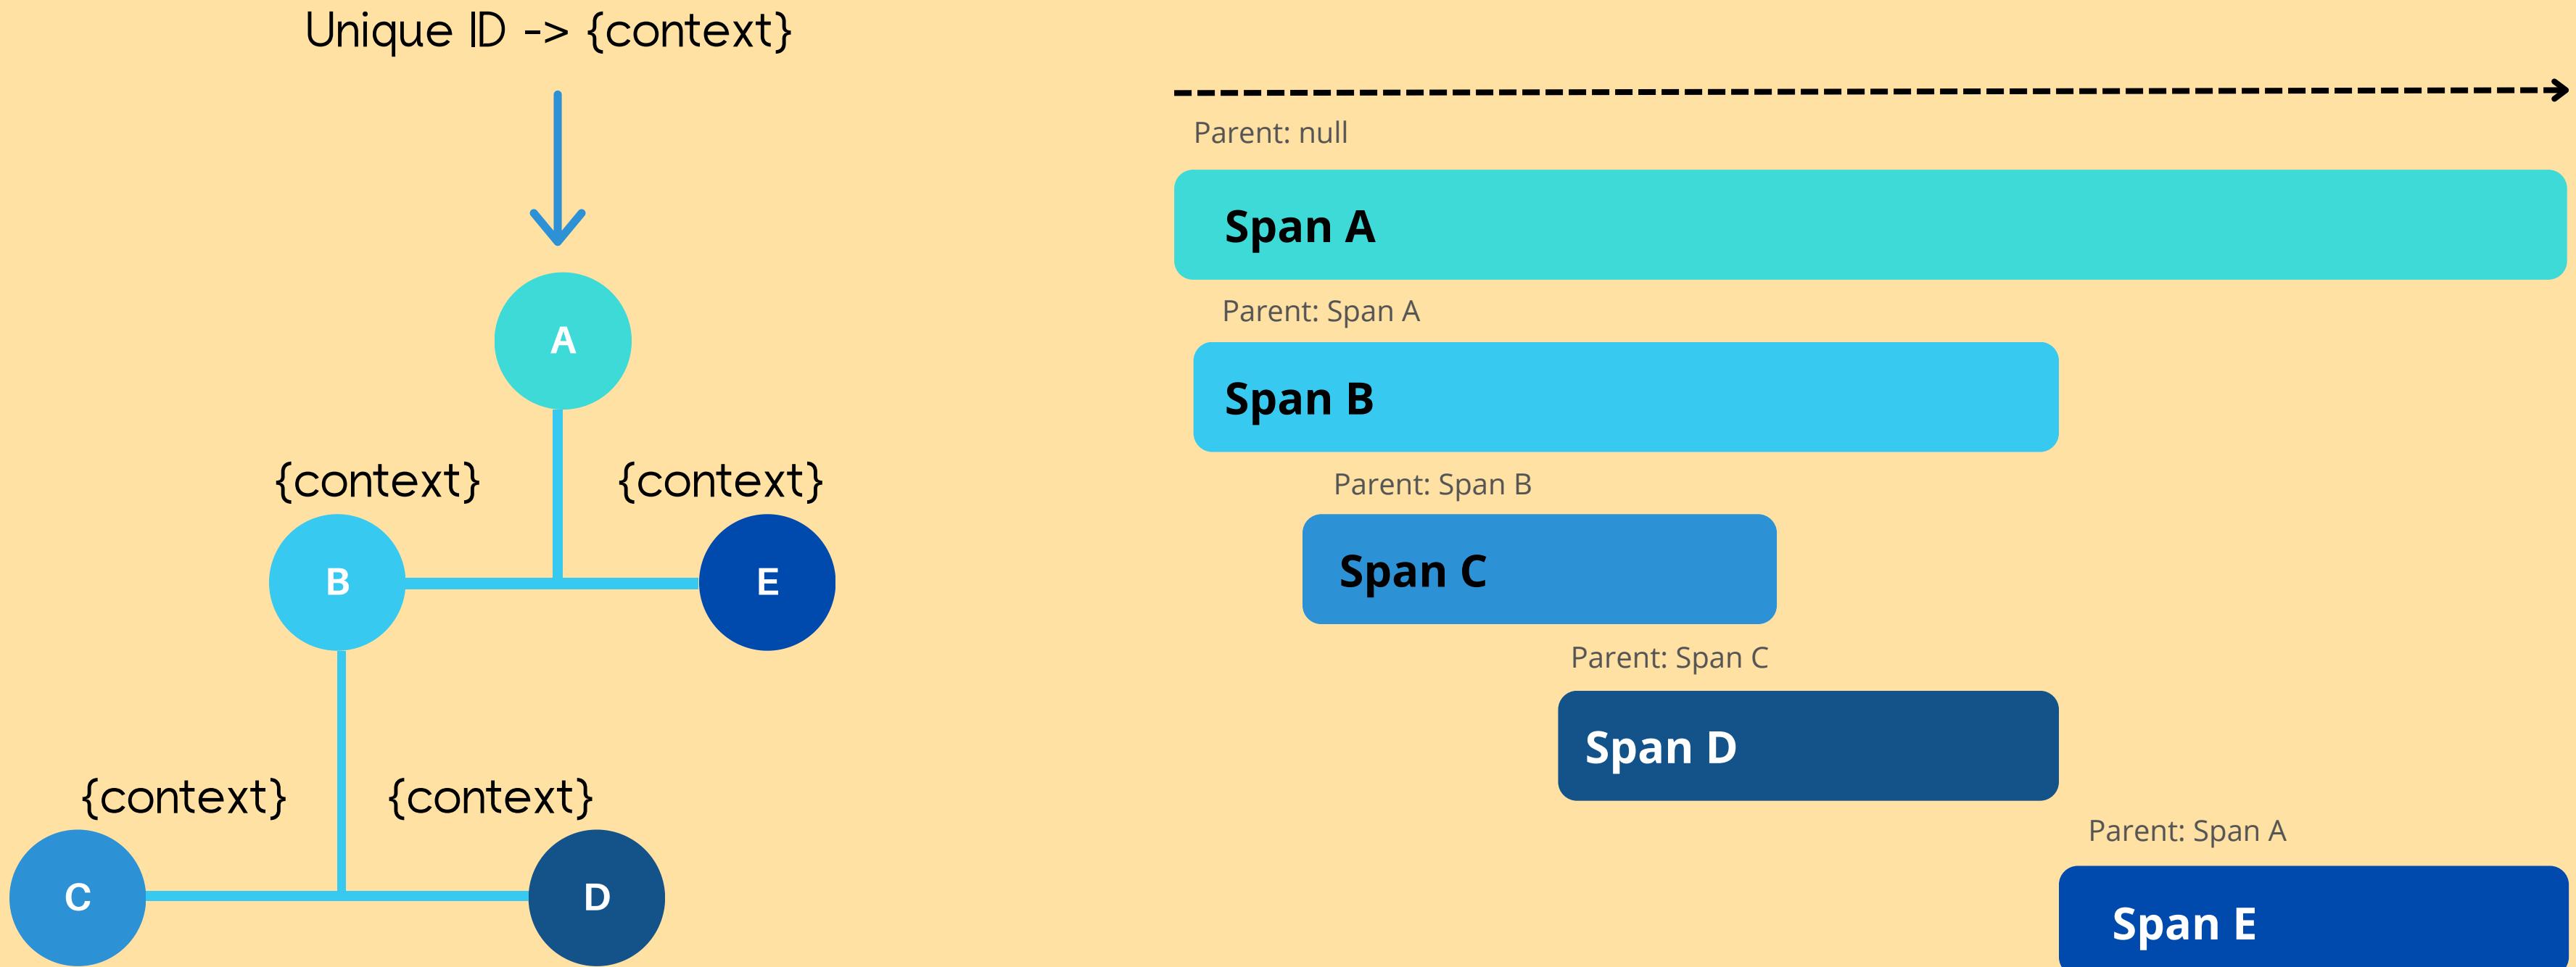
# REQUEST



# REQUEST

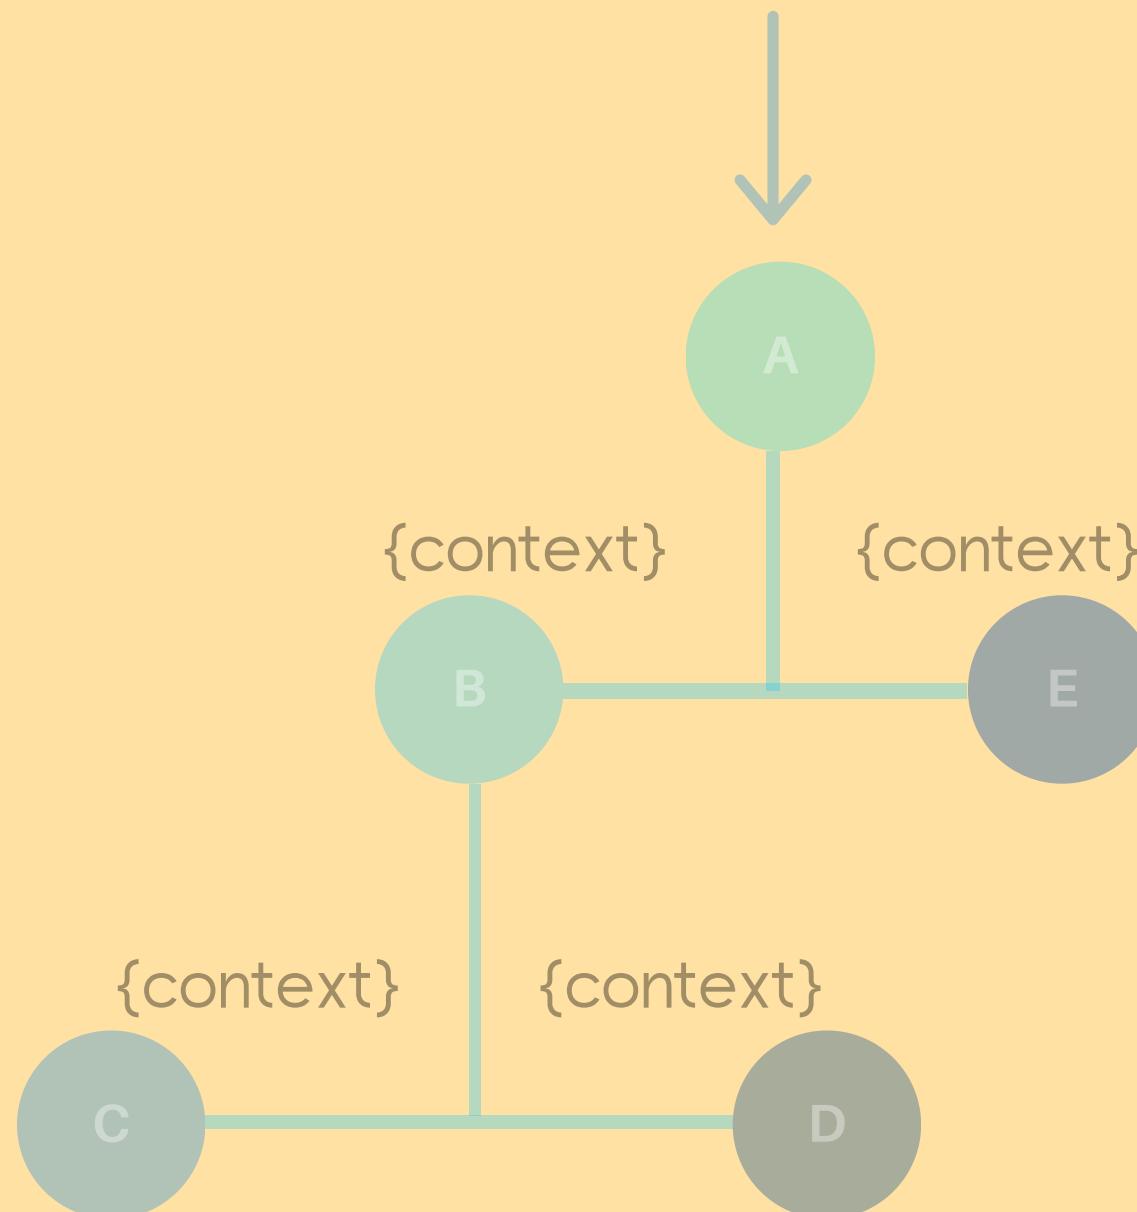


# SPANS



# TRACE

Unique ID -> {context}



TRACE

Parent: null

**Span A**

Trace ID: X

Parent: Span A

**Span B**

Trace ID: X

Parent: Span B

**Span C**

Trace ID: X

Parent: Span C

**Span D**

Trace ID: X

Parent: Span A

**Span E**

Trace ID: X

# How do we configure Istio?

## CRD (Custom Resource Definition)

CRDs allow you to define and manage Istio-specific configurations, such as VirtualServices, DestinationRules, and Gateways, providing fine-grained control over traffic management and policies.

## Traffic Management

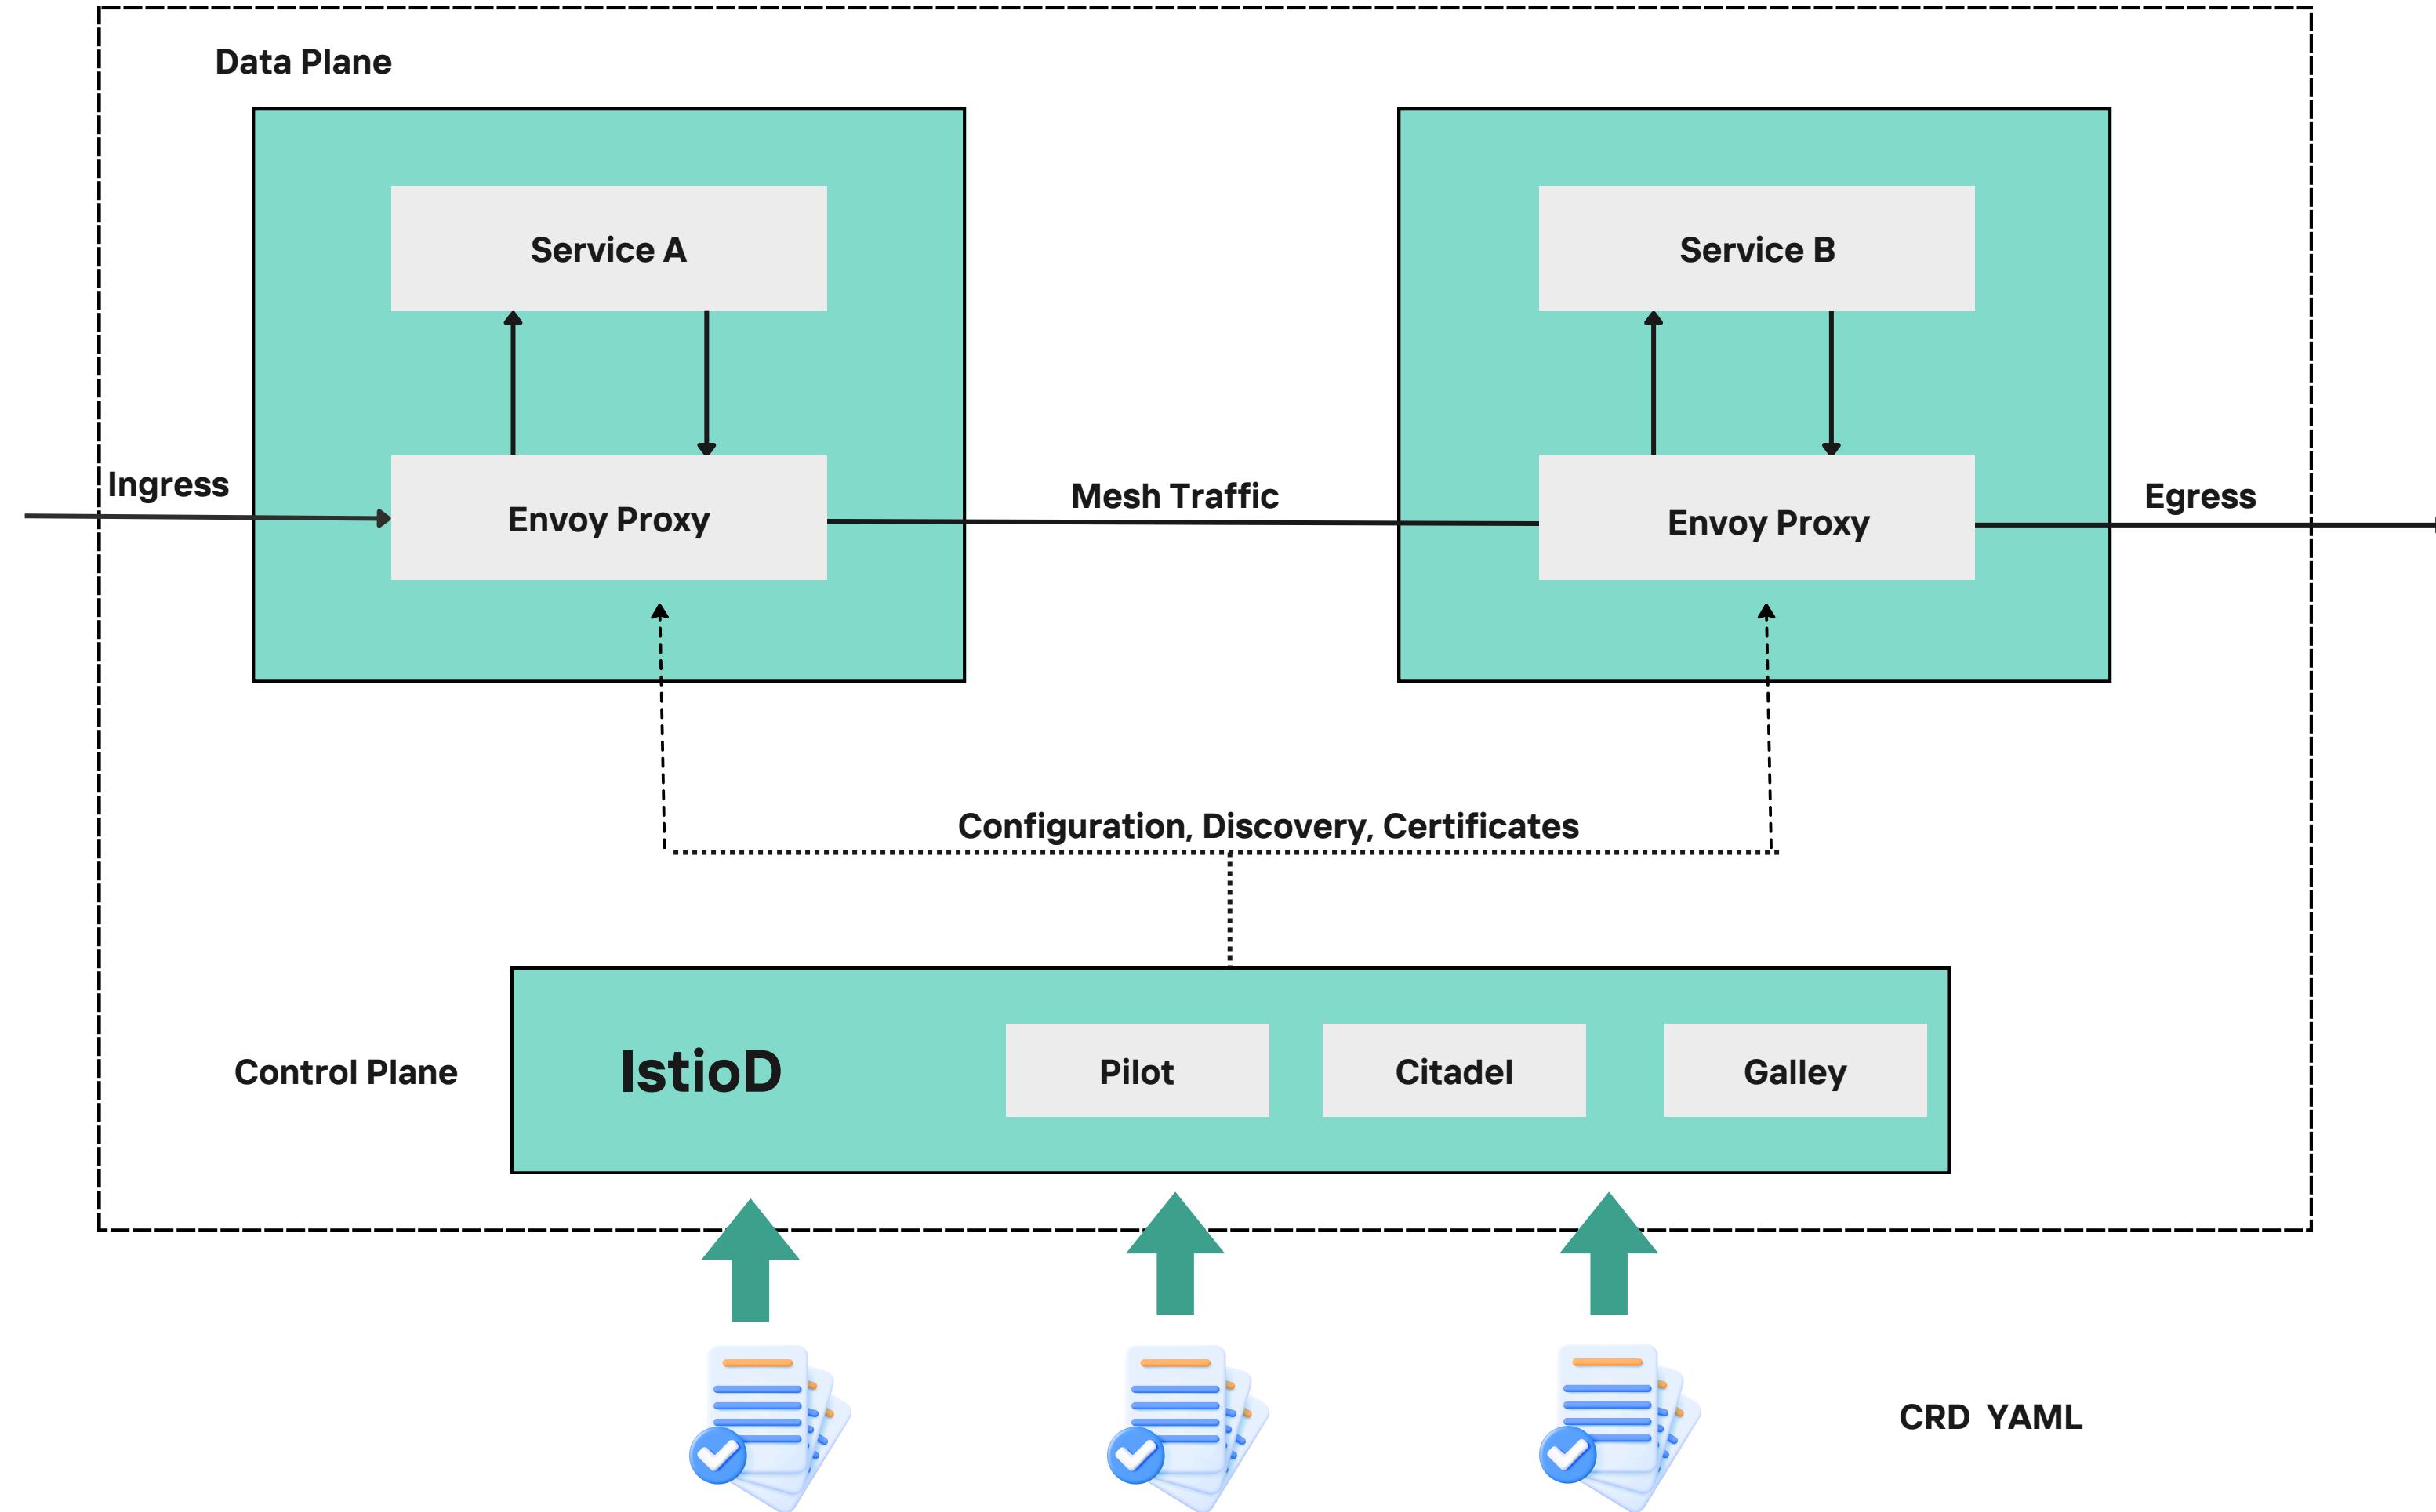
: Control the flow of traffic between services.

Istio provides powerful traffic management features, including routing, load balancing, retries, and circuit breaking. You can configure these settings using Istio CRDs to optimize service communication and performance.

## Observability Setup

: Enable monitoring, logging, and tracing.

Integrate Istio with observability tools like Prometheus, Grafana, Jaeger, and Kiali to gain insights into service behavior, monitor performance, and troubleshoot issues effectively.



# Kiali

Kiali provides observability, visualizations, and configuration capabilities for the Istio service mesh.

---

**Service Graph:** Visualizes the service mesh topology.

Shows how services are connected and their interactions.

---

**Health Monitoring:** Displays the health status of services.

Monitors service performance and detects anomalies.

It provides a visual representation of the service mesh topology, showing relationships between microservices.

Kiali displays real-time traffic flows, health metrics, and resource utilization within the mesh.

Overview

Traffic Graph

Applications

Workloads

Services

Istio Config

Mesh

Namespace: default ▾ | Traffic ▾ | App graph ▾

Replay

Last 1m ▾

Every 15s ▾



Display ▾

Find... ▾

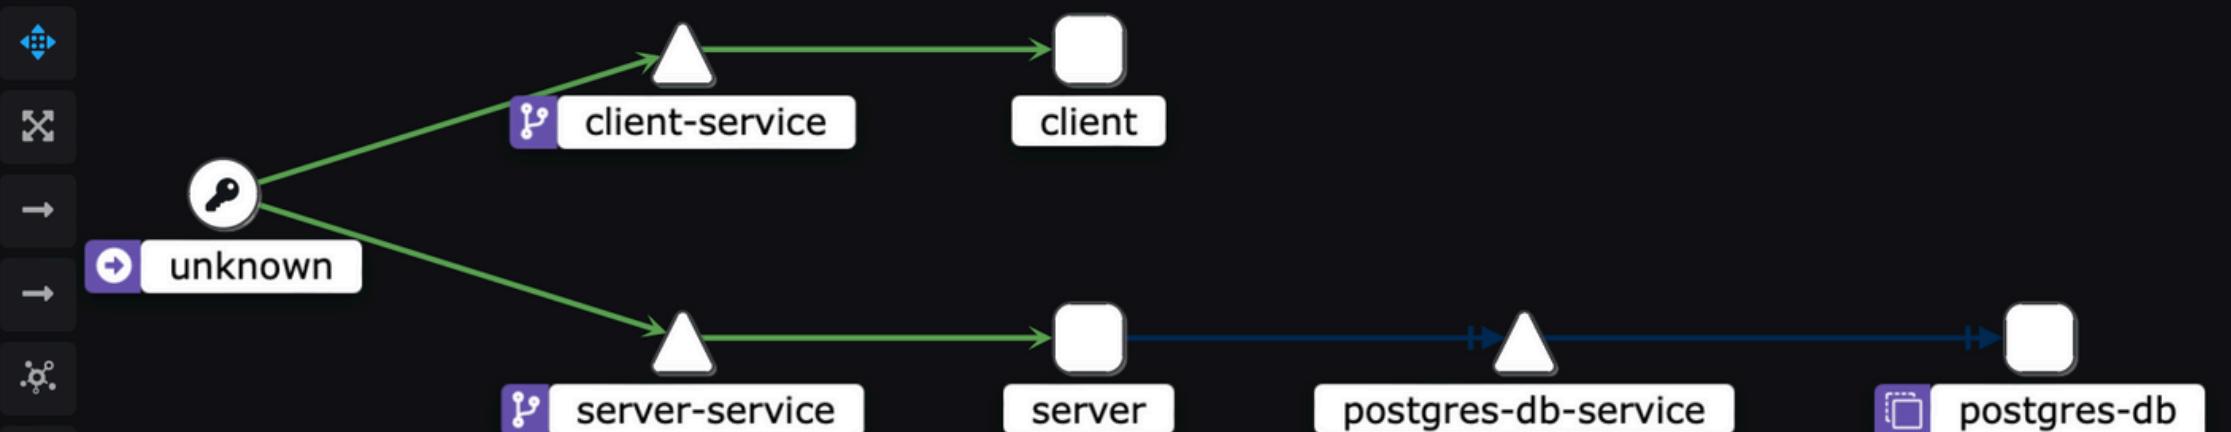
Hide... ▾



Help

Reset

Jul 11 at 08:34:28 AM ... 08:35:28 AM



## Current Graph

NS default ✓

3 apps (3 versions)

3 services

6 edges

Inbound

Outbound

Total

## HTTP (requests per second):

Total	% Success	% Error
-------	-----------	---------

0.31	100.00	0.00
------	--------	------

OK	3xx	4xx	5xx	NR
----	-----	-----	-----	----

█	█	█	█	█
---	---	---	---	---

0	25	50	75	100
---	----	----	----	-----

OK	3xx	4xx	5xx	NR
----	-----	-----	-----	----

█	█	█	█	█
---	---	---	---	---

0	25	50	75	100
---	----	----	----	-----

OK	3xx	4xx	5xx	NR
----	-----	-----	-----	----

█	█	█	█	█
---	---	---	---	---

0	25	50	75	100
---	----	----	----	-----

OK	3xx	4xx	5xx	NR
----	-----	-----	-----	----

█	█	█	█	█
---	---	---	---	---

0	25	50	75	100
---	----	----	----	-----

OK	3xx	4xx	5xx	NR
----	-----	-----	-----	----

█	█	█	█	█
---	---	---	---	---

0	25	50	75	100
---	----	----	----	-----

OK	3xx	4xx	5xx	NR
----	-----	-----	-----	----

█	█	█	█	█
---	---	---	---	---

0	25	50	75	100
---	----	----	----	-----

OK	3xx	4xx	5xx	NR
----	-----	-----	-----	----

█	█	█	█	█
---	---	---	---	---

0	25	50	75	100
---	----	----	----	-----

OK	3xx	4xx	5xx	NR
----	-----	-----	-----	----

█	█	█	█	█
---	---	---	---	---

0	25	50	75	100
---	----	----	----	-----

OK	3xx	4xx	5xx	NR
----	-----	-----	-----	----

█	█	█	█	█
---	---	---	---	---

0	25	50	75	100
---	----	----	----	-----

OK	3xx	4xx	5xx	NR
----	-----	-----	-----	----

█	█	█	█	█
---	---	---	---	---

0	25	50	75	100
---	----	----	----	-----

OK	3xx	4xx	5xx	NR
----	-----	-----	-----	----

█	█	█	█	█
---	---	---	---	---

0	25	50	75	100
---	----	----	----	-----

OK	3xx	4xx	5xx	NR
----	-----	-----	-----	----

█	█	█	█	█
---	---	---	---	---

0	25	50	75	100
---	----	----	----	-----

OK	3xx	4xx	5xx	NR
----	-----	-----	-----	----

█	█	█	█	█
---	---	---	---	---

0	25	50	75	100
---	----	----	----	-----

OK	3xx	4xx	5xx	NR
----	-----	-----	-----	----

█	█	█	█	█
---	---	---	---	---

0	25	50	75	100
---	----	----	----	-----

OK	3xx	4xx	5xx	NR
----	-----	-----	-----	----

█	█	█	█	█
---	---	---	---	---

0	25	50	75	100
---	----	----	----	-----

OK	3xx	4xx	5xx	NR
----	-----	-----	-----	----

█	█	█	█	█
---	---	---	---	---

0	25	50	75	100
---	----	----	----	-----

OK	3xx	4xx	5xx	NR
----	-----	-----	-----	----

█	█	█	█	█
---	---	---	---	---

0	25	50	75	100
---	----	----	----	-----

OK	3xx	4xx	5xx	NR
----	-----	-----	-----	----

█	█	█	█	█
---	---	---	---	---

0	25	50	75	100
---	----	----	----	-----

OK	3xx	4xx	5xx	NR
----	-----	-----	-----	----

█	█	█	█	█
---	---	---	---	---

0	25	50	75	100
---	----	----	----	-----

OK	3xx	4xx	5xx	NR
----	-----	-----	-----	----

█	█	█	█	█
---	---	---	---	---

0	25	50	75	100
---	----	----	----	-----

OK	3xx	4xx	5xx	NR
----	-----	-----	-----	----

█	█	█	█	█
---	---	---	---	---

0	25	50	75	100
---	----	----	----	-----

OK	3xx	4xx	5xx	NR
----	-----	-----	-----	----

█	█	█	█	█
---	---	---	---	---

0	25	50	75	100
---	----	----	----	-----

OK	3xx	4xx	5xx	NR
----	-----	-----	-----	----

█	█	█	█	█
---	---	---	---	---

0	25	50	75	100
---	----	----	----	-----

OK	3xx	4xx	5xx	NR
----	-----	-----	-----	----

█	█	█	█	█
---	---	---	---	---

0	25	50	75	100
---	----	----	----	-----

OK	3xx	4xx	5xx	NR
----	-----	-----	-----	----

█	█	█	█	█
---	---	---	---	---

0	25	50	75	100
---	----	----	----	-----

OK	3xx	4xx	5xx	NR
----	-----	-----	-----	----

█	█	█	█	█
---	---	---	---	---

0	25	50	75	100
---	----	----	----	-----

OK	3xx	4xx	5xx	NR
----	-----	-----	-----	----

█	█	█	█	█
---	---	---	---	---

0	25	50	75	100
---	----	----	----	-----

OK	3xx	4xx	5xx	NR
----	-----	-----	-----	----

# Jaeger

Jaeger is a tool for monitoring and troubleshooting transactions in complex distributed systems.

---

**Trace Collection:** Captures detailed tracing information.

Collects data about individual requests and their journey through the services.

---

**Root Cause Analysis:** Identifies performance issues and bottlenecks.

Helps pinpoint where latency and failures occur.

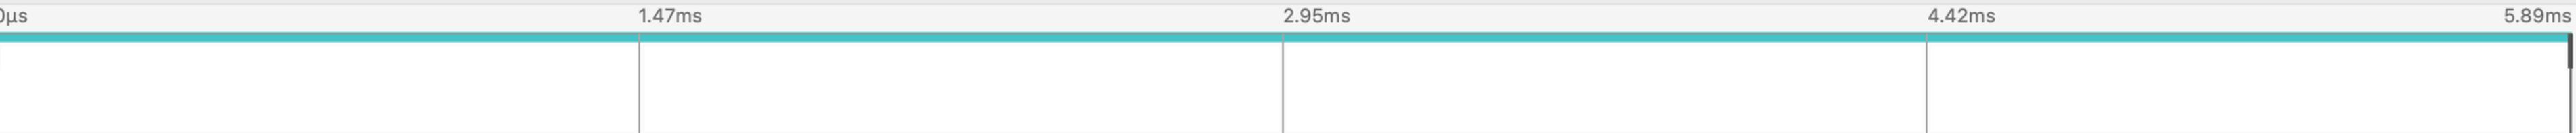
Istio automatically injects tracing information into service calls, allowing Jaeger to capture the complete request flow.

## server.default: server-service.default.svc.cluster.local: 8000/\* f6d3600

Find...



Trace Timeline ▾

Trace Start **July 11 2024, 08:12:55.537** | Duration **5.89ms** | Services **1** | Depth **1** | Total Spans **1**

### Service & Operation

▼ > ▼ >> ||

0µs

1.47ms

2.95ms

4.42ms

5.89ms

server.default server-service.default.svc.cluster.local:80...

### server-service.default.svc.cluster.local:8000/\*

Service: **server.default** | Duration: **5.89ms** | Start Time: **0µs**

#### ▼ Tags

internal.span.format	zipkin
istio.canonical_revision	latest
istio.canonical_service	server
istio.cluster_id	Kubernetes
istio.mesh_id	cluster.local
istio.namespace	default
net.host.ip	10.244.0.74
node_id	sidecar~10.244.0.74~server-deployment-58bb4fc485-45qh2.default~default.svc.cluster.local
peer.address	10.244.0.1
request_size	0
response_flags	-
response_size	326
span.kind	server

# Grafana

Grafana is a powerful tool for creating dashboards and visualizing metrics collected from Prometheus.

---

**Custom Dashboards:** Build interactive and customizable dashboards.

Provides real-time insights into system performance and resource usage.

---

**Alerting:** Set up alerts based on specific metrics.

Notifies when certain thresholds are reached, enabling proactive management.

Grafana allows us to create custom dashboards that visualize these Istio metrics alongside tracing data from Jaeger.



Q Search or jump to...

cmd+k

+ | ? | ⌂ | Sign in

Home > Dashboards > istio > Istio Control Plane Dashboard



Add

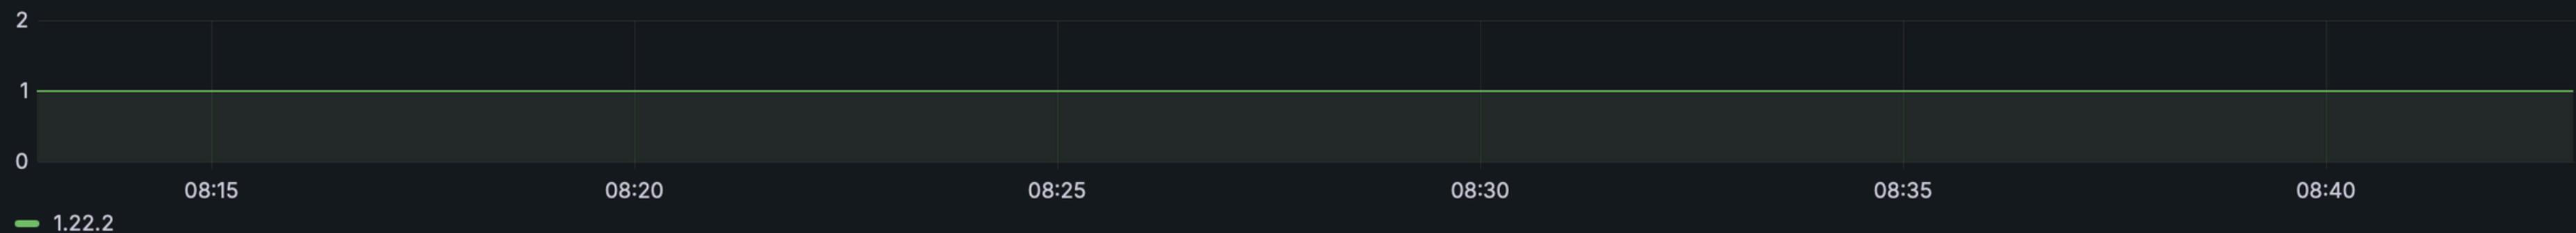
Share

Last 30 minutes



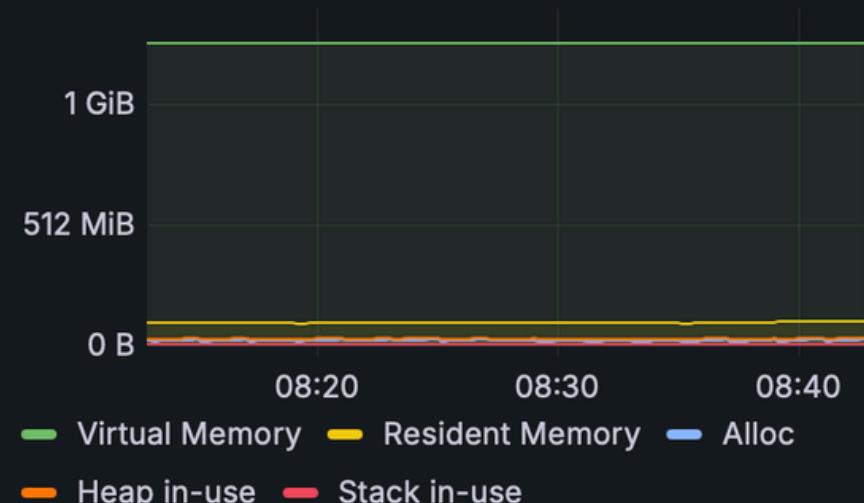
## Deployed Versions

### Pilot Versions

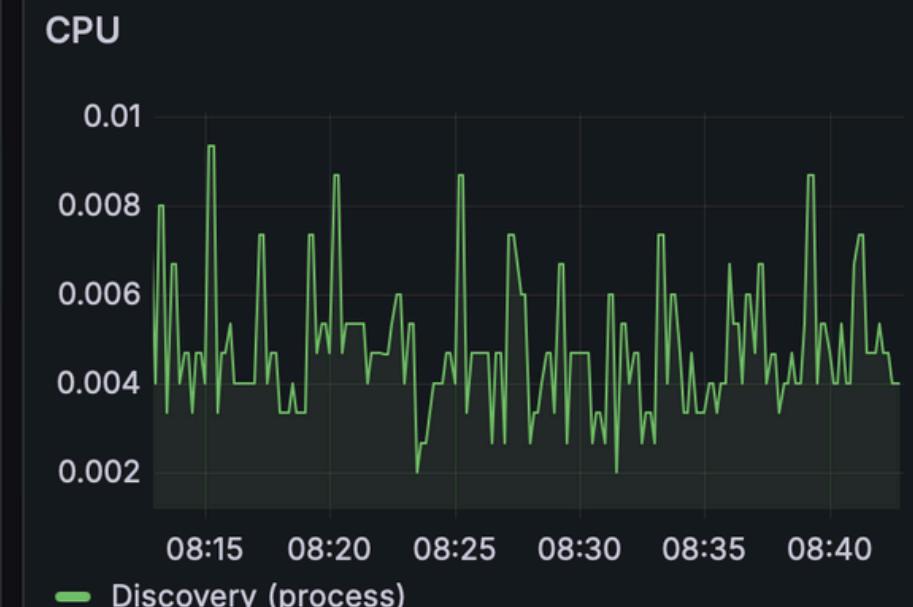


## Resource Usage

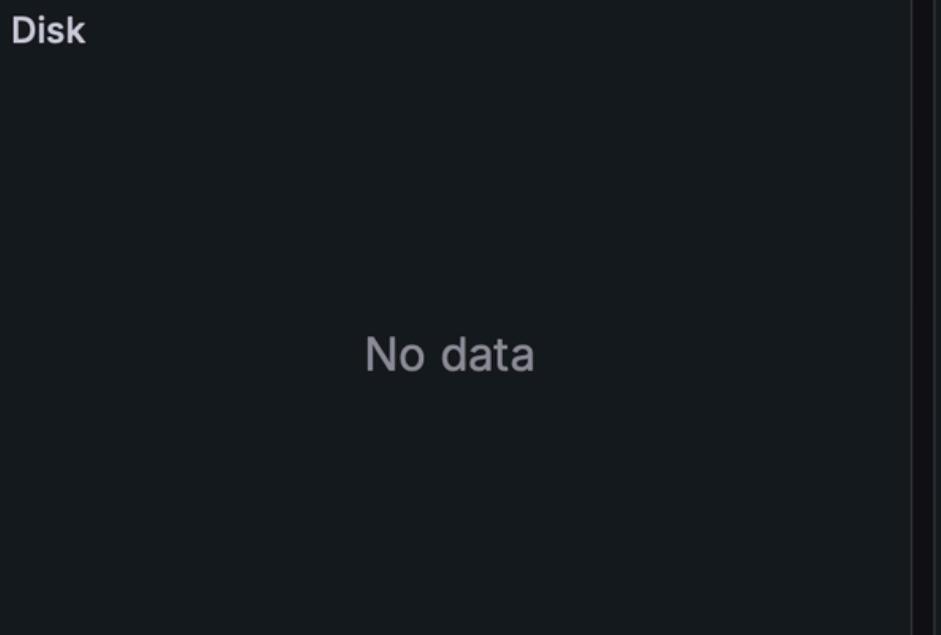
### Memory



### CPU



### Disk

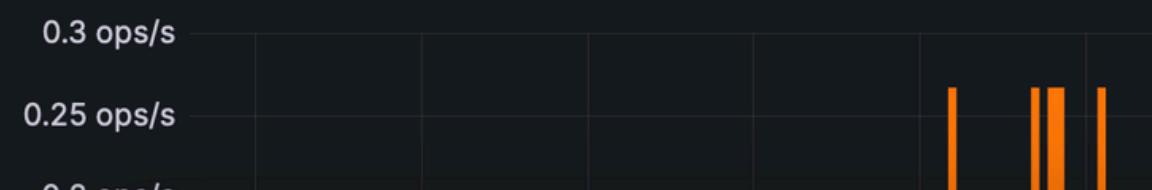


### Goroutines

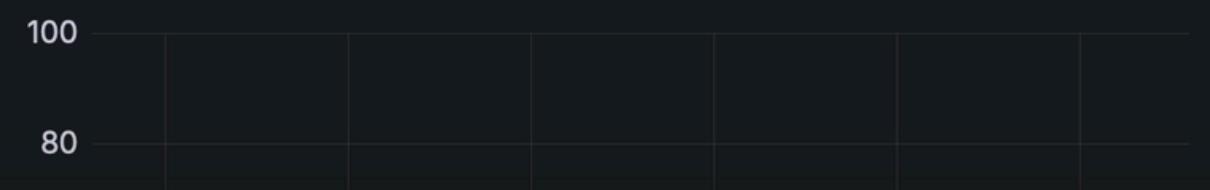


## Pilot Push Information

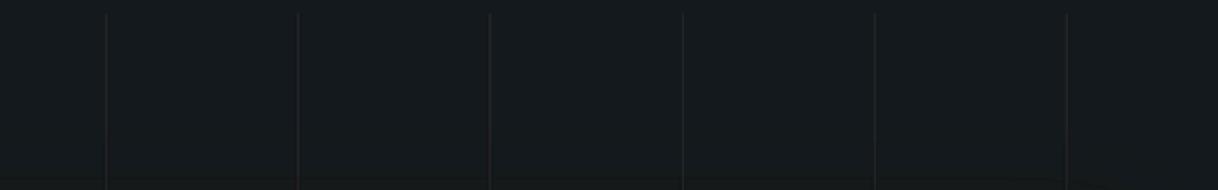
### Pilot Pushes



### Pilot Errors



### Proxy Push Time

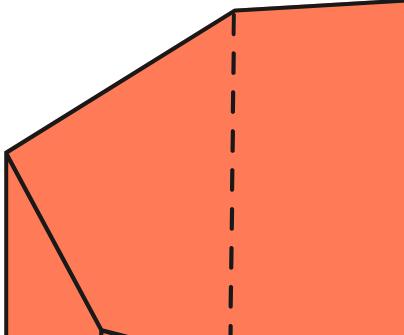
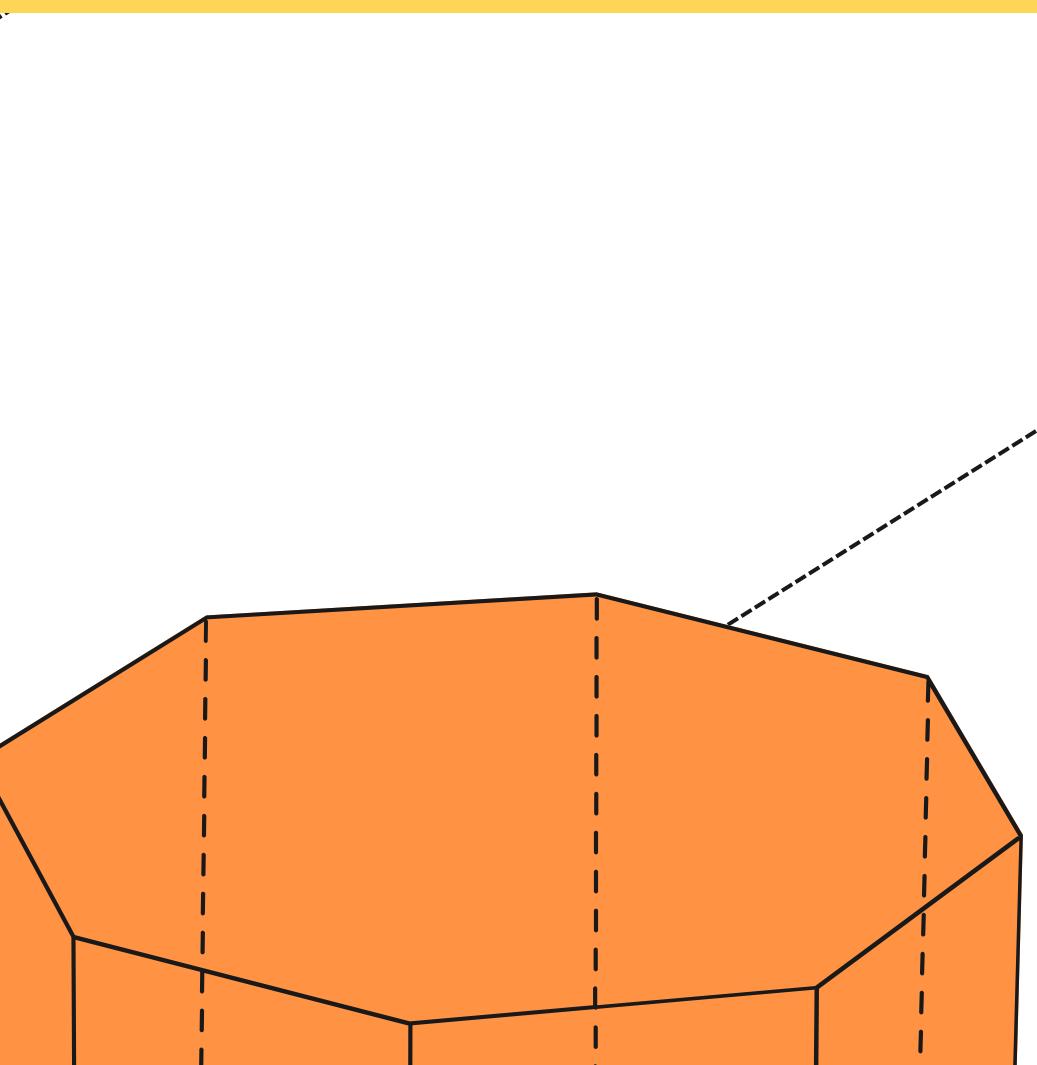
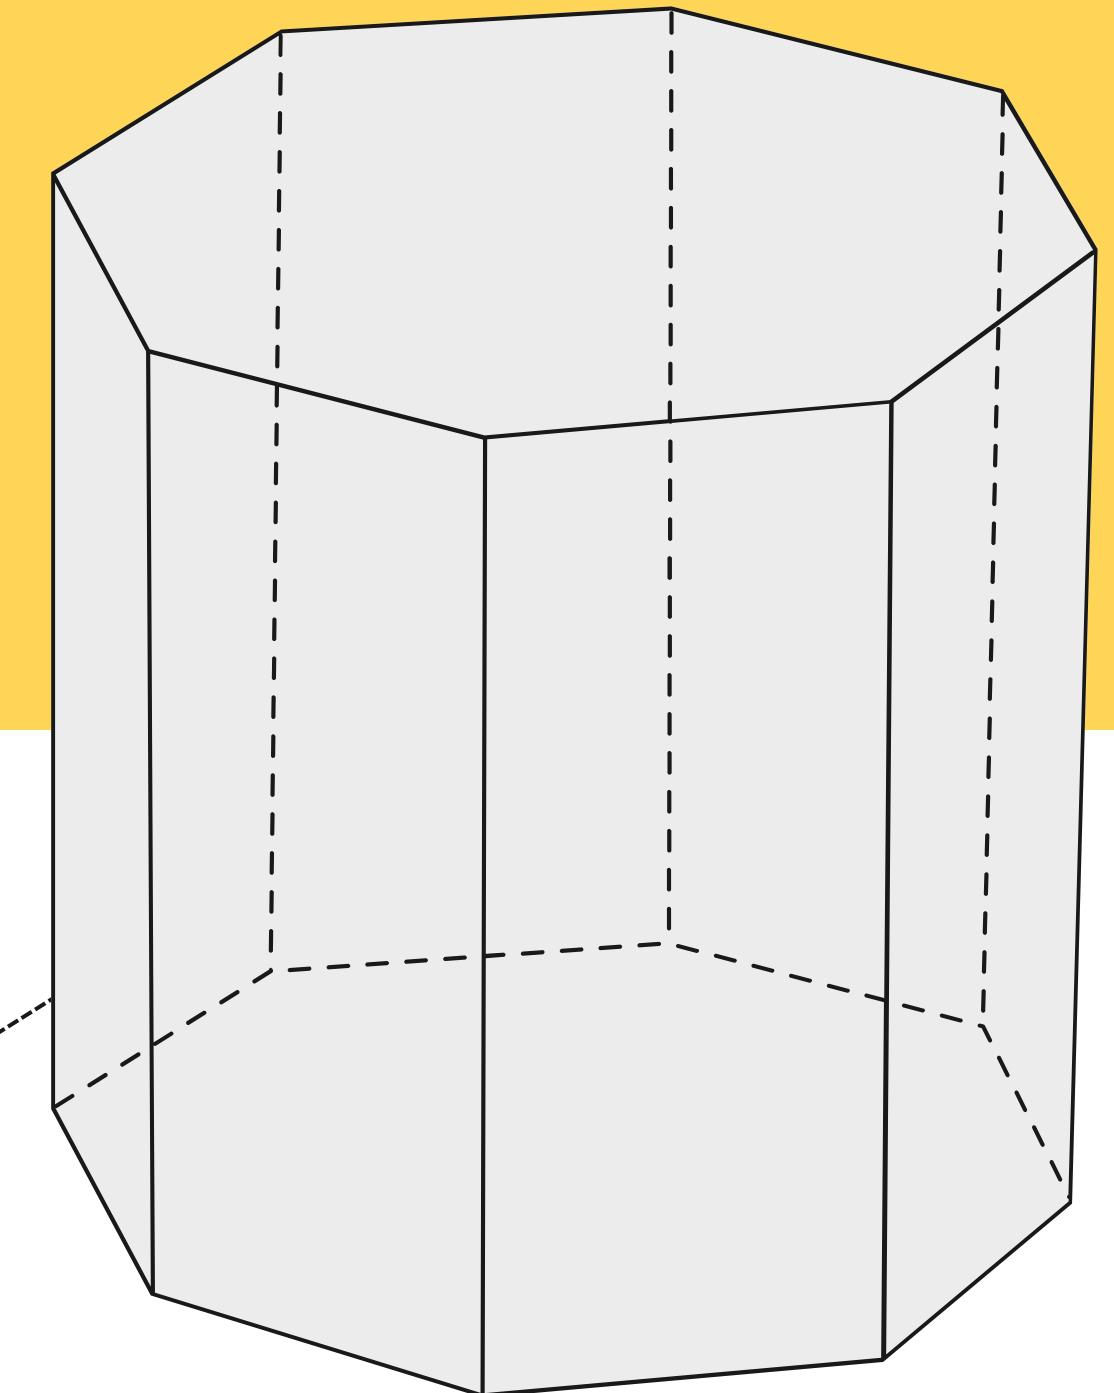


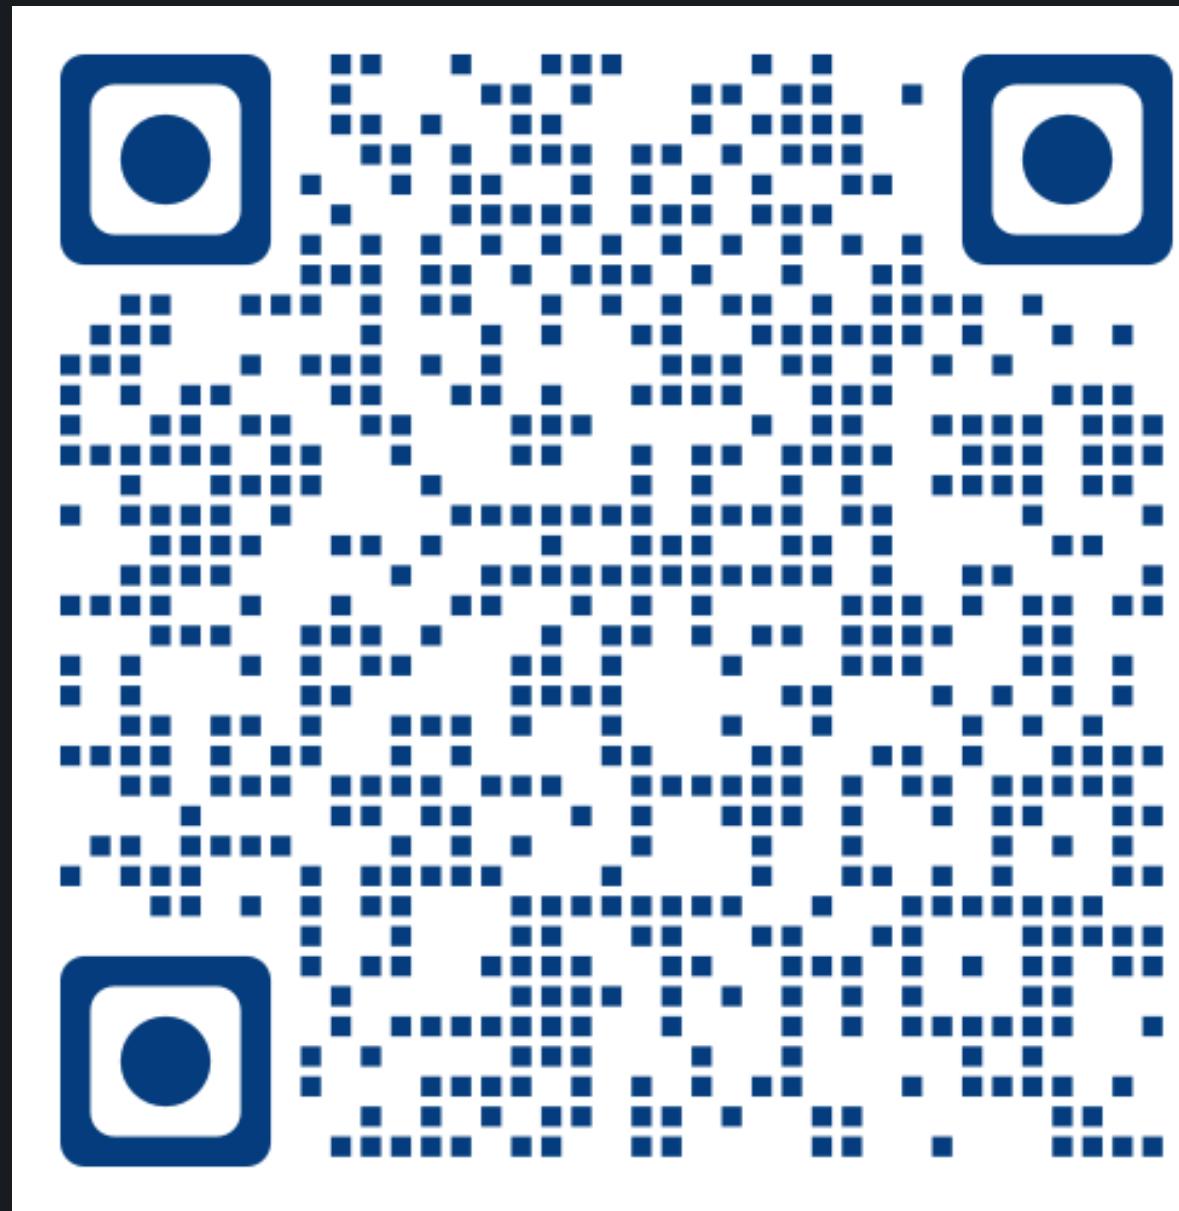
# THANK YOU!



@n4jp4y

<http://n4jp4y.com>





For presentation  
and Code Repository



<http://n4jp4y.com>